

# Computer Networks and Internet Technology

2021W703033 VO Rechnernetze und Internettechnik  
Winter Semester 2021/22

Jan Beutel

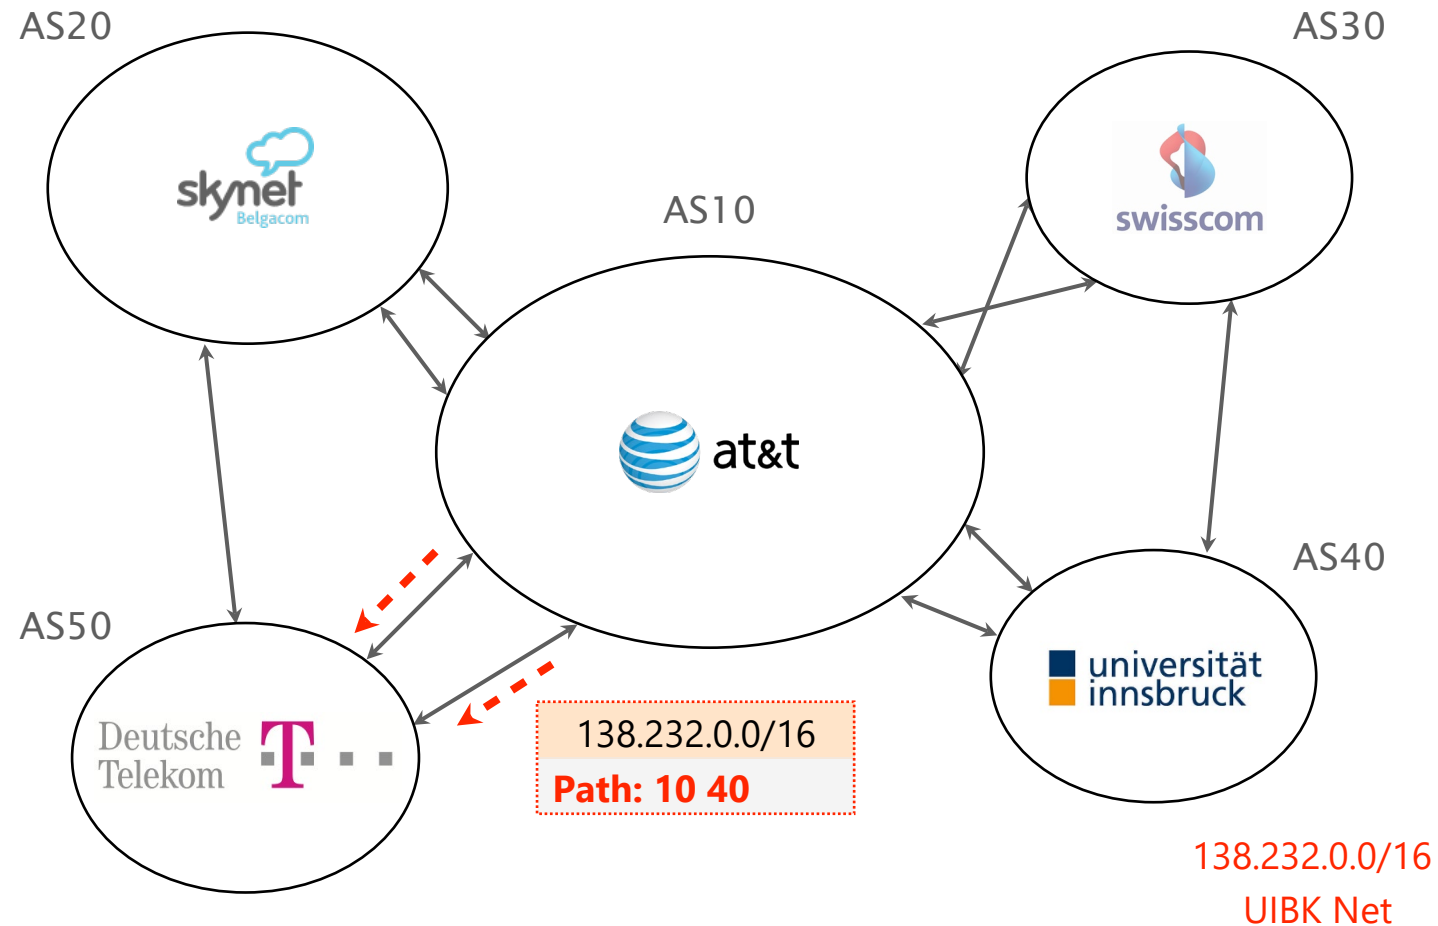
# Communication Networks and Internet Technology

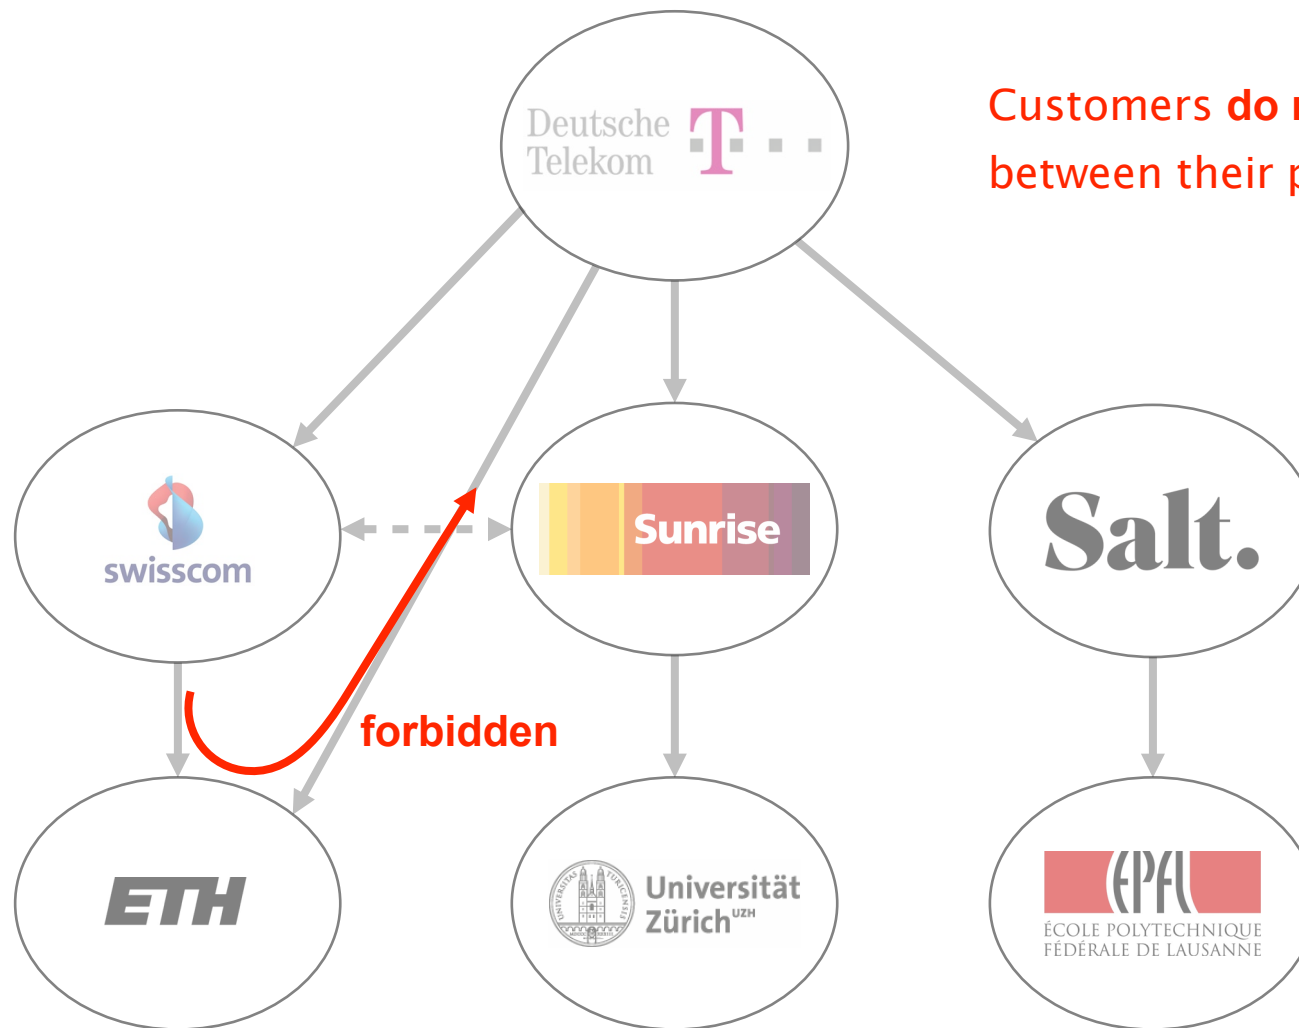
**Recap of last weeks lecture**

BGP relies on **path-vector routing** to support flexible routing policies and avoid count-to-infinity

key idea      advertise the **entire path** instead of distances

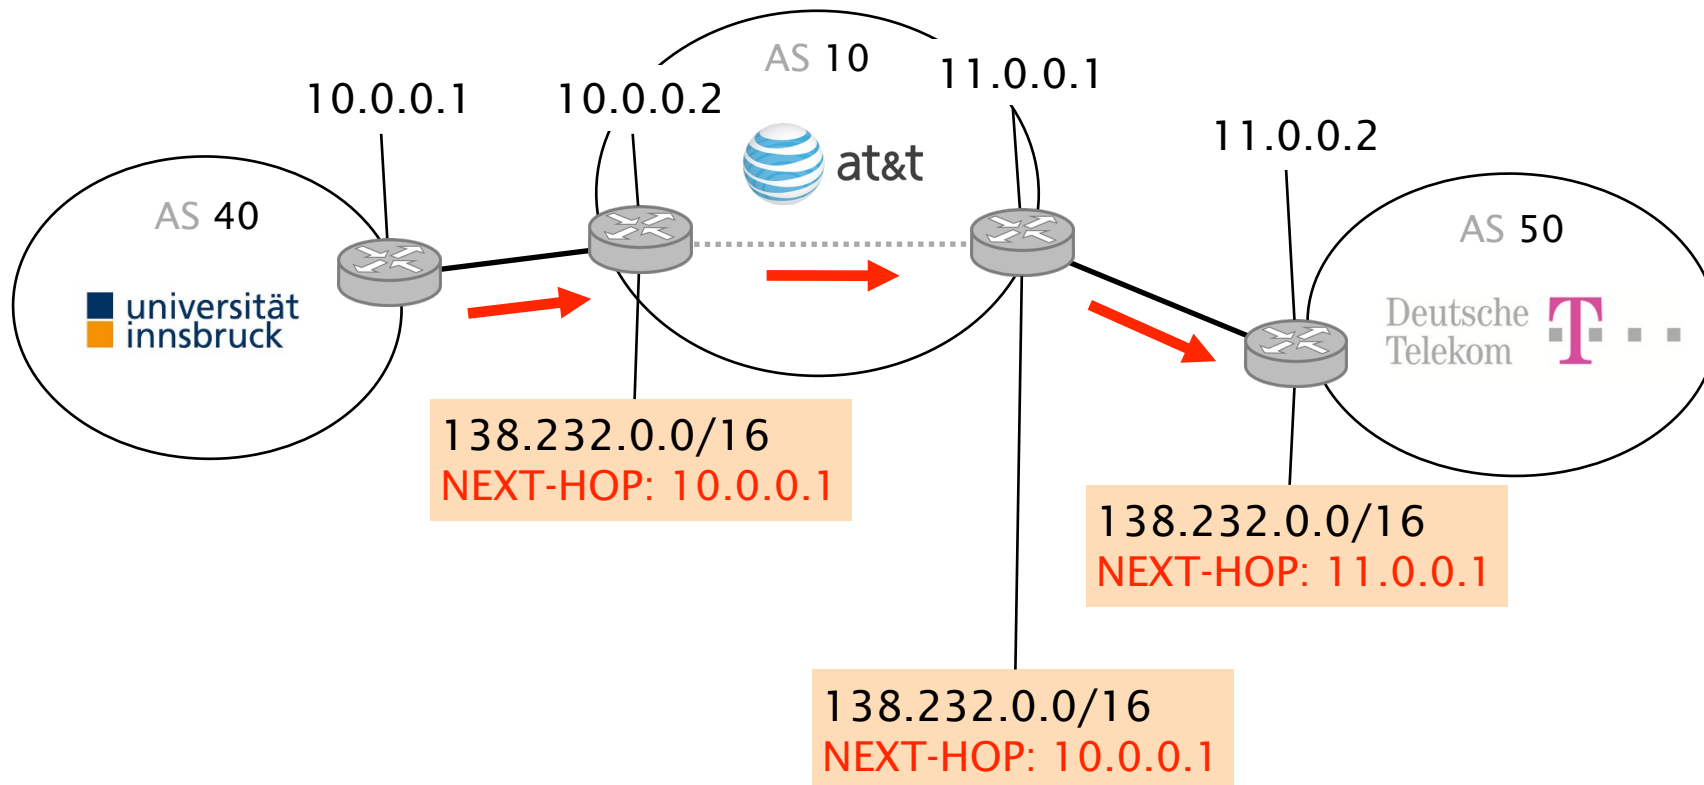
Each AS appends itself to the path  
when it propagates announcements





Customers **do not** transit traffic  
between their providers


The NEXT-HOP is set when the route enters an AS,  
it does **not** change within the AS



# Communication Networks and Internet Technology

**This weeks lecture**

We're continuing our journey up the layers,  
now looking at **the transport layer**





UDP / TCP

Congestion  
Control

Un/reliable Transport



The diagram consists of two rectangular boxes. The box on the left is light green and contains the text 'UDP / TCP'. The box on the right is light orange and contains the text 'Congestion Control'. Below the green box is the text 'Un/reliable Transport' in a grey font. There are no arrows or other graphical elements connecting the boxes.

UDP / TCP

Congestion  
Control

Un/reliable Transport

# What do we need in the Transport layer?

- Functionality implemented in **network**
  - Keep minimal (easy to build, broadly applicable)
- Functionality implemented in the **application**
  - Keep minimal (easy to write)
  - Restricted to application-specific functionality
- Functionality implemented in the “**network stack**”
  - The shared networking code on the host
  - This relieves burden from both app and network
  - **The transport layer is a key component here**

# What do we need in the Transport layer?

- **Application layer**
  - Communication for specific applications
  - e.g., HyperText Transfer Protocol (HTTP),  
File Transfer Protocol (FTP)
- **Network layer**
  - Global communication between hosts
  - Hides details of the link technology
  - e.g., Internet Protocol (IP)

# What Problems Should Be Solved Here?

- Data delivering, to the *correct* application
  - IP just points towards next protocol
  - *Transport needs to demultiplex incoming data (ports)*
- Files or bytestreams abstractions for the applications
  - Network deals with packets
  - *Transport layer needs to translate between them*
- Reliable transfer (if needed)
- Not overloading the receiver
- Not overloading the network

# What Is Needed to Address These?

- *Demultiplexing*: identifier for application process
  - Going from host-to-host (IP) to process-to-process
- *Translating between bytestreams and packets*:
  - Do segmentation and reassembly
- *Reliability*: ACKs and all that stuff
- *Corruption*: Checksum
- *Not overloading receiver*: “Flow Control”
  - Limit data in receiver’s buffer
- *Not overloading network*: “Congestion Control”

# UDP: Datagram messaging service

UDP provides a **connectionless**, **unreliable** transport service

- No-frills extension of “best-effort” IP
- UDP provides **only two services** to the App layer
  - Multiplexing/Demultiplexing among processes
  - Discarding corrupted packets (optional)

# TCP: Reliable, in-order delivery

- TCP provides a connection-oriented, reliable, bytestream transport service
- ***What UDP provides, plus:***
  - Retransmission of lost and corrupted packets
  - Flow control (to not overflow receiver)
  - Congestion control (to not overload network)
  - “Connection” set-up & tear-down



# Connections (or sessions)

- Reliability requires keeping state
  - Sender: packets sent but not ACKed, and related timers
  - Receiver: noncontiguous packets
- Each bytestream is called a connection or session
  - Each with their own connection state
  - State is in hosts, not network!

# What transport protocols do **not** provide

- Delay and/or bandwidth guarantees
  - This cannot be offered by transport
  - Requires support at IP level (*and let's not go there*)
- Sessions that survive change-of-IP-address
  - This is an artifact of current implementations
  - As we shall see....

# Important Context: Sockets and Ports

- **Sockets:** an operating system abstraction
- **Ports:** a networking abstraction
  - This is not a port on a switch (which is an interface)
  - Think of it as a *logical interface* on a host

# Sockets

- A socket is a software abstraction by which an application process exchanges network messages with the (transport layer in the) operating system
  - `socketID = socket(..., socket.TYPE)`
  - `socketID.sendto(message, ...)`
  - `socketID.recvfrom(...)`
- Two important types of sockets
  - UDP socket: TYPE is `SOCK_DGRAM`
  - TCP socket: TYPE is `SOCK_STREAM`

# Ports

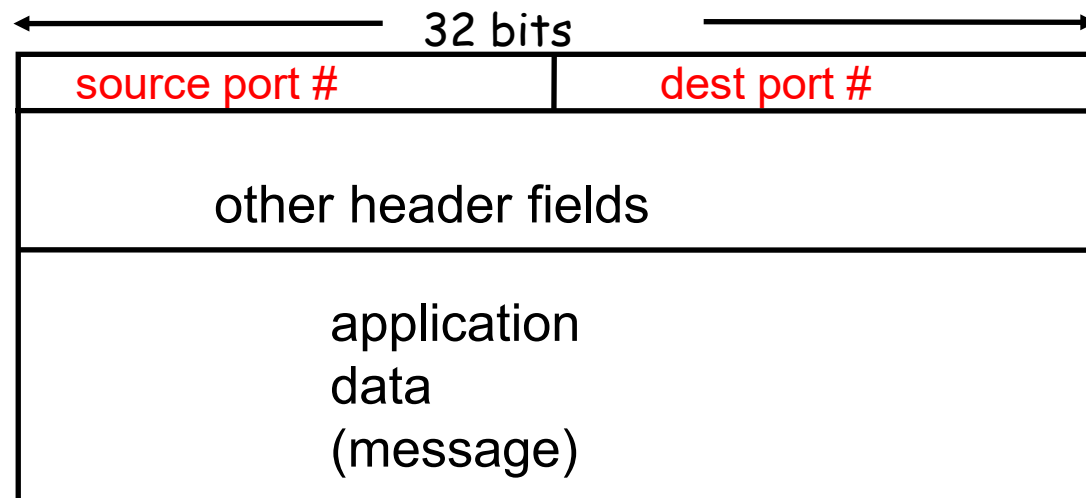
- **Problem:** which app (socket) gets which packets
- **Solution:** port as transport layer identifier (16 bits)
  - Packet carries source/destination port numbers in transport header
- OS stores mapping between sockets and ports
  - Port: in packets
  - Socket: in OS

# More on Ports

- Separate 16-bit port address space for UDP, TCP
- “Well known” ports (0-1023)
  - Agreement on which services run on these ports
  - *e.g.*, ssh:22, http:80
  - Client (app) knows appropriate port on server
  - Services can listen on well-known port
- Ephemeral ports (most 1024-65535):
  - Given to clients (at random)

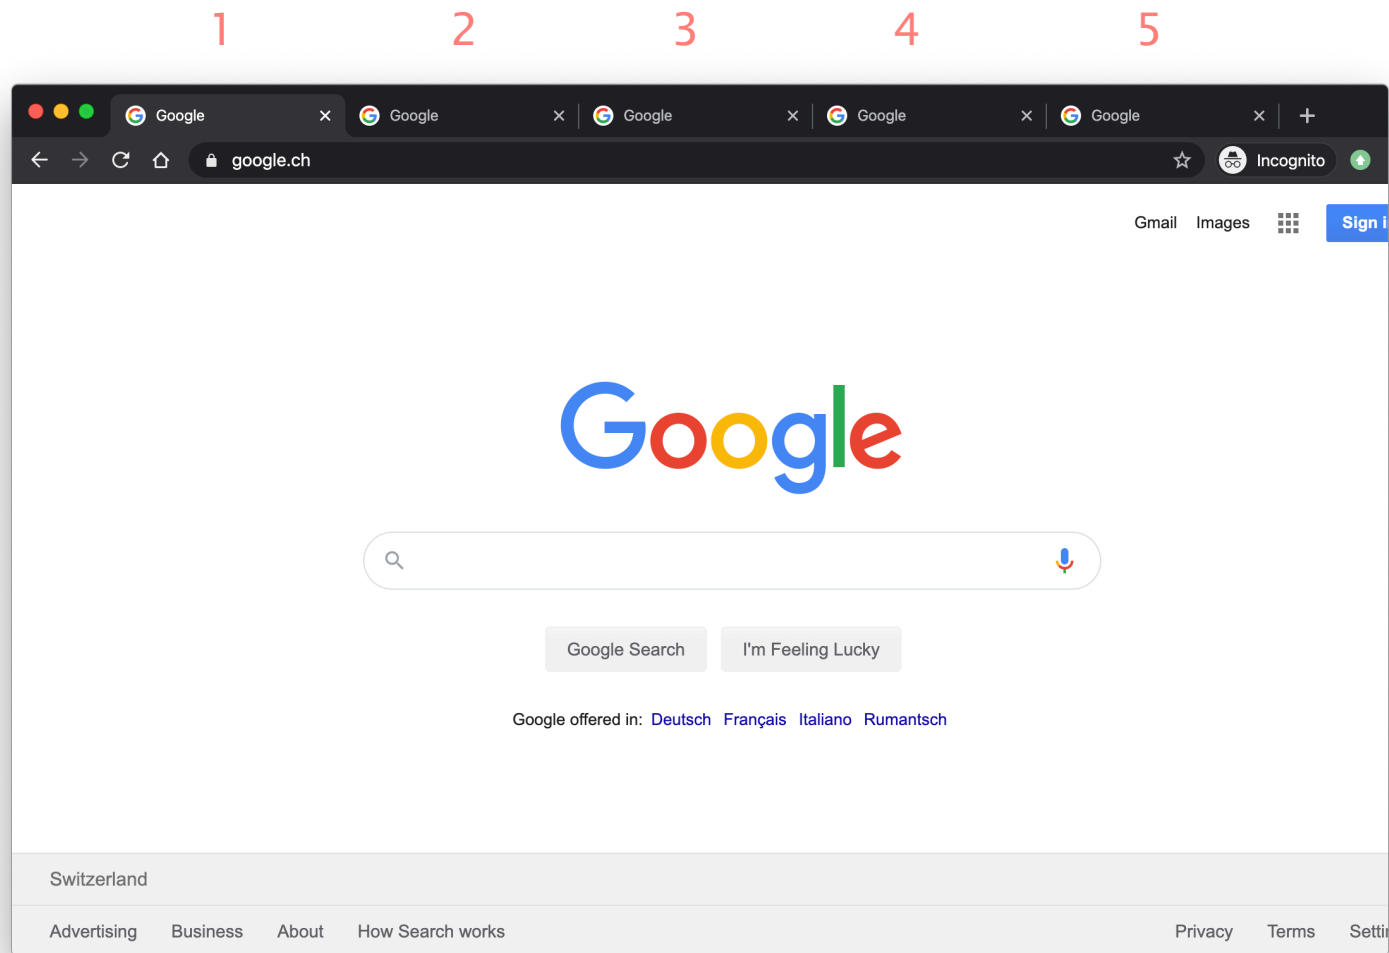
# Multiplexing and Demultiplexing

- Host receives IP datagrams
  - Each datagram has source and destination IP address,
  - Each segment has source and destination port number
- Host uses IP addresses and port numbers to direct the segment to appropriate socket



A TCP/UDP socket is identified by a 4-tuple:  
(src IP, src port, dst IP, dest port)





Let's say you open 5 tabs to [google.ch](https://google.ch)

Your IP: 129.132.19.1

Google's IP: 172.217.168.3

Client OS

src IP

src port

dest IP

dest port

socket

1

129.132.19.1

54001

172.217.168.3

443

2

129.132.19.1

55240

172.217.168.3

443

3

129.132.19.1

48472

172.217.168.3

443

4

129.132.19.1

35456

172.217.168.3

443

5

129.132.19.1

42001

172.217.168.3

443



Server OS

src IP

src port

dest IP

dest port

socket

1

172.217.168.3

443

129.132.19.1

54001

2

172.217.168.3

443

129.132.19.1

55240

3

172.217.168.3

443

129.132.19.1

48472

4

172.217.168.3

443

129.132.19.1

35456

5

172.217.168.3


443

129.132.19.1

42001



# The life of a TCP connection is a sequence of states, described with a Finite State Machine



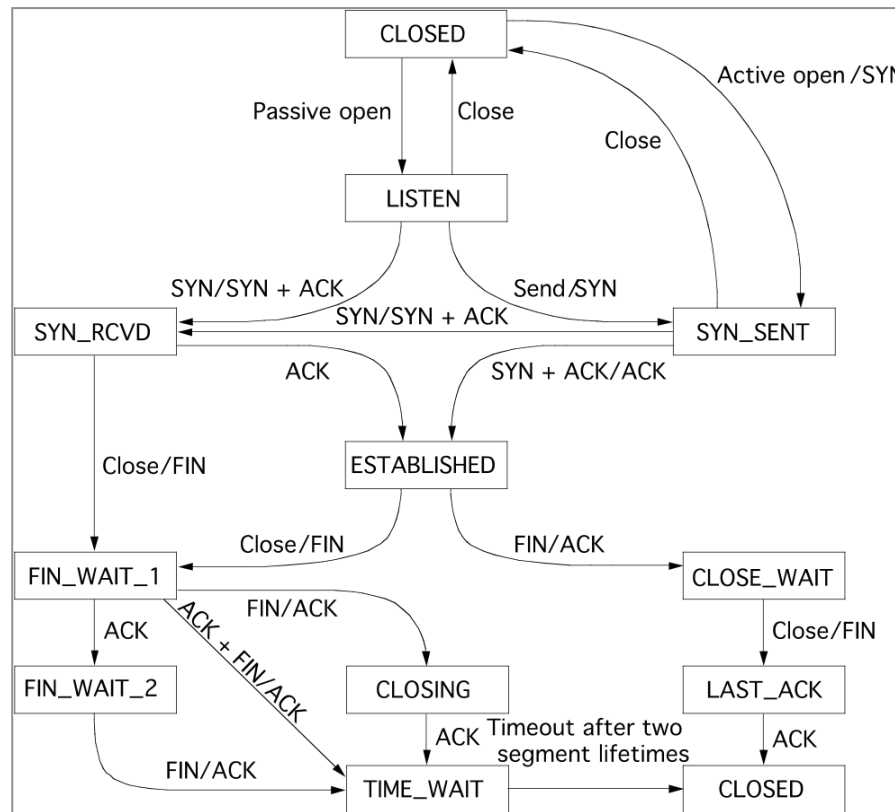
TCP connections start/end in the **CLOSED** state

Most of states relate to

- the connection establishment (three-way handshake)
- the connection termination (ensuring reliability)

Data is exchanged in the **ESTABLISHED** state

The TCP connection moves from one state to another in response of events (timeouts, "flagged" segments, ...)

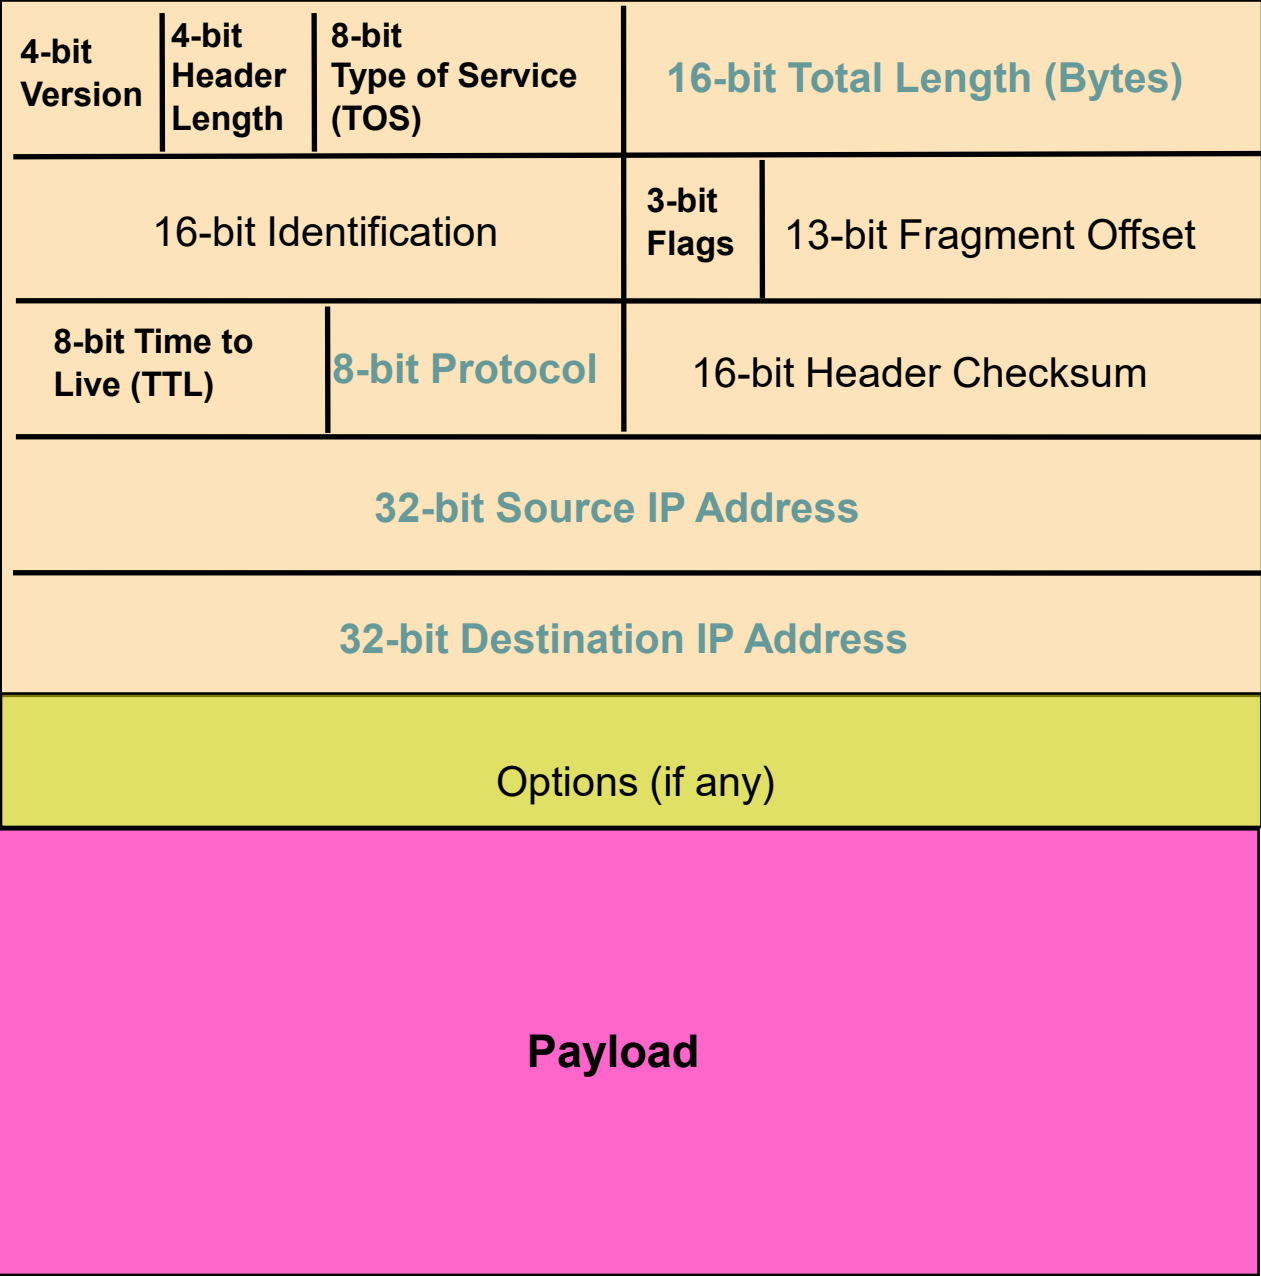


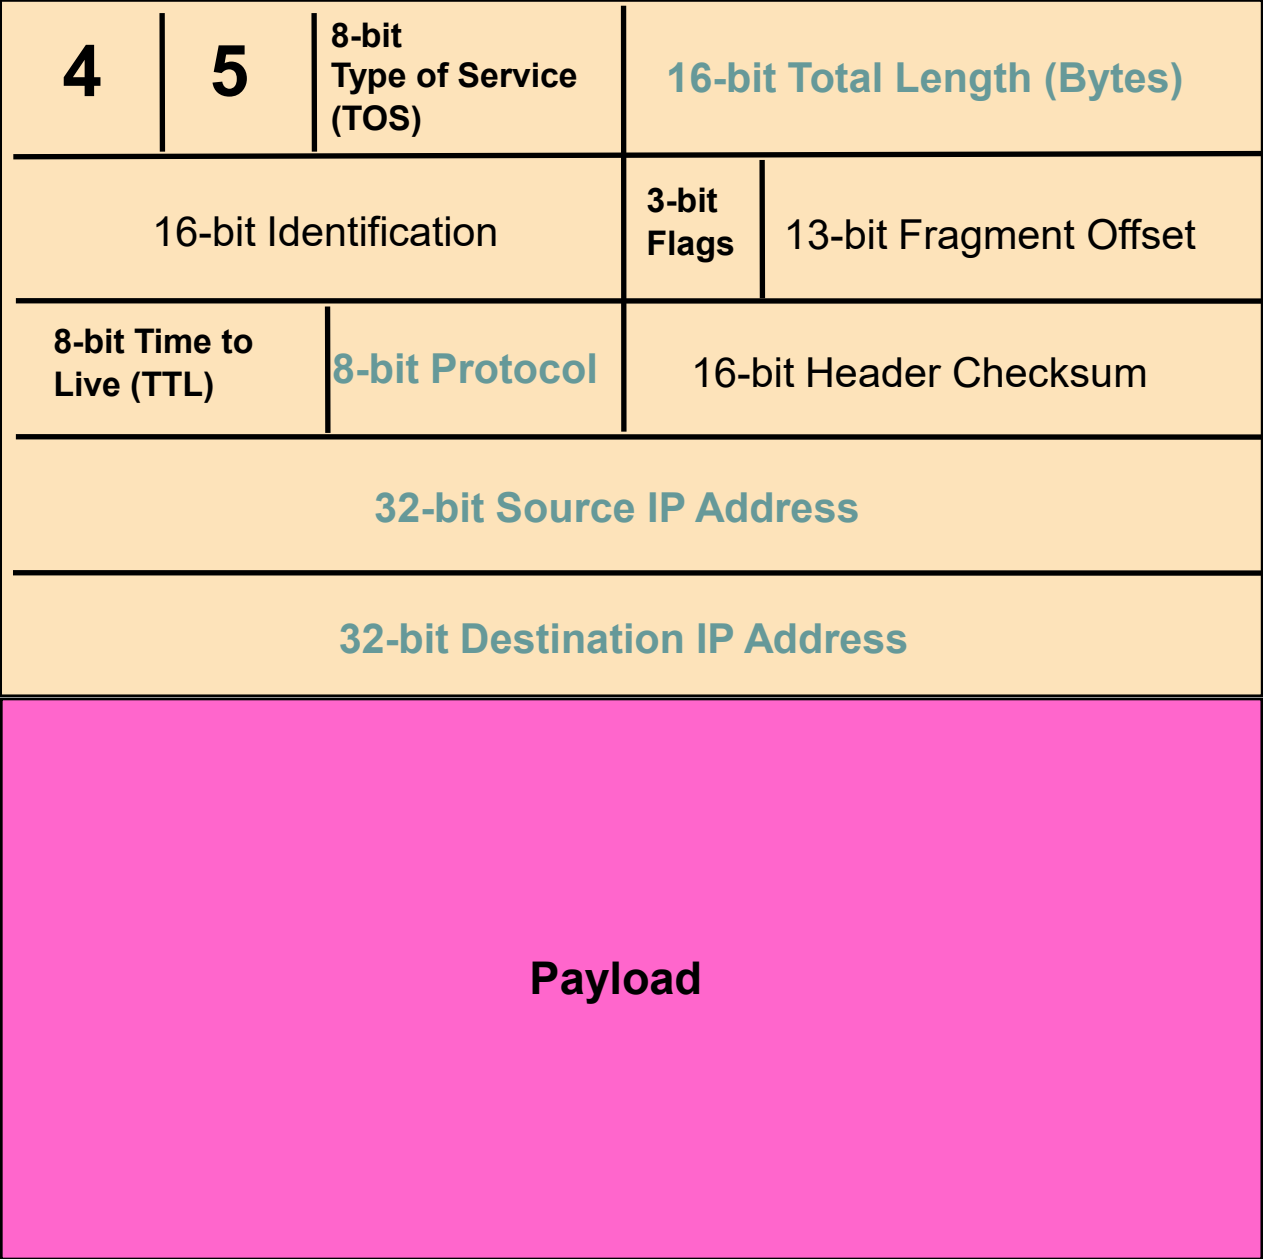
TCP connections start/end in the **CLOSED** state

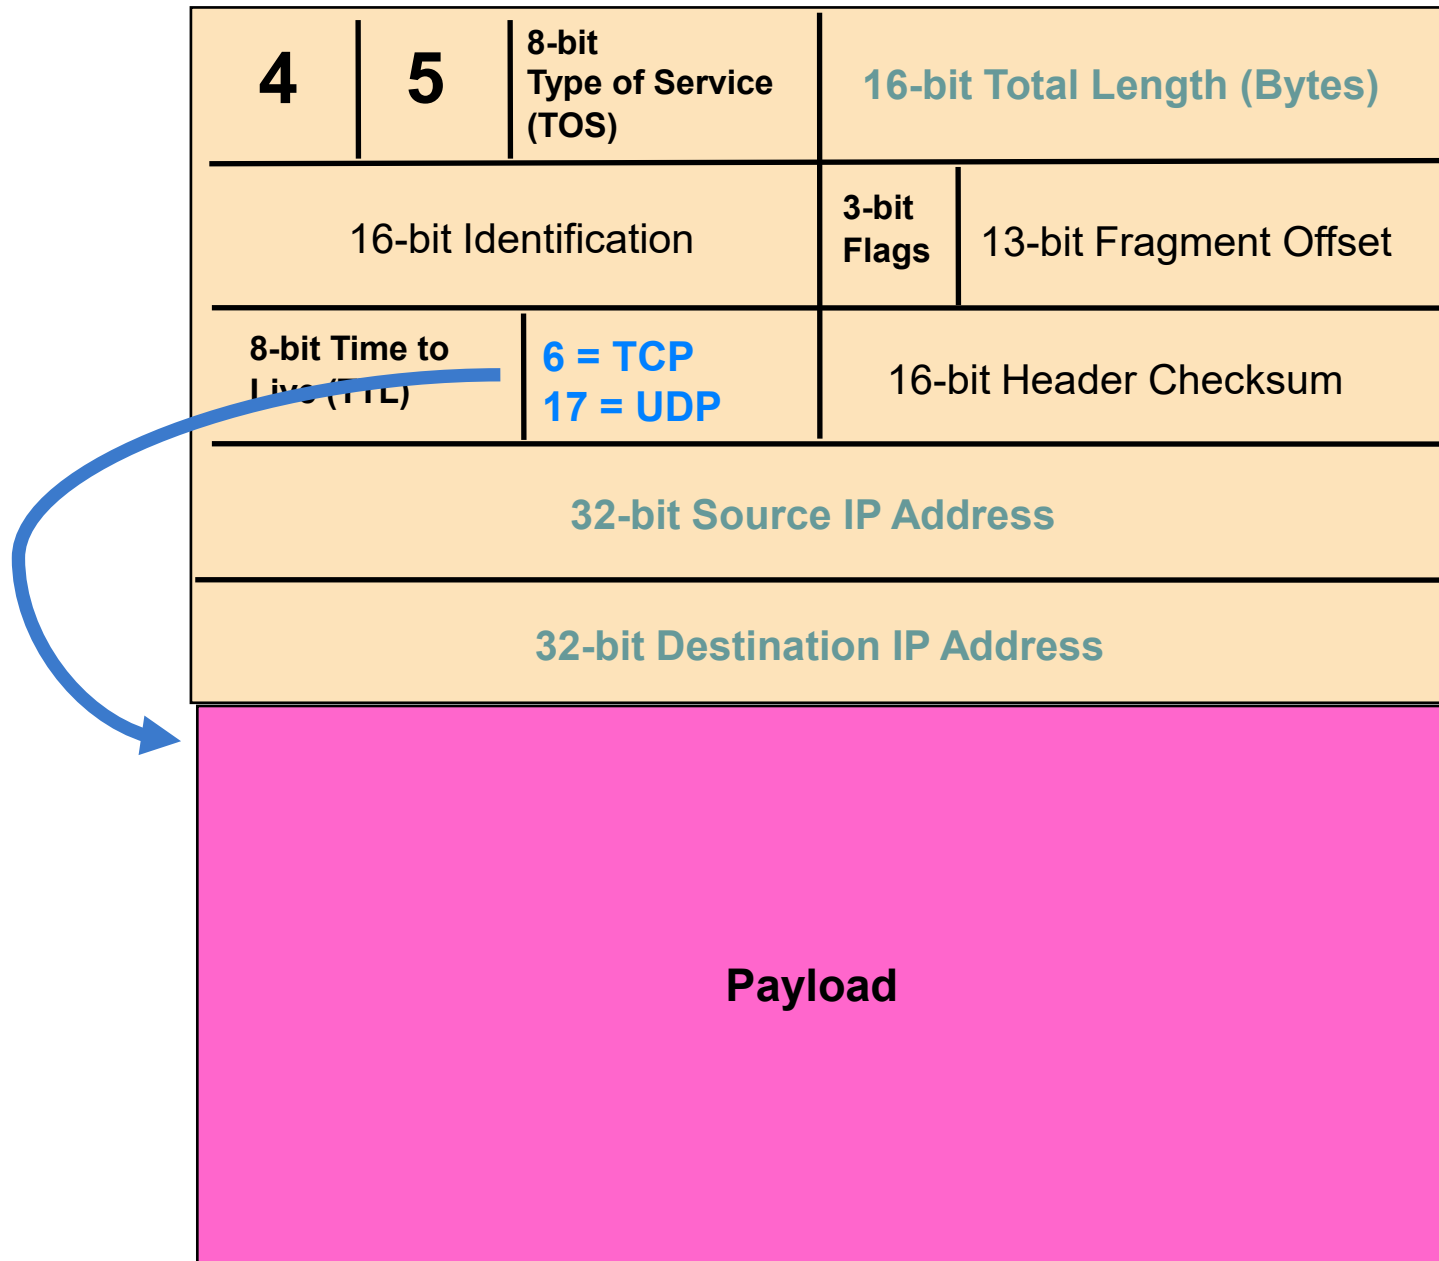
Most of states relate to

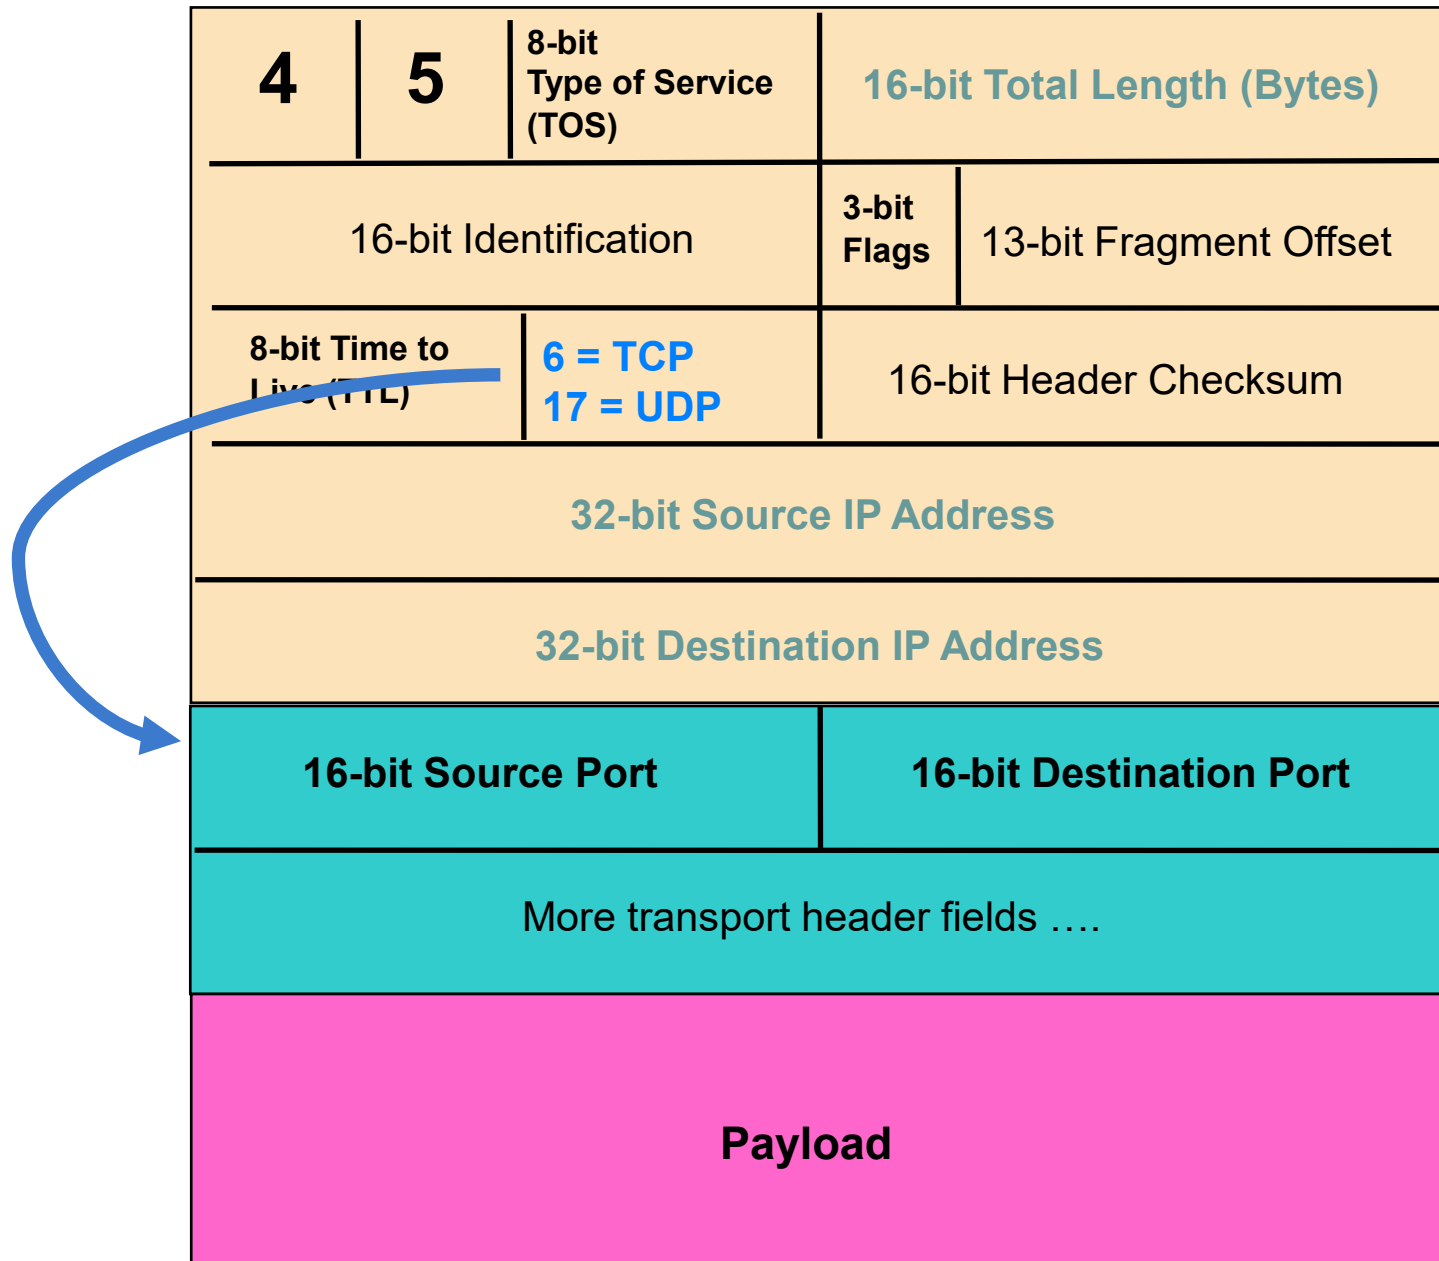
- the connection establishment (three-way handshake)
- the connection termination (ensuring reliability)

Data is exchanged in the **ESTABLISHED** state







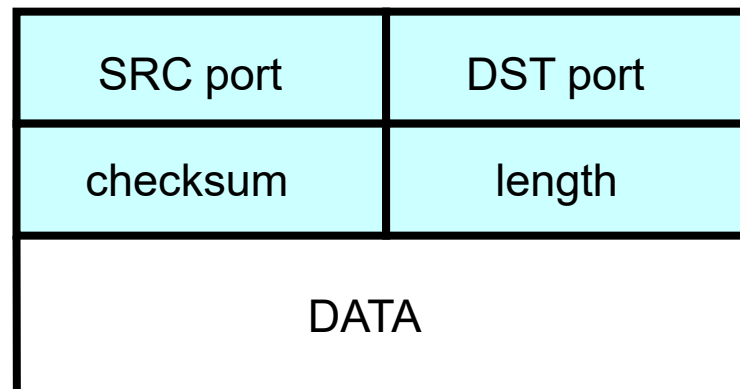




# UDP

# UDP: User Datagram Protocol

- Lightweight communication between processes
  - Avoid overhead and delays of ordered, reliable delivery
  - Send messages to and receive them from a socket
- UDP described in RFC 768 – (1980!)
  - IP plus port numbers to support (de)multiplexing
  - Optional error checking on the packet contents
    - (checksum field = 0 means “don’t verify checksum”)

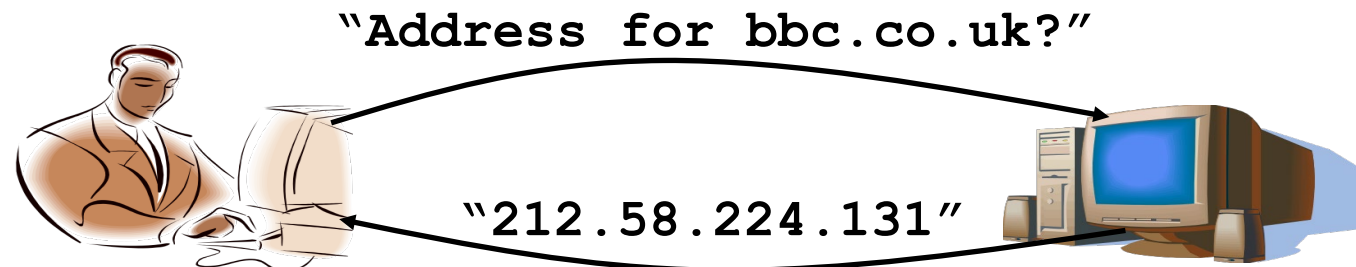


# Why Would Anyone Use UDP?

- Finer control over what data is sent and when
  - As soon as an application process writes into the socket
  - ... UDP will package the data and send the packet
- No delay for connection establishment
  - UDP just blasts away without any formal preliminaries
  - ... which avoids introducing any unnecessary delays
- No connection state
  - No allocation of buffers, sequence #s, timers ...
  - ... making it easier to handle many active clients at once
- Small packet header overhead
  - UDP header is only 8 bytes

# Popular Applications That Use UDP

- Some **interactive streaming** apps
  - Retransmitting lost/corrupted packets often pointless: by the time the packet is retransmitted, it's too late
  - telephone calls, video conferencing, gaming...
  - ***Modern streaming protocols using TCP (and HTTP)***
- Simple query protocols like Domain Name System (DNS)
  - Connection establishment overhead would double cost
  - Easier to have **application** retransmit if needed



# TCP

# Transmission Control Protocol (TCP)

- Reliable, in-order delivery *(previously, but quick review)*
  - Ensures byte stream (eventually) arrives intact
    - In the presence of **corruption** and **loss**
- Connection oriented *(today)*
  - Explicit set-up and tear-down of TCP session
- Full duplex stream-of-bytes service *(today)*
  - Sends and receives a stream of bytes, not messages
- Flow control *(previously, but quick review)*
  - Ensures that sender doesn't overwhelm receiver
- Congestion control *(next week)*
  - Dynamic adaptation to network path's capacity

# Basic Components of Reliability

- ACKs
  - Can't be reliable without knowing whether data has arrived
  - ***TCP uses byte sequence numbers to identify payloads***
- Checksums
  - Can't be reliable without knowing whether data is corrupted
  - ***TCP does checksum over TCP and pseudoheader***
- Timeouts and retransmissions
  - Can't be reliable without retransmitting lost/corrupted data
  - ***TCP retransmits based on timeouts and duplicate ACKs***
  - *Timeout based on estimate of RTT*

# Other TCP Design Decisions

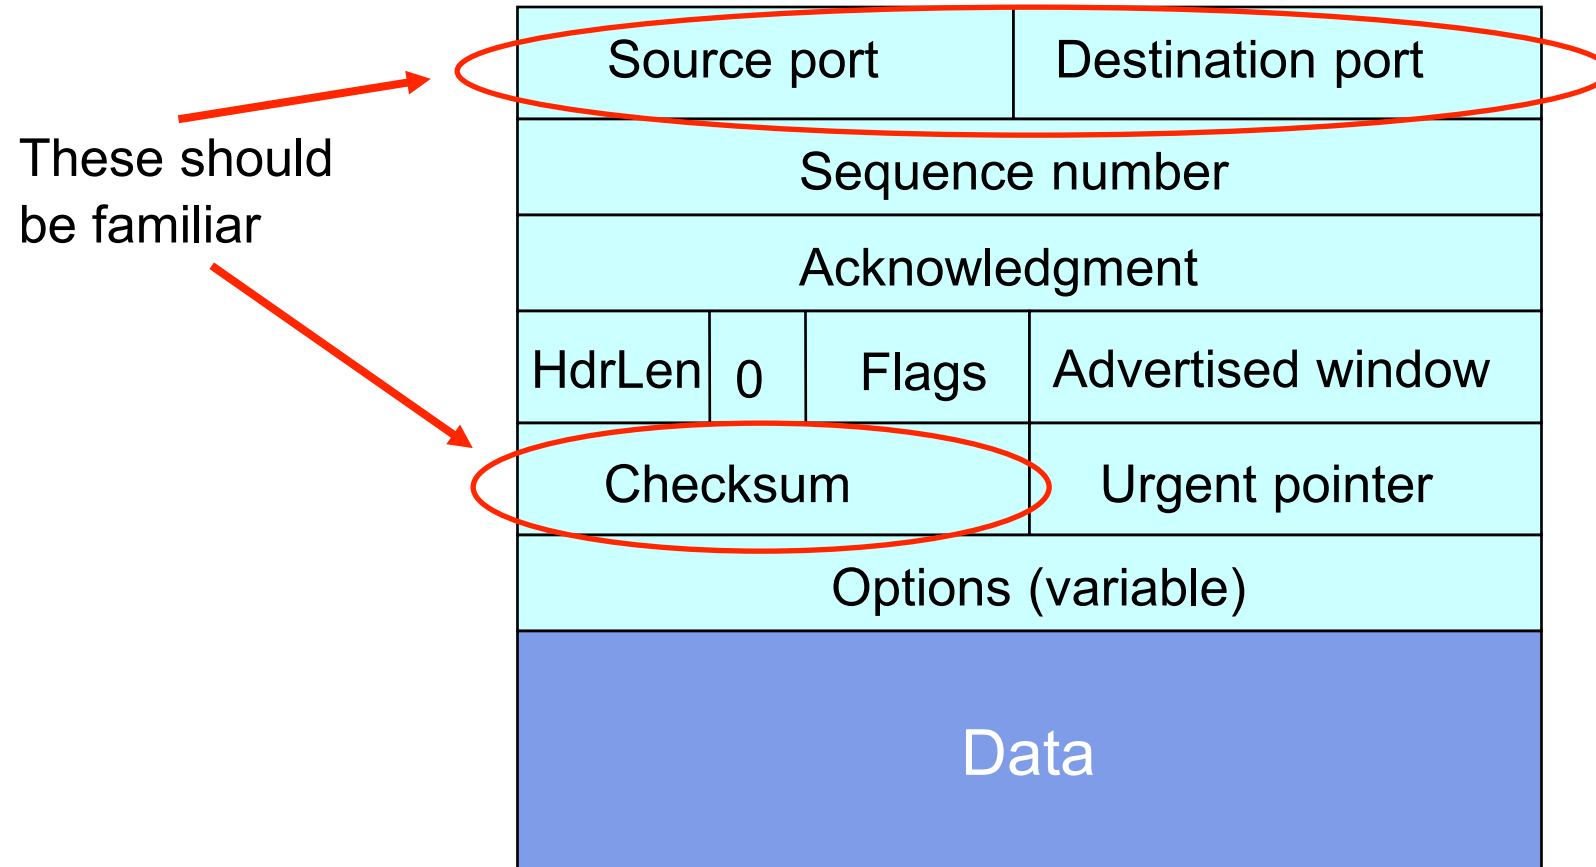
- Sliding window flow control
  - Allow  $W$  contiguous bytes to be in flight
- Cumulative acknowledgements
  - Selective ACKs (full information) also supported (ignore)
- Single timer set after each payload is ACKed
  - Timer is effectively for the “next expected payload”
  - When timer goes off, resend that payload and wait
  - And double timeout period
- Various tricks related to “fast retransmit”
  - Using duplicate ACKs to trigger retransmission



# TCP Header

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			


# TCP Header



# Segments and Sequence Numbers

# TCP “Stream of Bytes” Service...


Application @ Host A



Application @ Host B

# ... Provided Using TCP “Segments”

Host A




Host B

# TCP Segment

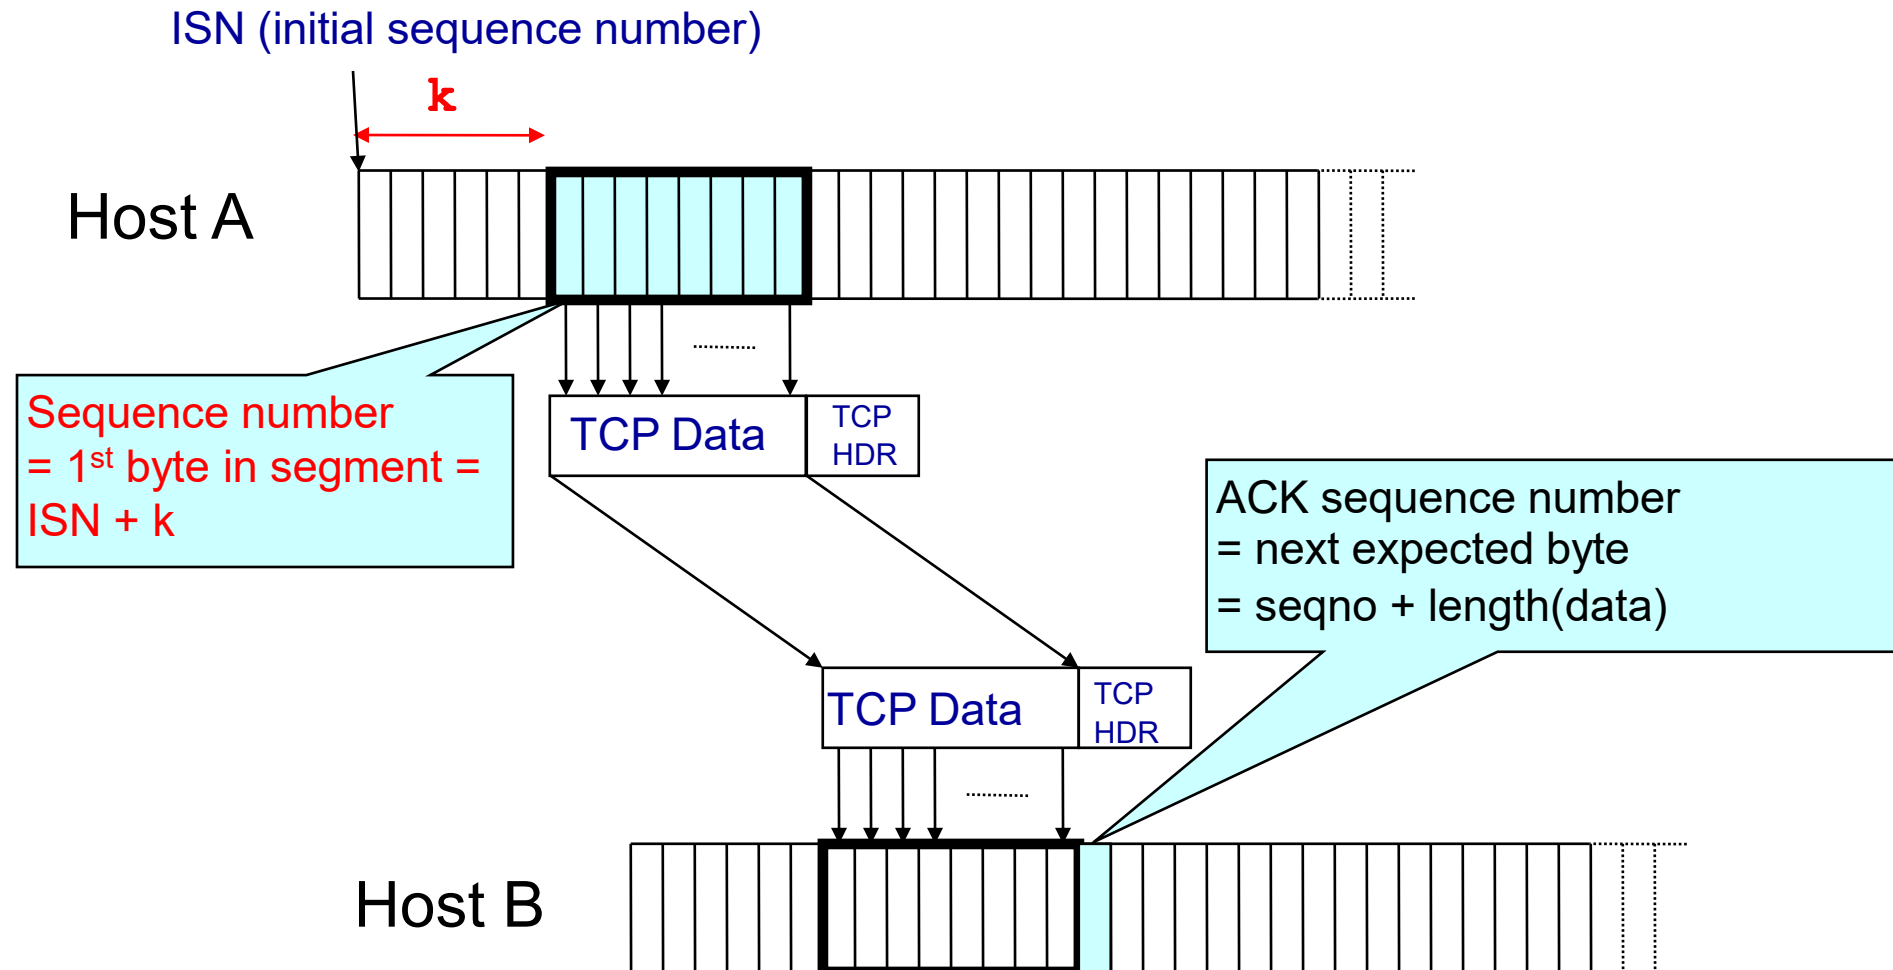


- **IP packet**
  - No bigger than Maximum Transmission Unit (**MTU**)
  - E.g., up to 1500 bytes with Ethernet
- **TCP packet**
  - IP packet with a TCP header and data inside
  - TCP header  $\geq 20$  bytes long
- **TCP segment**
  - No more than **Maximum Segment Size** (MSS) bytes
  - E.g., up to 1460 consecutive bytes from the stream
  - $MSS = MTU - (IP \text{ header}) - (TCP \text{ header})$

# Sequence Numbers



# Sequence Numbers





# ACKing and Sequence Numbers

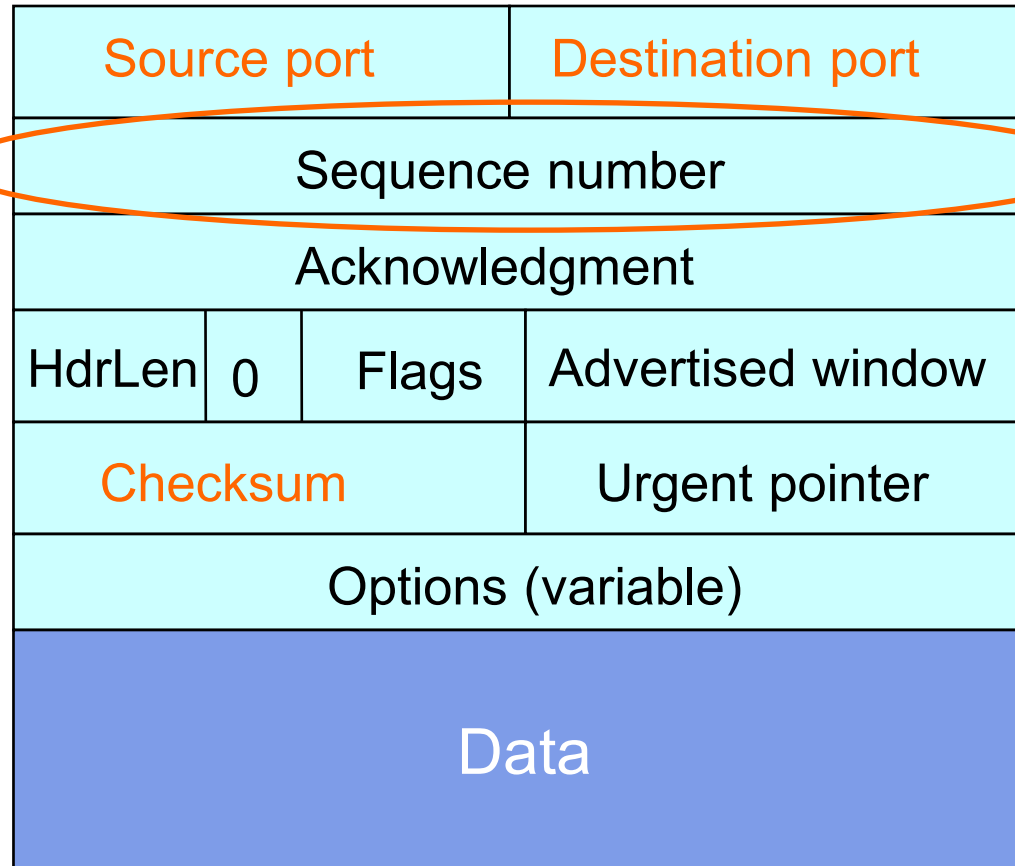
- Sender sends packet
  - Data starts with sequence number  $X$
  - Packet contains  $B$  bytes
    - $X, X+1, X+2, \dots, X+B-1$
- Upon receipt of packet, receiver sends an ACK
  - If all data prior to  $X$  already received:
    - ACK acknowledges  $X+B$  (because that is next expected byte)
  - If highest contiguous byte received is smaller value  $Y$ 
    - ACK acknowledges  $Y+1$
    - Even if this has been ACKed before

# Normal Pattern

- Sender:  $\text{seqno}=X$ ,  $\text{length}=B$
- Receiver:  $\text{ACK}=X+B$
- Sender:  $\text{seqno}=X+B$ ,  $\text{length}=B$
- Receiver:  $\text{ACK}=X+2B$
- Sender:  $\text{seqno}=X+2B$ ,  $\text{length}=B$
- ...
- Seqno of next packet is same as last ACK field

# TCP Header

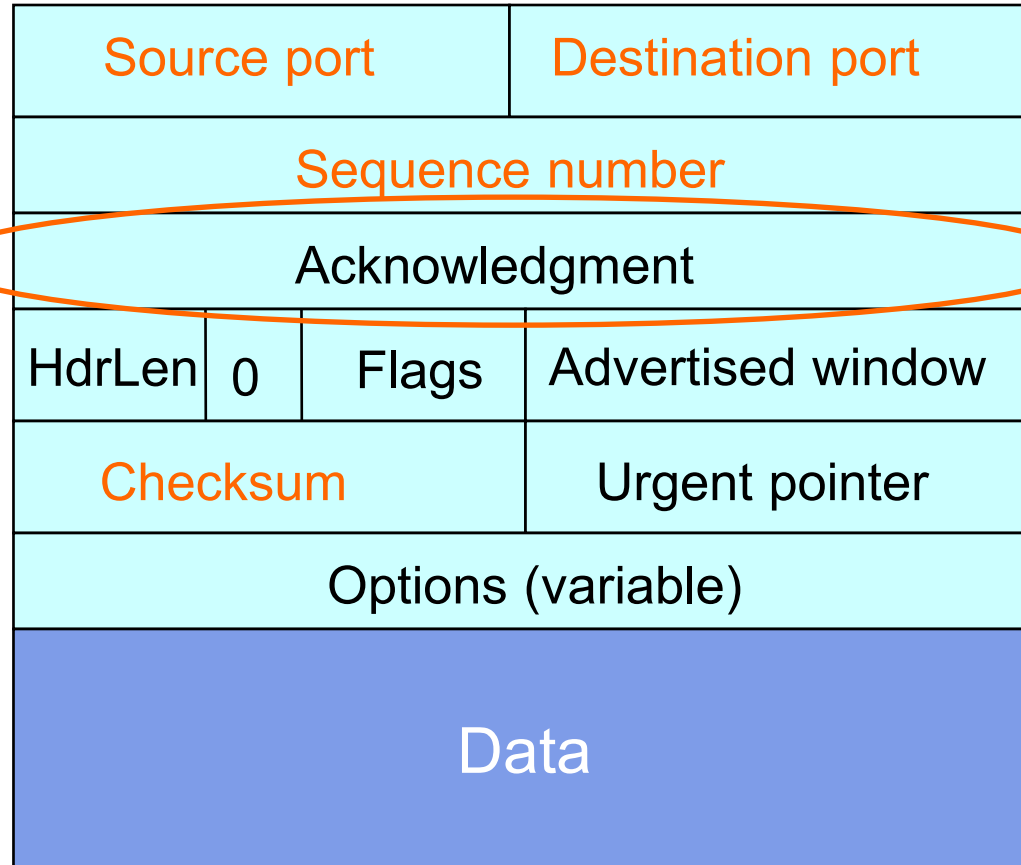
Starting **byte offset of data** carried in this segment



# TCP Header

Acknowledgment  
gives seqno just  
beyond highest  
seqno received  
**in order**

*“What Byte is Next”*



# TCP Header

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			

# Sliding Window Flow Control

- Advertised Window:  $W$ 
  - Can send  $W$  bytes beyond the next expected byte
- Receiver uses  $W$  to prevent sender from overflowing buffer
- Limits number of bytes sender can have in flight

# Implementing Sliding Window

- Both sender & receiver maintain a **window**
  - Sender: not yet ACK'ed
  - Receiver: not yet delivered to application
- **Left edge** of window:
  - Sender: beginning of **unacknowledged** data
  - Receiver: beginning of **undelivered** data
- For the sender:
  - Window size = maximum amount of data in flight
- For the receiver:
  - Window size = maximum amount of undelivered data

# *Sliding Window Summary*

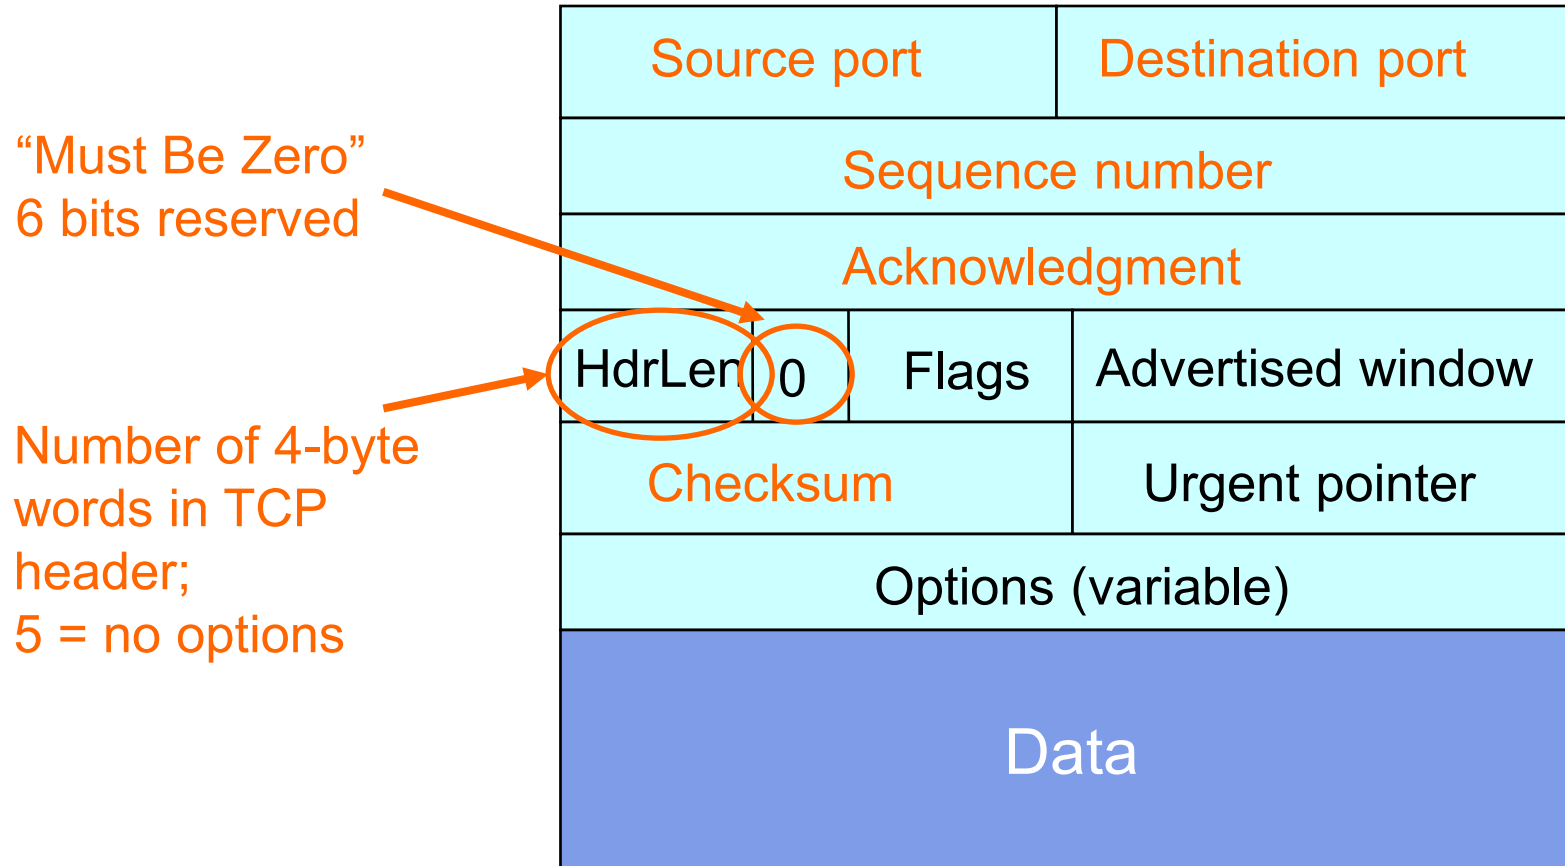
- Sender: window **advances** when new data ack'd
- Receiver: window advances as receiving process **consumes** data
- Receiver **advertises** to the sender where the receiver window currently ends (“righthand edge”)
  - Sender agrees not to exceed this amount
  - It makes sure by setting its own window size to a value that can't send beyond the receiver's righthand edge



# Advertised Window Limits Rate

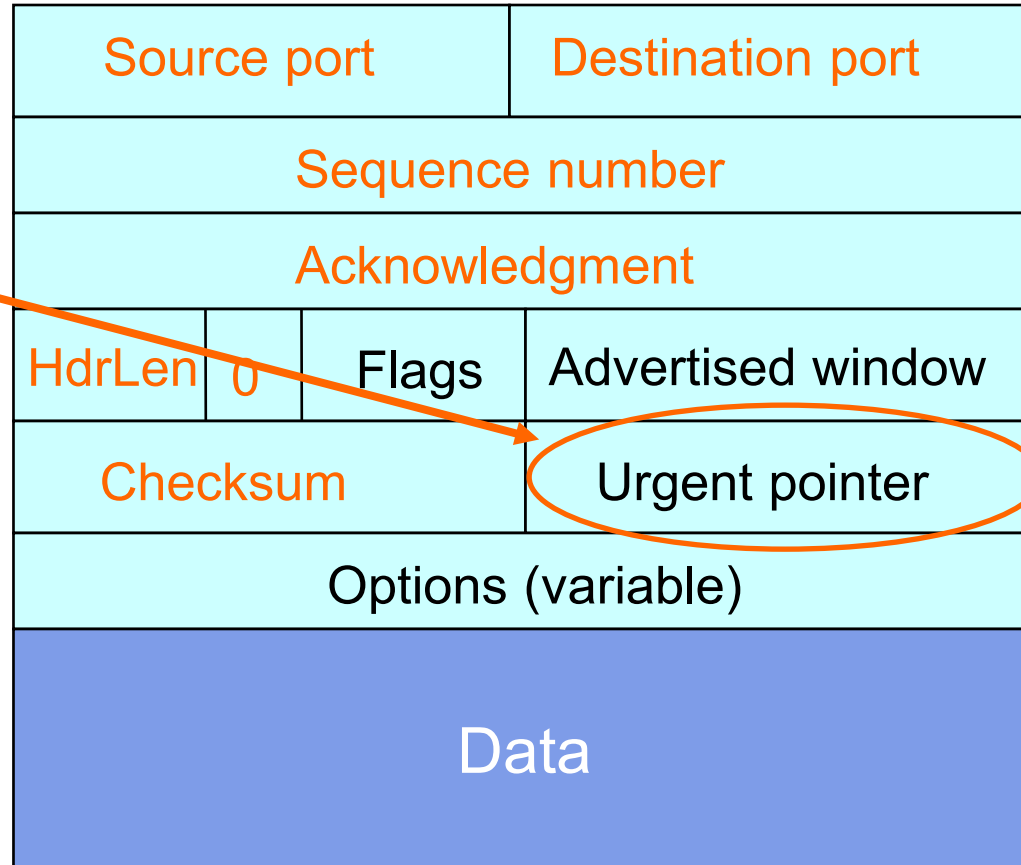
- Sender can send no faster than  $W/RTT$  bytes/sec
- Receiver only advertises more space when it has consumed old arriving data
- In original TCP design, that was the **sole** protocol mechanism controlling sender's rate
- What's missing?

# TCP Header: What's left?



# TCP Header: What's left?

Used with **URG**  
flag to indicate  
urgent data (not  
discussed further)



# TCP Header: What's left?


Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			

# TCP Connection Establishment and Initial Sequence Numbers

# Initial Sequence Number (ISN)

- Sequence number for the very first byte
  - E.g., Why not just use ISN = 0?
- Practical issue
  - IP addresses and port #s uniquely identify a connection
  - Eventually, though, these port #s do get **used again**
  - ... small chance an old packet is **still in flight**
- TCP therefore **requires** changing ISN
  - initially set from 32-bit clock that ticks every 4 microseconds
  - now drawn from a pseudo random number generator (security)
- To establish a connection, hosts exchange ISNs
  - **How does this help?**

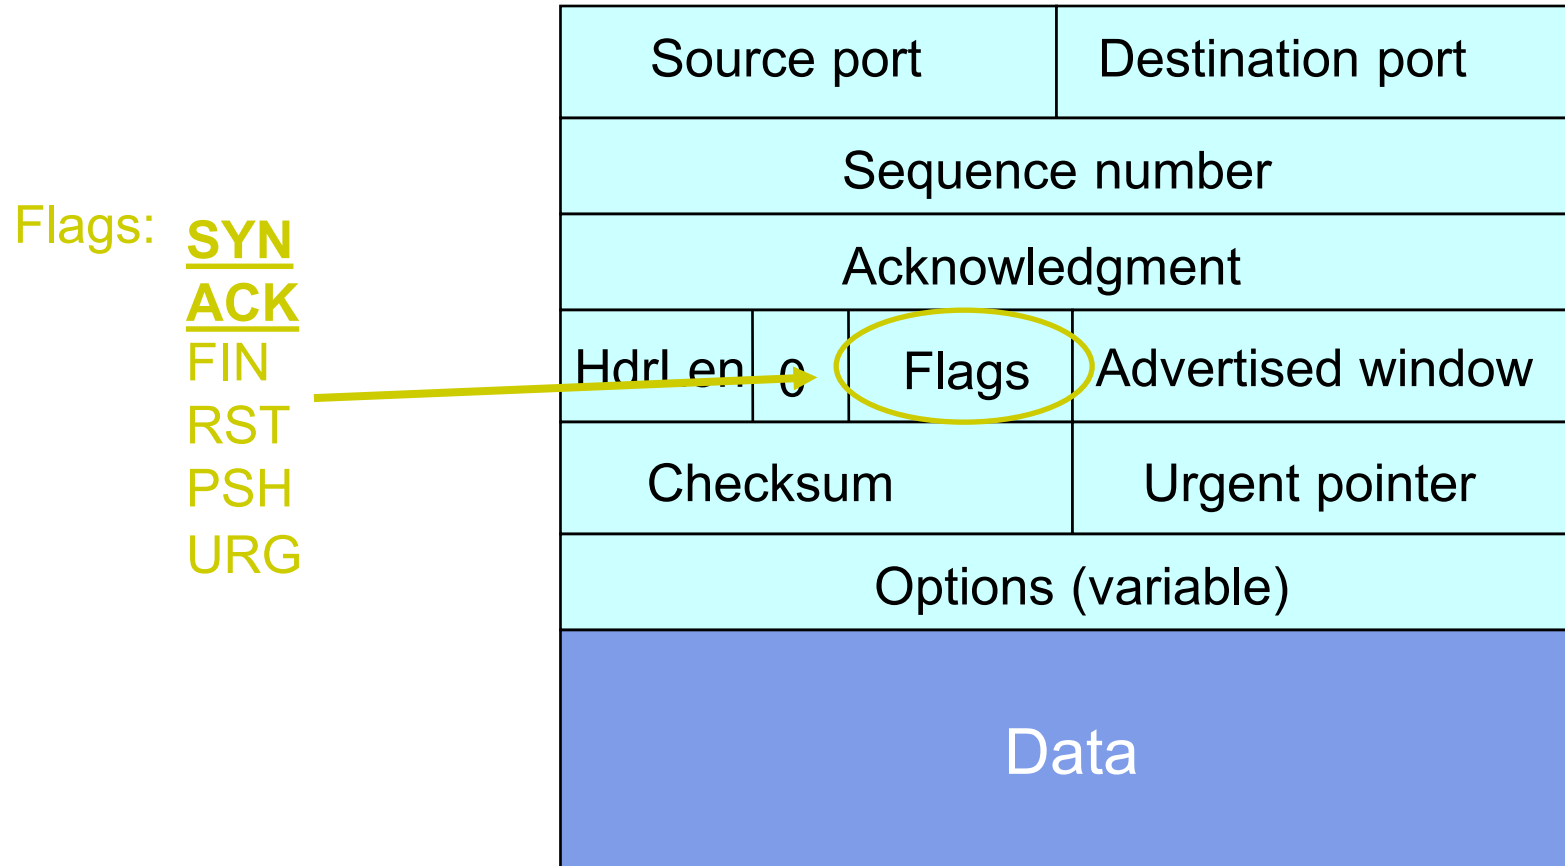
# Establishing a TCP Connection



Each host tells its ISN to the other host.

- Three-way handshake to establish connection
  - Host A sends a **SYN** (open; “synchronize sequence numbers”)
  - Host B returns a SYN acknowledgment (**SYN ACK**)
  - Host A sends an **ACK** to acknowledge the SYN ACK

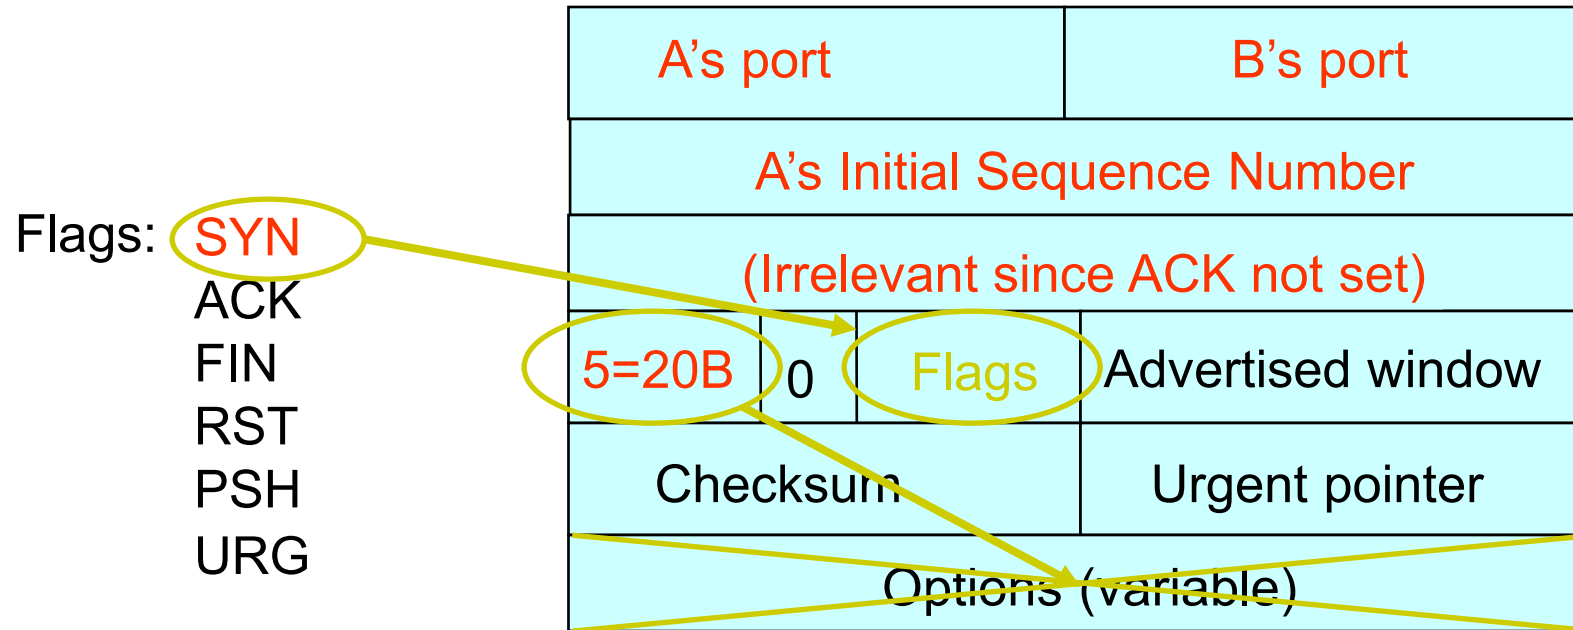
# TCP Header



See `/usr/include/netinet/tcp.h` on Unix Systems

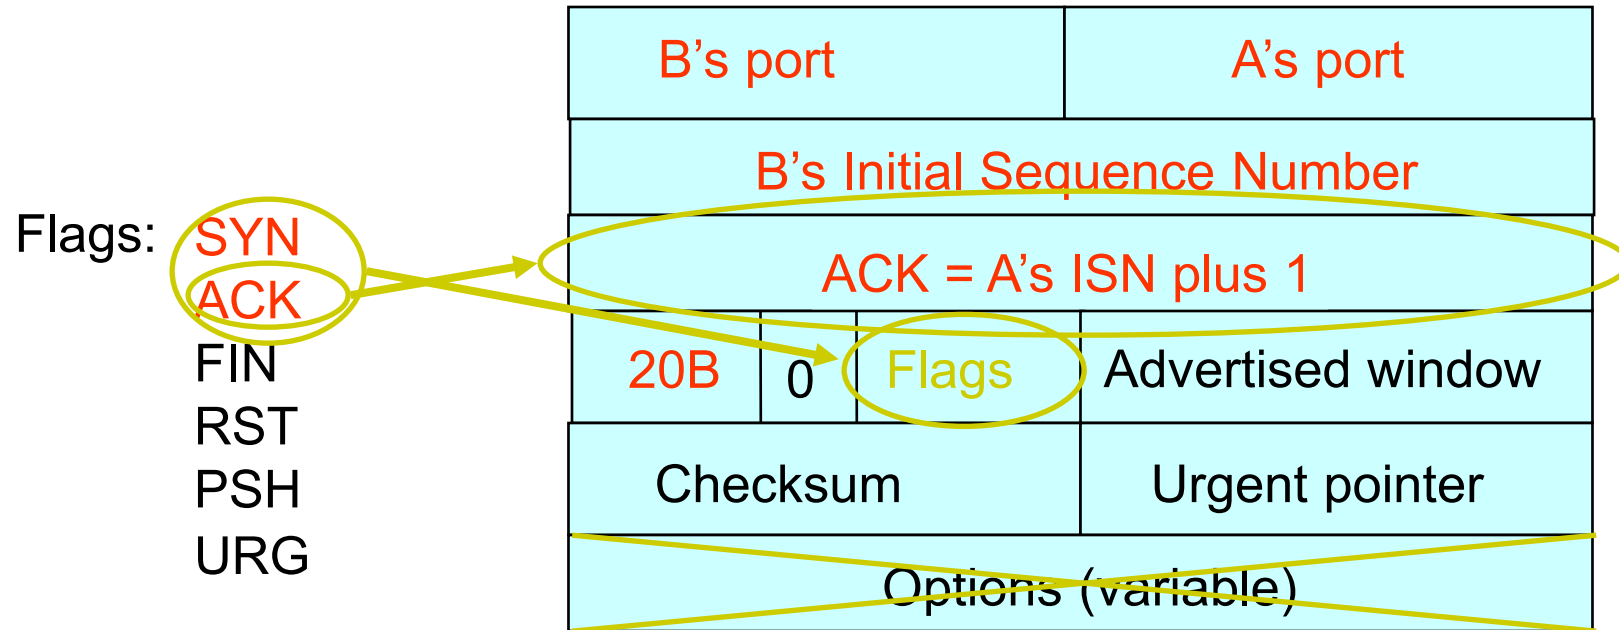


# Step 1: A's Initial SYN Packet



**A tells B it wants to open a connection...**

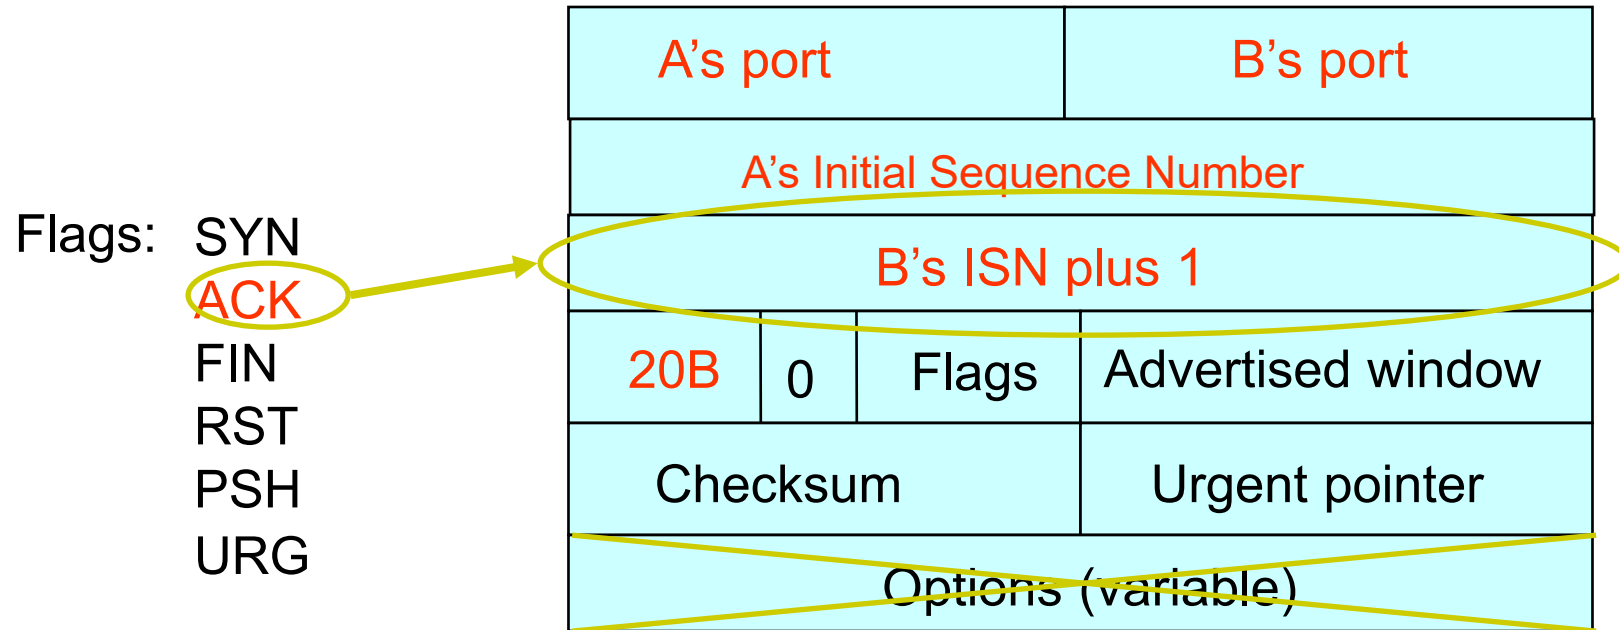
## Step 2: B's SYN-ACK Packet



**B tells A it accepts, and is ready to hear the next byte...**

**... upon receiving this packet, A can start sending data**


# Step 3: A's ACK of the SYN-ACK



**A tells B it's likewise okay to start sending**

**... upon receiving this packet, B can start sending data**

# Timing Diagram: 3-Way Handshaking



# What if the SYN Packet Gets Lost?


- Suppose the SYN packet gets lost
  - Packet is lost inside the network, or:
  - Server **discards** the packet (e.g., listen queue is full)
- Eventually, no SYN-ACK arrives
  - Sender sets a **timer** and **waits** for the SYN-ACK
  - ... and retransmits the SYN if needed
- How should the TCP sender set the timer?
  - Sender has **no idea** how far away the receiver is
  - Hard to guess a reasonable length of time to wait
  - **SHOULD** (RFCs 1122 & 2988) use default of **3 seconds**
    - Other implementations instead use 6 seconds

# SYN Loss and Web Downloads

- User clicks on a hypertext link
  - Browser creates a socket and does a “connect”
  - The “connect” triggers the OS to transmit a SYN
- If the SYN is lost...
  - 3-6 seconds of delay: can be **very long**
  - User may become impatient
  - ... and click the hyperlink again, or click “reload”
- User triggers an “abort” of the “connect”
  - Browser creates a **new** socket and another “connect”
  - Essentially, forces a faster send of a new SYN packet!
  - Sometimes very effective, and the page comes quickly


# Tearing Down the Connection

# Normal Termination, One Side At A Time






# Normal Termination, Both Together



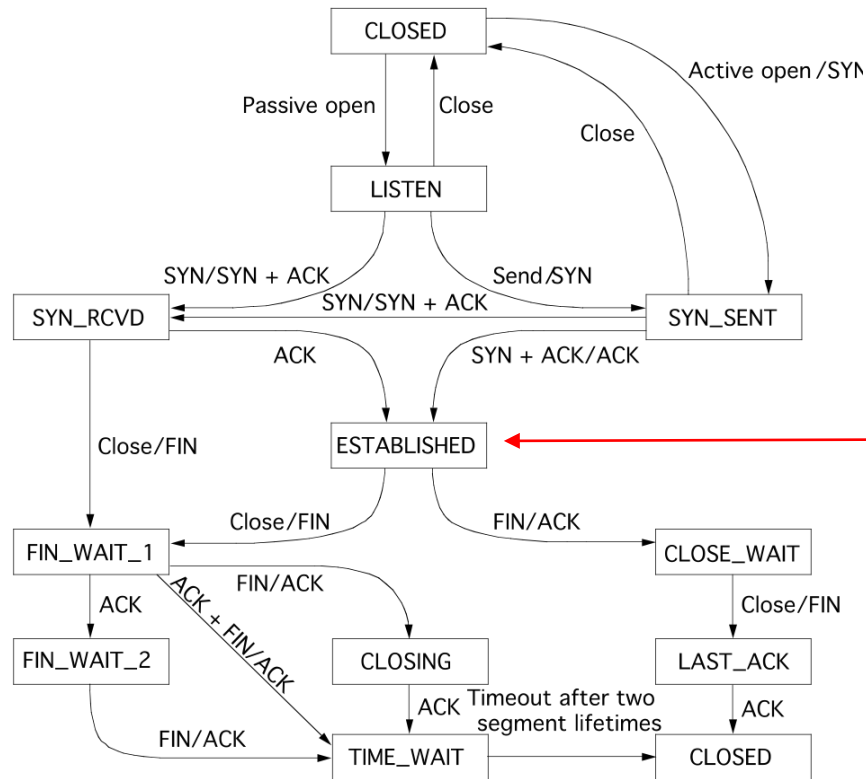
- Same as before, but B sets **FIN** with their ack of A's **FIN**

# Abrupt Termination



- A sends a RESET (**RST**) to B
  - E.g., because app. process on A **crashed**
- That's it
  - B does **not** ack the **RST**
  - Thus, **RST** is **not** delivered **reliably**
  - And: any data in flight is **lost**
  - But: if B sends anything more, will elicit **another RST**

# TCP State Transitions





Data, ACK exchanges are in here

# Reliability: TCP Retransmission

# Timeouts and Retransmissions

- Reliability requires retransmitting lost data
- Involves setting timer and retransmitting on timeout
- TCP resets timer whenever new data is ACKed
  - Retx of packet containing “next byte” when timer goes off

# Setting the Timeout Value




# RTT Estimation

- Use exponential averaging of RTT samples

$$\text{SampleRTT} = \text{AckRcvdTime} - \text{SendPacketTime}$$

$$\text{EstimatedRTT} = \alpha \times \text{EstimatedRTT} + (1 - \alpha) \times \text{SampleRTT}$$


$$0 < \alpha \leq 1$$



# Exponential Averaging Example

$$\text{EstimatedRTT} = \alpha * \text{EstimatedRTT} + (1 - \alpha) * \text{SampleRTT}$$



Assume RTT is constant  $\rightarrow$   $\text{SampleRTT} = \text{RTT}$





# Problem: Ambiguous Measurements

- How do we differentiate between the real ACK, and ACK of the retransmitted packet?



# Karn/Partridge Algorithm

- Measure *SampleRTT* only for original transmissions
  - Once a segment has been retransmitted, do not use it for any further measurements
  - Computes *EstimatedRTT* using  $\alpha = 0.875$
- Timeout value (RTO) =  $2 \times \text{EstimatedRTT}$
- Use exponential backoff for repeated retransmissions
  - Every time RTO timer expires, set  $\text{RTO} \leftarrow 2 \cdot \text{RTO}$ 
    - (Up to maximum  $\geq 60$  sec)
  - Every time new measurement comes in (= successful original transmission), collapse RTO back to  $2 \times \text{EstimatedRTT}$

# This is all very interesting, but.....

- Implementations often use a coarse-grained timer
  - 500 msec is typical
- So what?
  - Above algorithms are largely irrelevant
  - **Incurring a timeout is expensive**
- So we rely on duplicate ACKs

# Loss with cumulative ACKs

- Sender sends packets with 100B and seqnos.:
  - 100, 200, 300, 400, 500, 600, 700, 800, 900, ...
- Assume the fifth packet (seqno 500) is lost, but no others
- Stream of ACKs will be:
  - 200, 300, 400, 500, 500, 500, 500,...

# Loss with cumulative ACKs

- “Duplicate ACKs” are a sign of an *isolated* loss
  - The lack of ACK progress means 500 hasn't been delivered
  - Stream of ACKs means some packets are being delivered
- Therefore, could trigger resend upon receiving k duplicate ACKs
  - TCP uses  $k=3$
- We will revisit this in congestion control

# Reading: Book Kurose & Ross


Class textbook:

*Computer Networking: A Top-Down Approach (8<sup>th</sup> ed.)*

J.F. Kurose, K.W. Ross

Pearson, 2020

[http://gaia.cs.umass.edu/kurose\\_ross](http://gaia.cs.umass.edu/kurose_ross)



- Week 09
  - 3.5 (Connection-Oriented Transport: TCP)
  - 3.6 (Principles of Congestion Control) and 3.7 (TCP Congestion Control)

# Check Your Knowledge

PROBLEM SOLVING HOME

TRY A RANDOM PROBLEM

## INTERACTIVE END-OF-CHAPTER EXERCISES

Supplement to Computer Networking: A Top Down Approach 8th Edition

*"Tell me and I forget. Show me and I remember. Involve me and I understand." Chinese proverb*



### CHAPTER 3: TRANSPORT LAYER

- Internet checksum (similar to Chapter 3, P3 and P4)
- Reliable data transfer: rdt22
- TCP sequence and ACK numbers, with segment loss (similar to Chapter 3, P27)
- TCP RTT and timeout (similar to Chapter 3, P31)
- TCP congestion window evolution (similar to Chapter 3, P40)
- TCP retransmissions (reliable data transmission with ACK loss)
- UDP Mux and Demux
- TCP Mux and Demux

[http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)