

# Computer Networks and Internet Technology

2021W703033 VO Rechnernetze und Internettechnik  
Winter Semester 2021/22

Jan Beutel

# Communication Networks and Internet Technology

## Part 1: General overview

### #1 What is a network made of?


How is it shared?

How is it organized?

How does communication happen?

How do we characterize it?

Networks are composed of three basic components



# Communication Networks and Internet Technology

## Part 1: General overview

What is a network made of?

#2

How is it shared?

How is it organized?

How does communication happen?

How do we characterize it?

So far, we've been discussing what  
the "last mile" of the Internet looks like

What about the rest of the network?

## 3 must-have requirements of a good network topology

### Tolerate failures

several paths between each source and destination

### Possess enough sharing to be feasible & cost-effective

number of links should not be too high

### Provide adequate per-node capacity

number of links should not be too small




Compare these three designs in terms of  
**sharing**, **resiliency**, and **per-node capacity**

design

full-mesh

chain

bus



advantages


disadvantages



Switched networks provide  
**reasonable** and **flexible** compromise

design

switched




advantages

sharing and per-node capacity can be adapted  
to fit the network needs

disadvantages

require smart devices to perform:  
forwarding, routing, **resource allocation**

Links and switches are shared between flows




Communication Networks and Internet Technology

**This weeks lecture**

There exist two approaches to sharing:  
**reservation** and **on-demand**



Reservation



On-demand

principle

reserve the bandwidth  
you need in advance

send data when you need

Both are examples of **statistical multiplexing**

Reservation


On-demand

multiplexing

at the flow-level

at the packet-level

Between reservation and on-demand:  
Which one do you pick?




Consider that each source  
needs 10 Mbps


What do they get with:

- reservation
- on-demand

Assume the following peak demand and flow duration



Assume the following peak demand and flow duration



What does each source get with **reservation** and on-demand?

- first-come first-served
- equal (10 Mbps)



## Peak vs average rates

Each flow has	Peak rate	$P$
	Average rate	$A$

Reservation must reserve  $P$ , but level of utilization is  $A/P$

$P=100$  Mbps,  $A=10$  Mbps, level of utilization=10%

On-demand can usually achieve higher level of utilization  
depends on degree of sharing and burstiness of flows

Ultimately, it depends on the application

Reservation **makes sense** when **P/A is small**

voice traffic has a ratio of 3 or so

Reservation **wastes capacity** when **P/A is big**

data applications are bursty, ratios  $>100$  are common

Reservation **makes sense** when  $P/A$  is small

voice traffic has a ratio of 3 or so

Reservation **wastes capacity** when  $P/A$  is big

data applications are bursty, ratios  $>100$  are common

**That's why the phone network used reservations  
... and why the Internet does not!**

The two approaches are implemented using circuit-switching or packet-switching, respectively



Reservation

On-demand

implementation

circuit-switching

packet-switching

Reservation


On-demand

implementation

circuit-switching

packet-switching

## Circuit switching relies on the Resource Reservation Protocol




- (1) **src** sends a reservation request for 10Mbps to **dst**
- (2) switches “establish a circuit”
- (3) **src** starts sending data
- (4) **src** sends a “teardown circuit” message

# There exist many kinds of circuits


## Time-based multiplexing

- divide time in slots
- allocate one slot per circuit




## Frequency-based multiplexing

- divide spectrum in frequency bands
- allocate one band per circuit



Let's walk through example of data transfer  
using circuit switching







The efficiency of the transfer depends on  
how utilized the circuit is once established

This is an example of poor efficiency


The circuit is mostly idle due to traffic bursts



This is another example of poor efficiency  
The circuit is used for a short amount of time



Another problem of circuit switching is that it doesn't route around trouble



A is forced to signal a new circuit to restore communication

# Pros and cons of circuit switching

## advantages

predictable performance

simple & fast switching  
once circuit established

## disadvantages

inefficient if traffic is bursty or short

complex circuit setup/teardown  
which adds delays to transfer

requires new circuit upon failure

What about packet switching?


Reservation

circuit-switching

On-demand

packet-switching

In packet switching,  
data transfer is done using independent packets




Each packet contains a destination (**dst**)

Since packets are sent without global coordination,  
they can “clash” with each other




To absorb transient overload,  
packet switching relies on buffers

To absorb transient overload,  
packet switching relies on buffers



## Packet switching routes around trouble



route is recomputed on the fly by s2

# Pros and cons of packet switching

## advantages

efficient use of resources

simpler to implement

route around trouble

## disadvantages

unpredictable performance

requires buffer management and  
congestion control

Packet switching beats circuit switching  
with respect to *resiliency* and *efficiency*

Internet  packets

Packet switching will be our focus for the rest of the course

# Communication Networks and Internet Technology

## Part 1: Overview

What is a network made of?


How is it shared?

#3 **How is it organized?**


How does communication happen?

How do we characterize it?

The **Internet** is a **network of networks**




## Internet Service Providers








So far our view of the Internet



The Internet is a tad more complex...









# The Internet has a hierarchical structure

Tier-1 international	have no provider
Tier-2 national	provide transit to Tier-3s have at least one provider
Tier-3 local	do not provide any transit have at least one provider



# The distribution of networks in Tiers is extremely skewed towards Tier-3s


		total	~60,000 networks
Tier-1 international	have no provider		~12
Tier-2 national	provide transit to Tier-3s have at least one provider		~1,000s
Tier-3 local	do not provide any transit have at least one provider		85-90%





Some networks have an incentive to connect directly,  
to reduce their bill with their own provider

**This is known as “peering”**



Interconnecting each network to its neighbors one-by-one is not cost effective

**Physical** costs

of provisioning or renting physical links


**Bandwidth** costs

a lot of links are not necessarily fully utilized

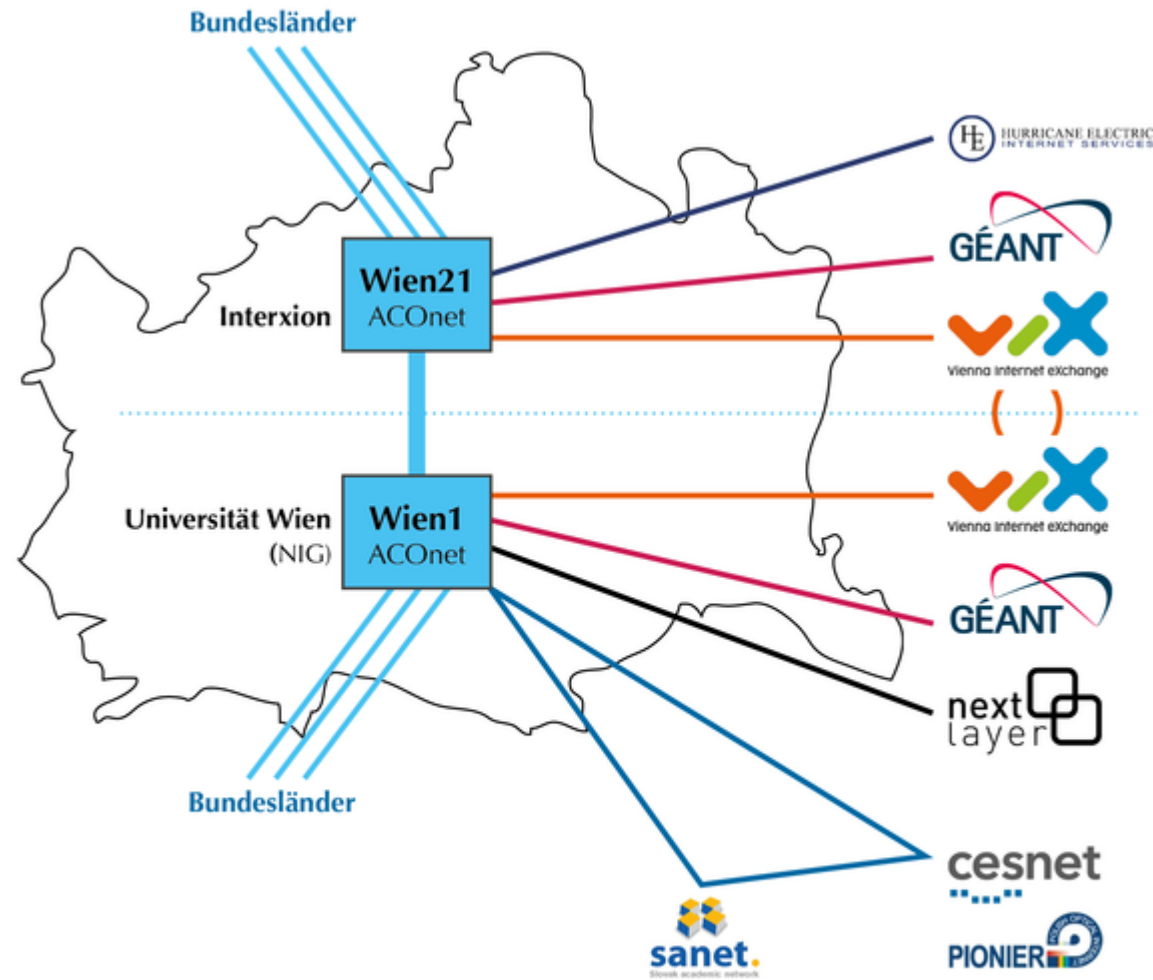
**Human** costs

to manage each connection individually

Internet eXchange Points (IXPs) solve these problems by letting *many* networks connect in one location



# Let's explore our network environment



# Communication Networks and Internet Technology

## Part 1: General overview

What is a network made of?

How is it shared?

How is it organized?

#4

How does communication happen?

How do we characterize it?


The Internet should allow


**processes on different hosts**  
to exchange data


everything else is just commentary...



How do you exchange data in a network as complex as  
**this?**








To exchange data, Alice and Bob use  
a set of network protocols

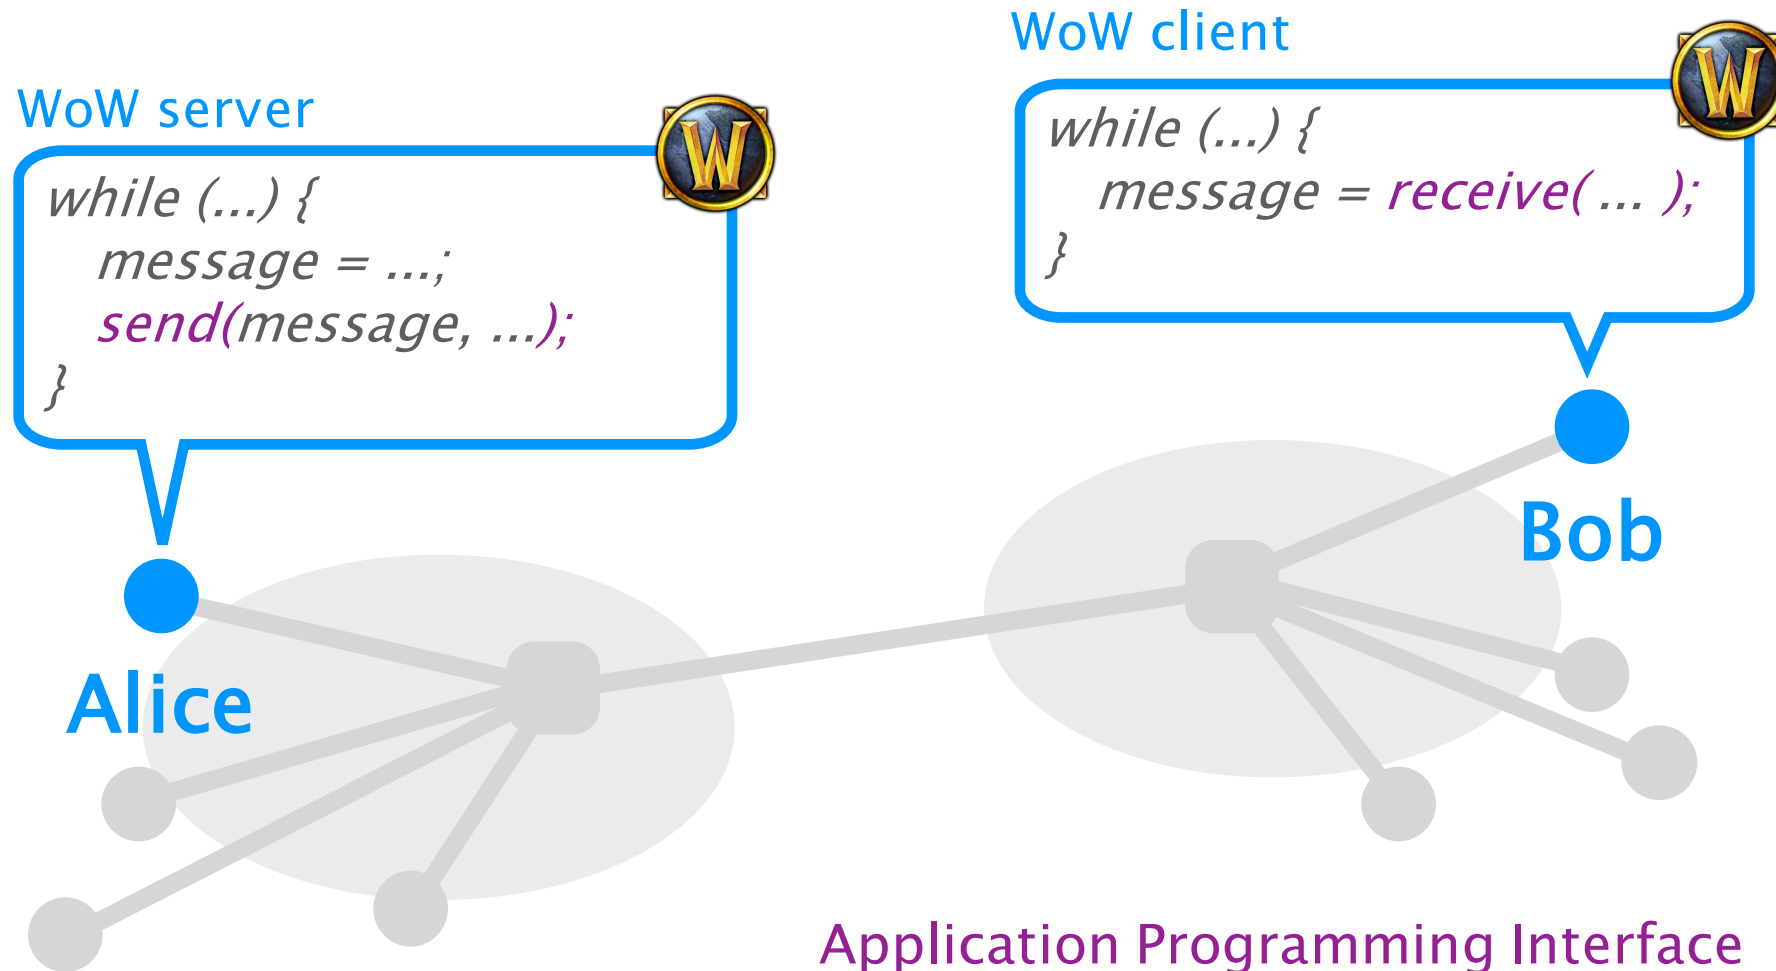
A protocol is like a conversational convention:  
who should talk next and how they should respond



Sometimes implementations are not compliant...



Each protocol is governed  
by a specific interface



In practice, there exists **a lot** of network protocols.  
How does the Internet organize **this**?






HOW STANDARDS PROLIFERATE:  
(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC.)

SITUATION:  
THERE ARE  
14 COMPETING  
STANDARDS.

14?! RIDICULOUS!  
WE NEED TO DEVELOP  
ONE UNIVERSAL STANDARD  
THAT COVERS EVERYONE'S  
USE CASES.



SOON:

SITUATION:  
THERE ARE  
15 COMPETING  
STANDARDS.

# Modularity is a key component of any good system

## Problem

can't build large systems out of spaghetti code  
hard (if not, impossible) to understand, debug, update

need to bound the scope of changes  
evolve the system without rewriting it from scratch

## Solution

**Modularity is how we do it**  
...and understand the system at a higher-level





Photo: Donna Coveney

Modularity,  
based on abstraction,  
is *the* way things get done

— *Barbara Liskov, MIT*

To provide structure to the design of network protocols,  
network designers organize protocols in layers



and the network hardware/software  
that implement them

Internet communication can be decomposed  
in 5 independent layers (or 7 layers for the OSI model)

layer

L5      Application

L4      Transport

L3      Network

L2      Link

L1      Physical

Each layer provides a service to the layer above

	layer	service provided:
L5	Application	network access
L4	Transport	end-to-end delivery (reliable or not)
L3	Network	global best-effort delivery
L2	Link	local best-effort delivery
L1	Physical	physical transfer of bits

Each layer provides a service to the layer above  
by using the services of the layer directly below it

Applications

...built on...

Reliable (or unreliable) transport

...built on...


Best-effort global packet delivery

...built on...

Best-effort local packet delivery

...built on...

Physical transfer of bits



Each layer (except for L3) is implemented with different protocols

	layer	protocol
L5	Application	HTTP, SMTP, FTP, SIP, ...
L4	Transport	TCP, UDP, SCTP
L3	Network	IP
L2	Link	Ethernet, Wifi, (A/V)DSL, WiMAX, LTE, ...
L1	Physical	Twisted pair, fiber, coaxial cable, ...



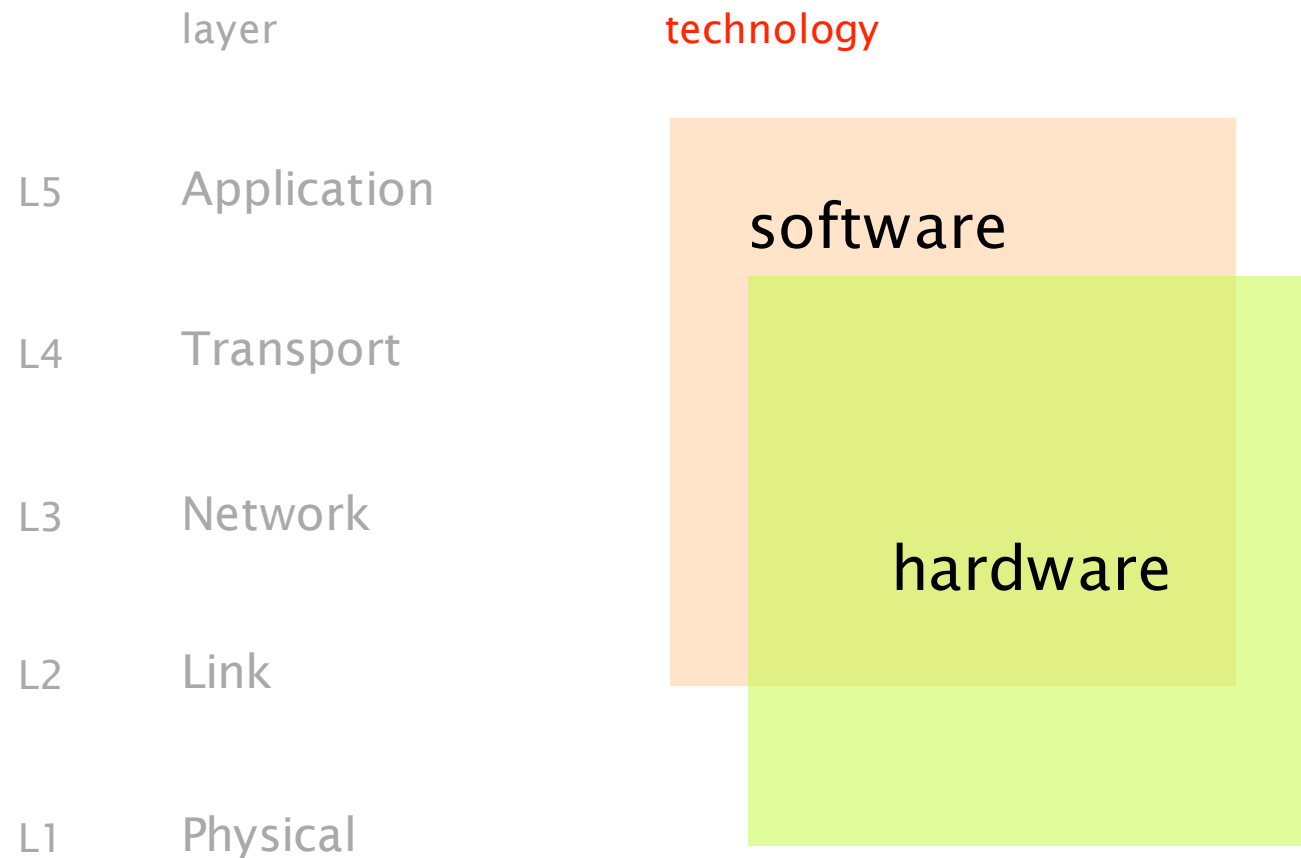
The Internet Protocol (IP) acts as  
an unifying, network, layer

	layer	protocol
L5	Application	HTTP, SMTP, FTP, SIP, ...
L4	Transport	TCP, UDP, SCTP
L3	Network	IP
L2	Link	Ethernet, Wifi, (A/V)DSL, Cable, LTE, ...
L1	Physical	Twisted pair, fiber, coaxial cable, ...

Each layer has a unit of **data**

	layer	role
L5	Application	exchanges <b>messages</b> between processes
L4	Transport	transports <b>segments</b> between end-systems
L3	Network	moves <b>packets</b> around the network
L2	Link	moves <b>frames</b> across a link
L1	Physical	moves <b>bits</b> across a physical medium

Each layer is implemented with different protocols  
and technologies



hardware

L1

L2

L3


L4

L5

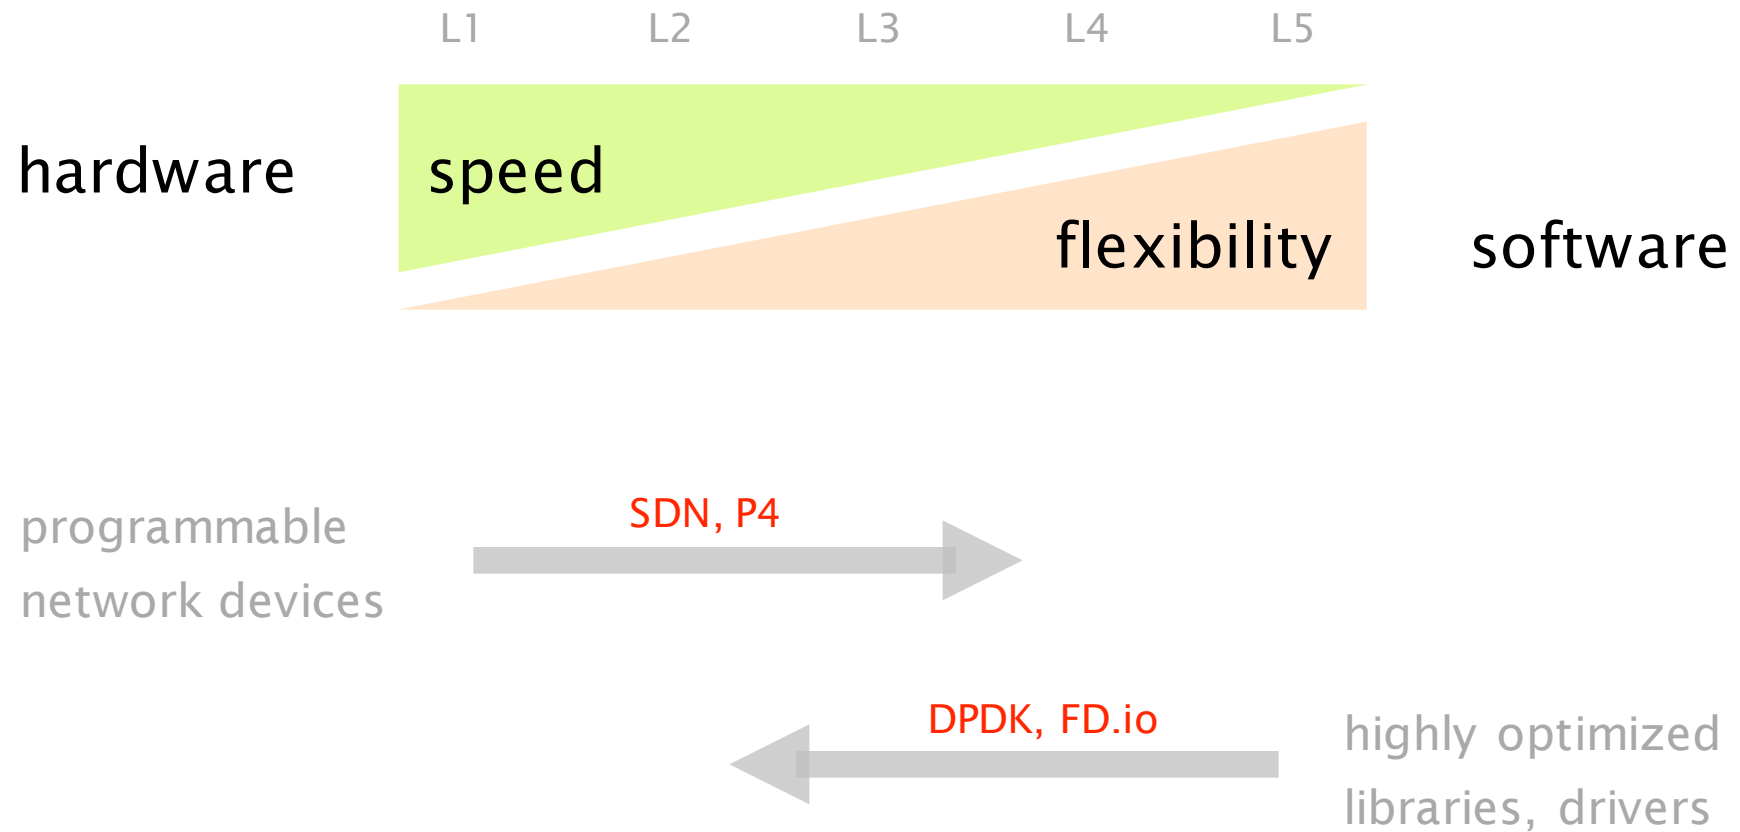
speed

flexibility

software



# Software and hardware advancements



Microsoft Supercharges Bing · X

Secure https://www.wired.com/2014/06/microsoft-fpga/

WIRE|D

Microsoft Supercharges Bing Search With Programmable Chips

SUBSCRIBE

BUSINESS

CULTURE

DESIGN

GEAR

SCIENCE

SECURITY

TRANSPORTATION

SHARE

f

SHARE  
1123

t

TWEET

COMMENT

103

EMAIL

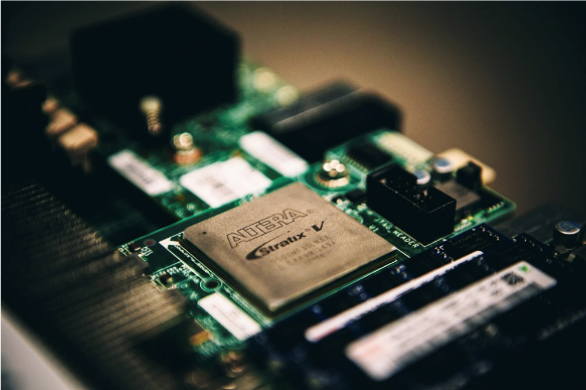
ROBERT MCMILLAN

BUSINESS

06.16.14

6:30 AM

# MICROSOFT SUPERCHARGES BING SEARCH WITH PROGRAMMABLE CHIPS




Microsoft


DOUG BURGER CALLED it Project Catapult.

Burger works inside Microsoft Research—the group where the tech giant explores blue-sky ideas—and in November 2012, he pitched a radical new concept to Qi Lu, the man who oversees Microsoft’s Bing web search engine. He wanted to


MOST POPULAR



MOBILE  
Android Can't Compete With iMessage. Google Is Changing That  
DAVID PIERCE



SCANDALS  
Google Accuses Uber of Stealing Its Self-Driving Car Tech  
ALEX DAVIES





PRODUCT REVIEW  
Review: Microsoft Surface Studio  
DAVID PIERCE

→

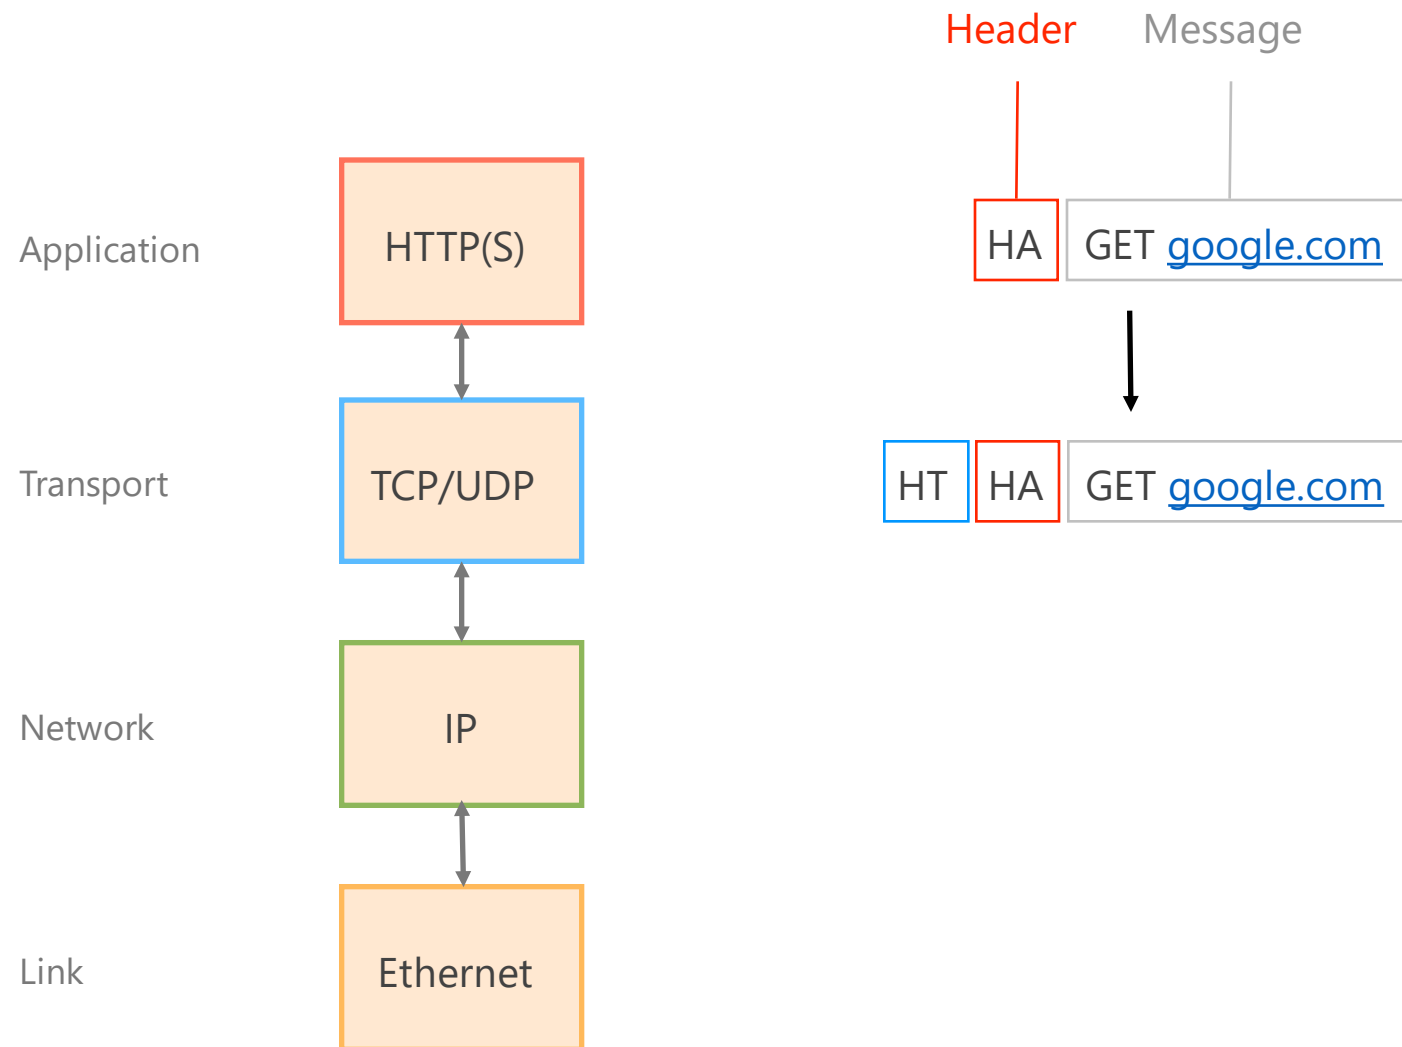
MORE STORIES

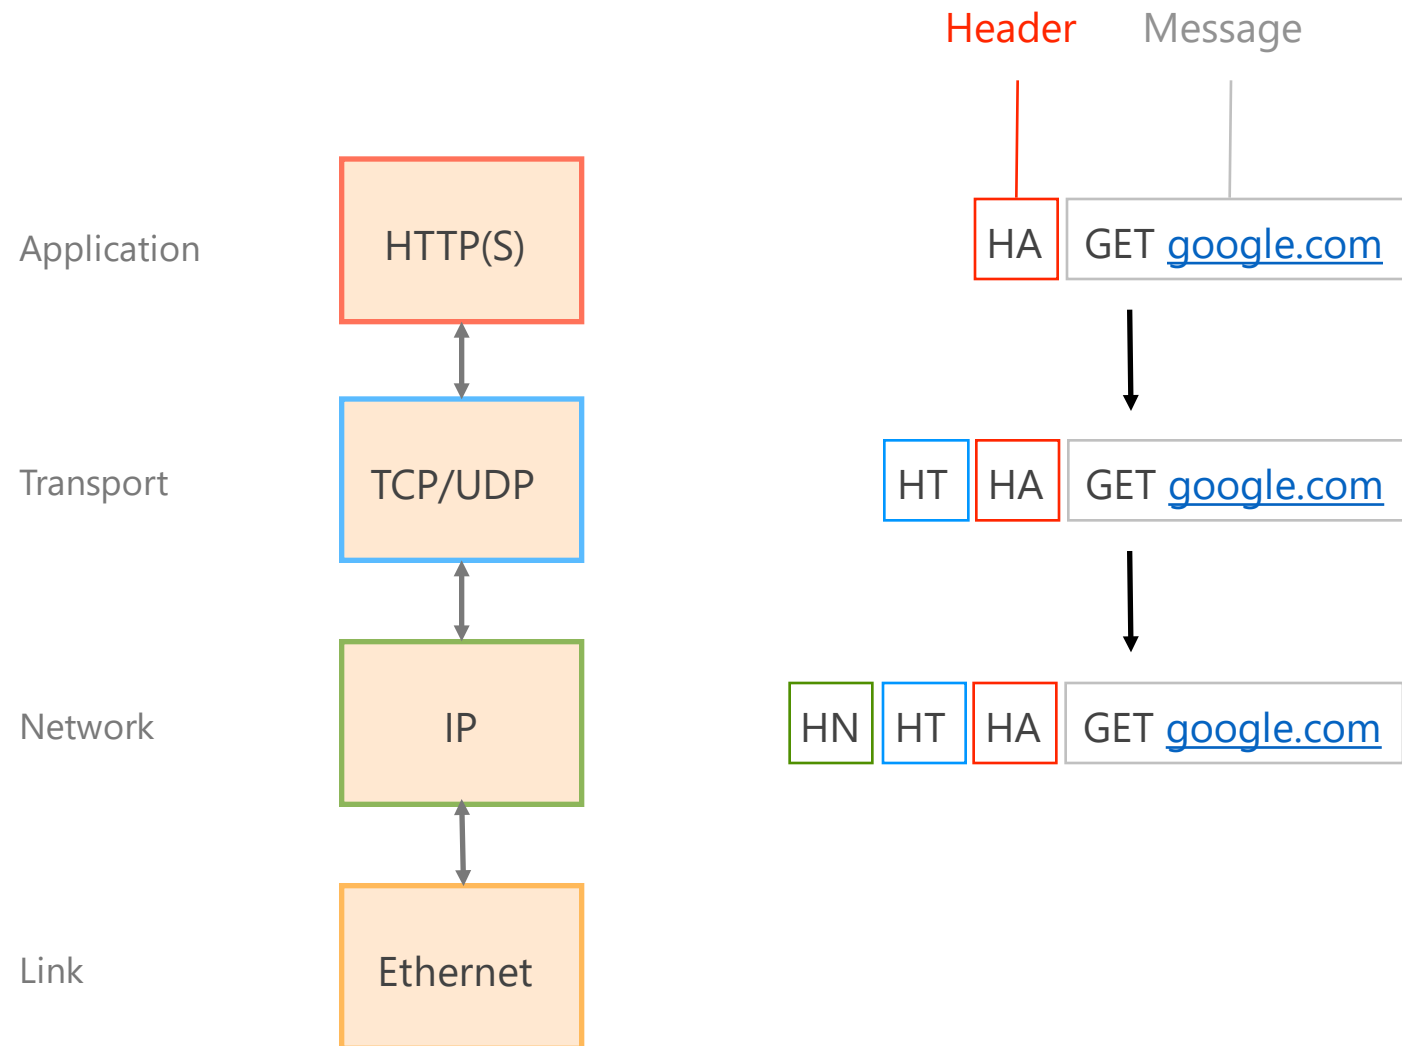
Each layer takes messages from the layer above,  
and *encapsulates* with its own header and/or trailer

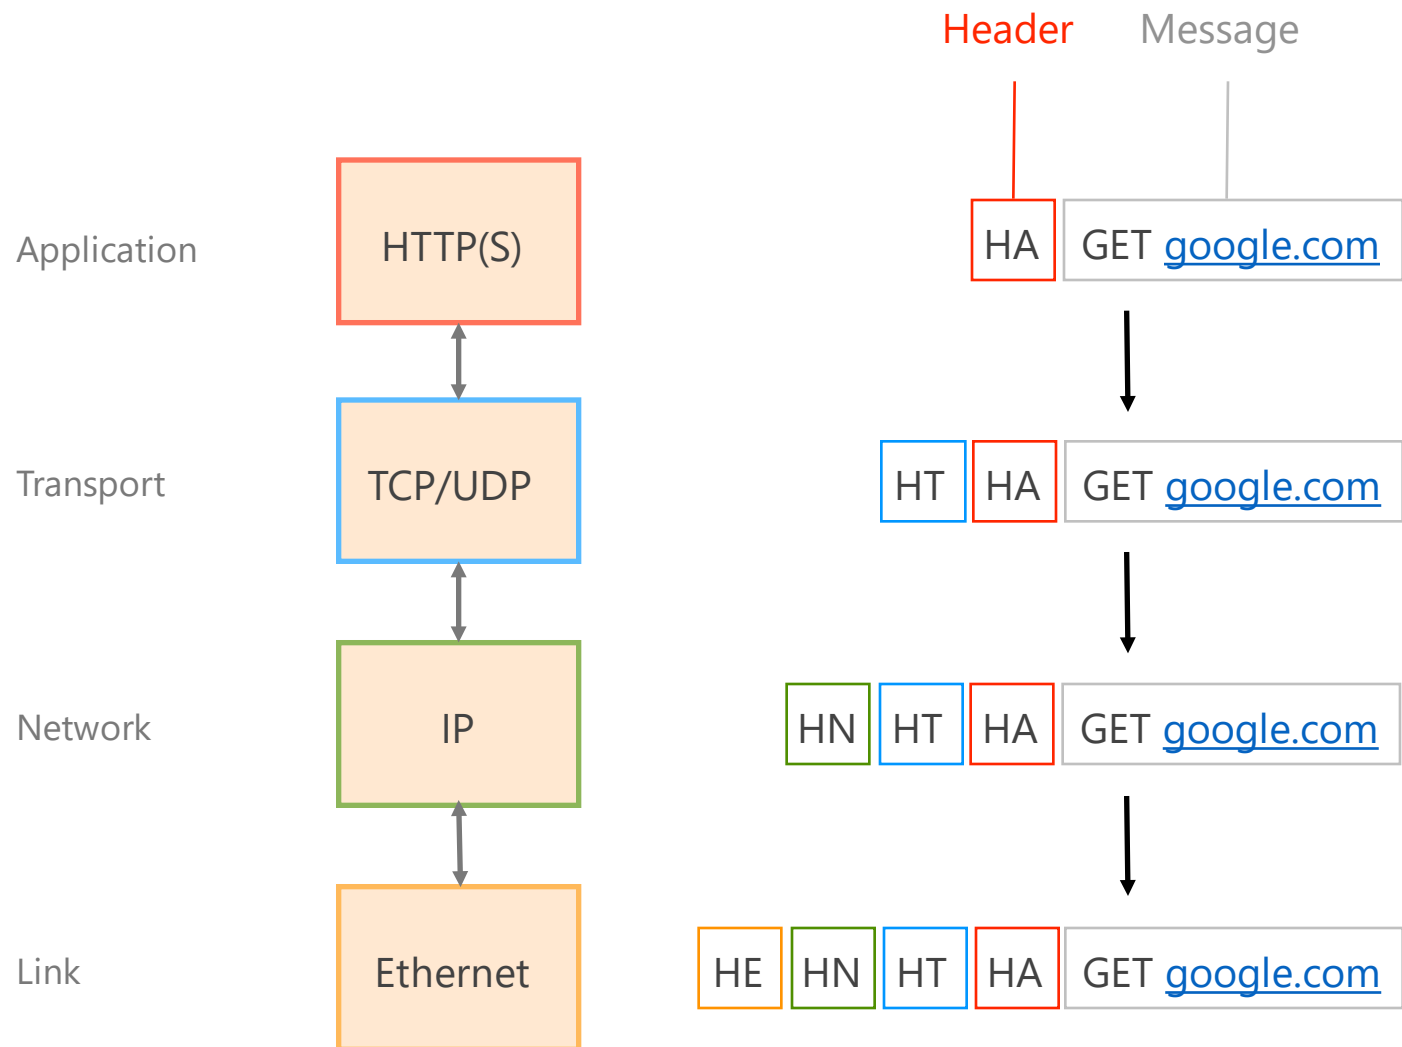




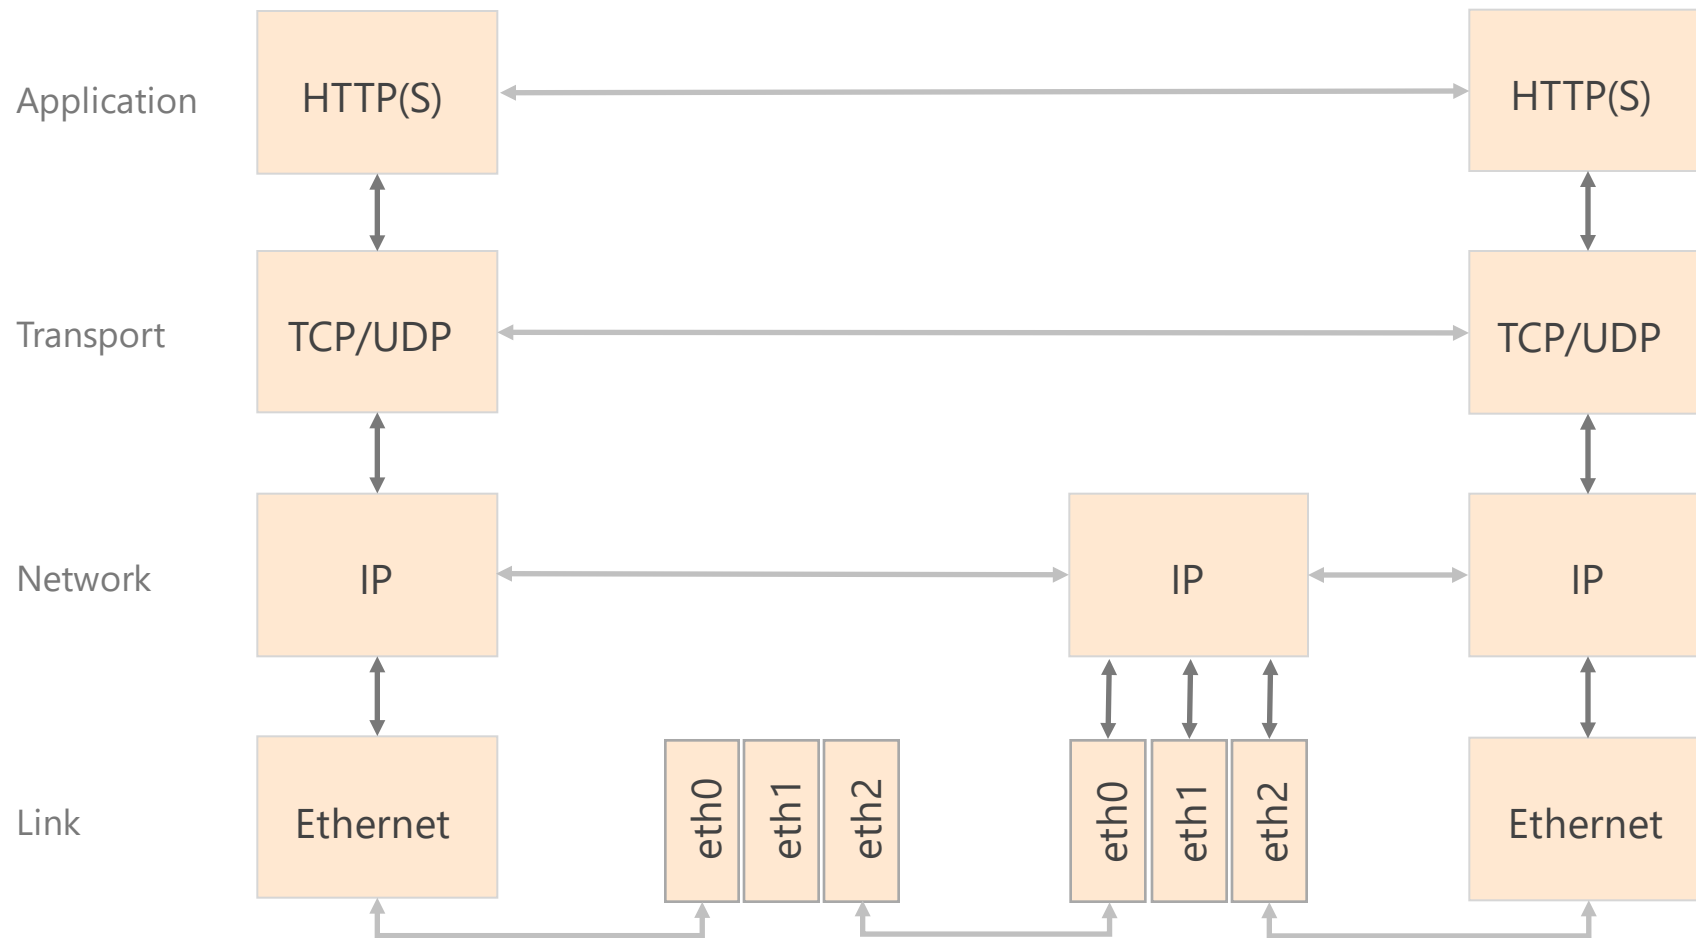




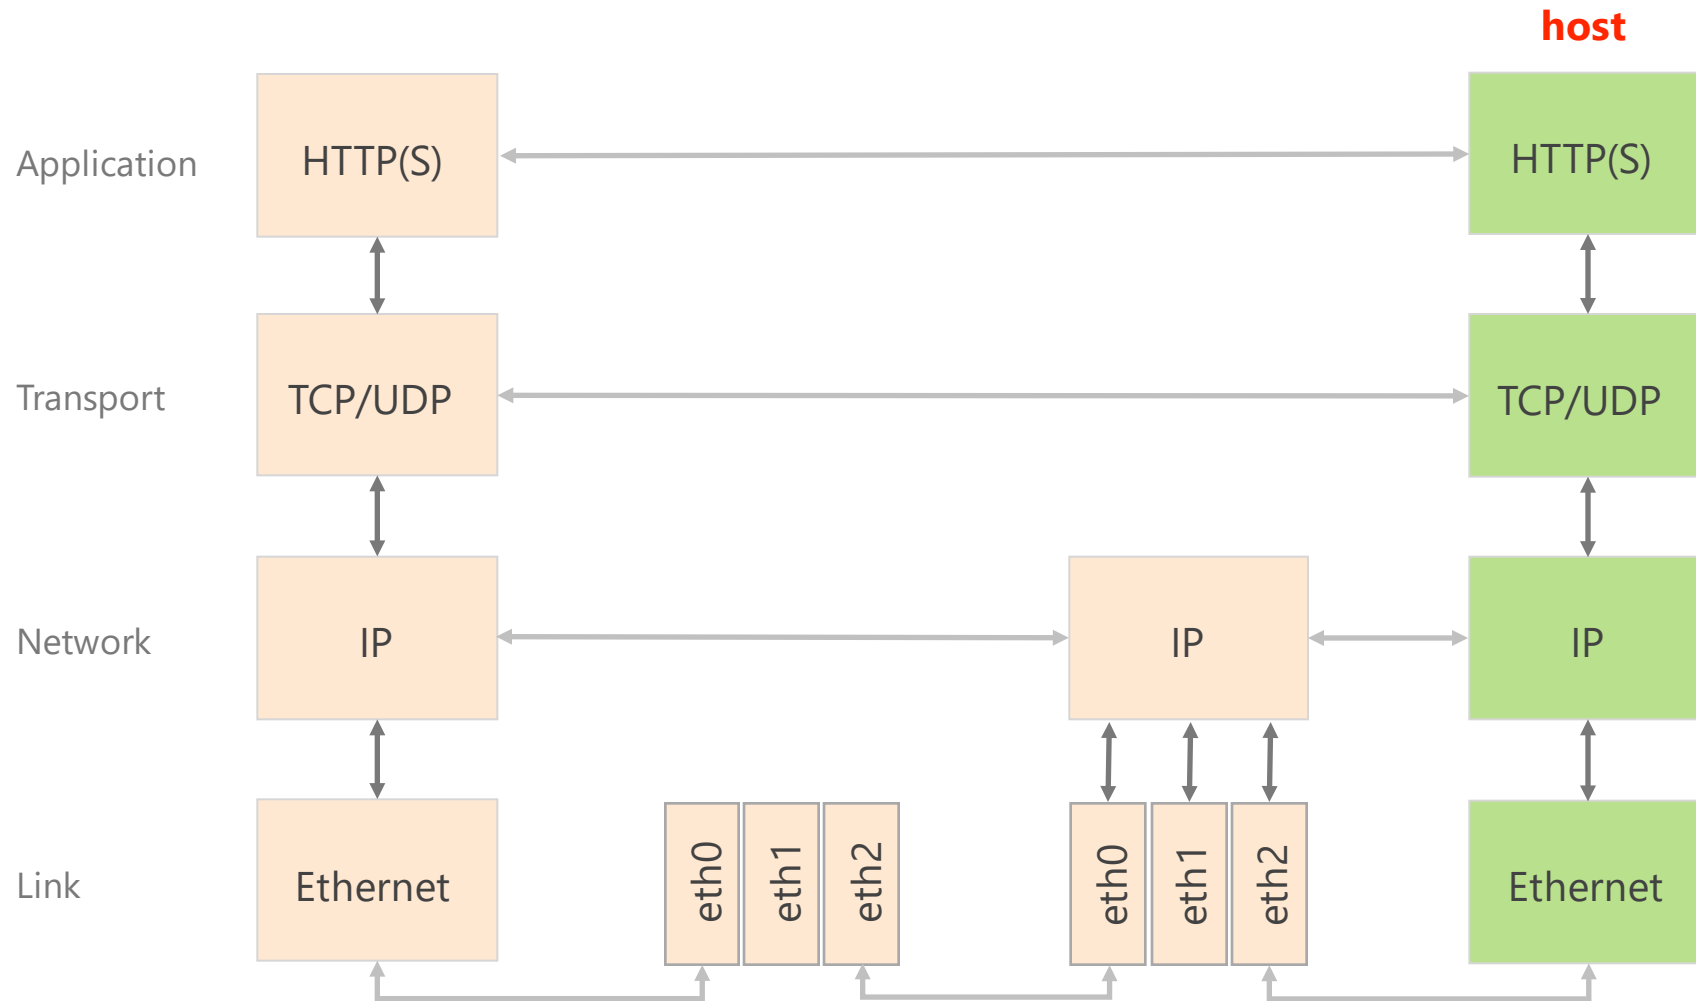




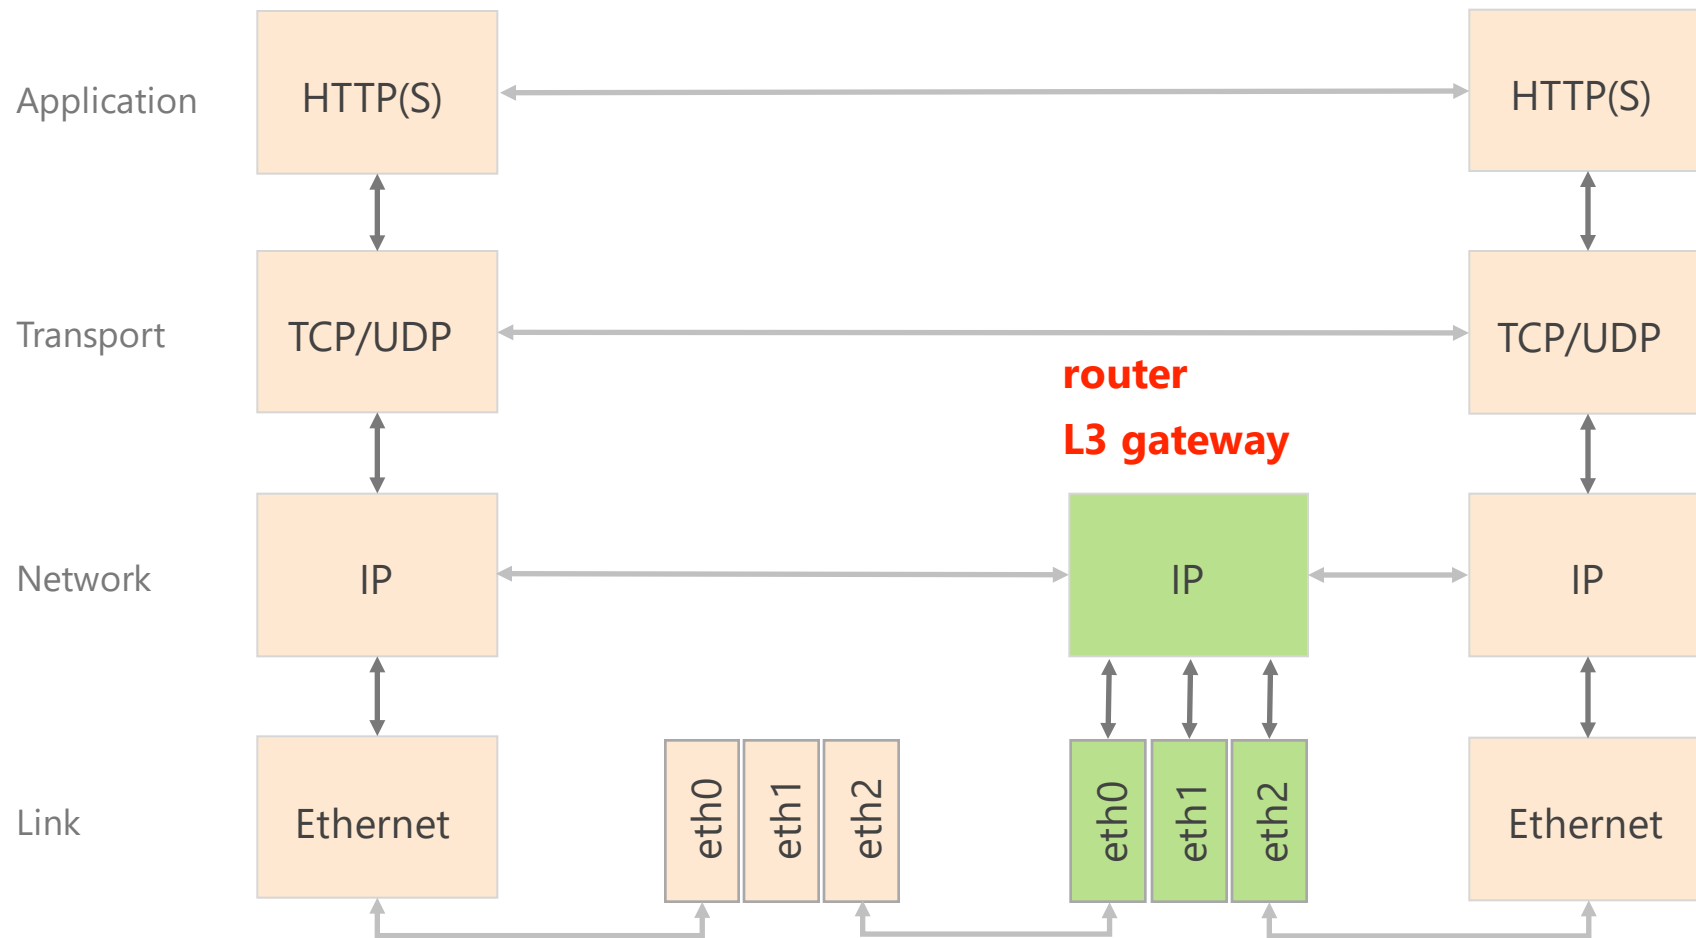
In practice, layers are distributed on every network device



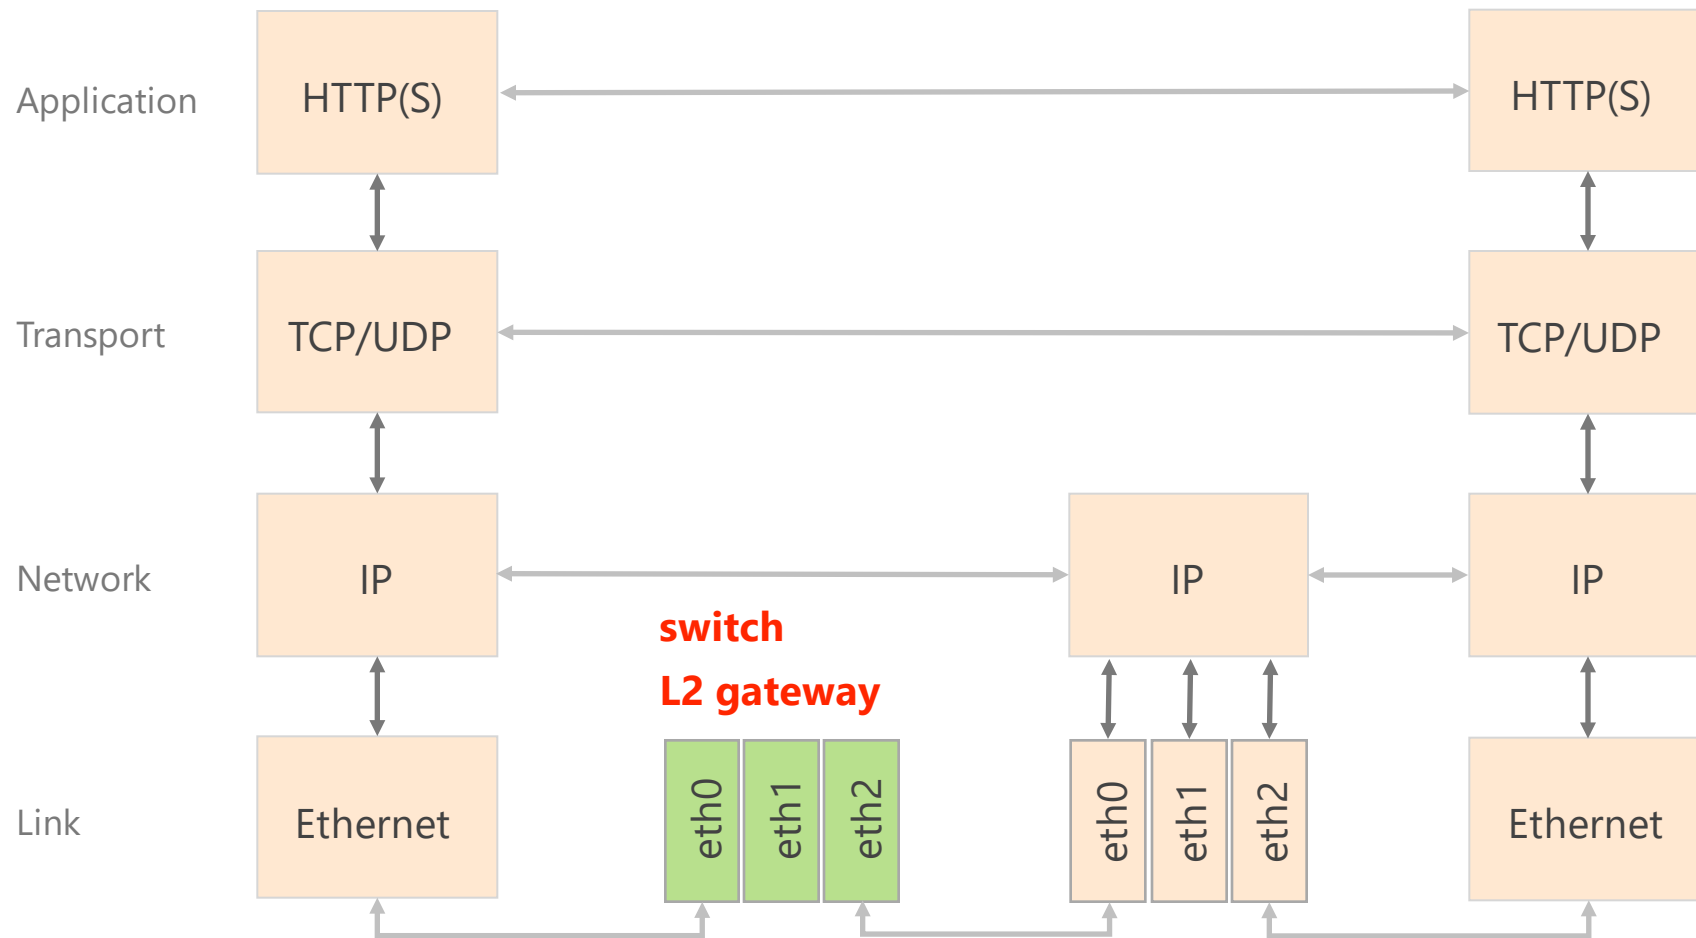
Since when bits arrive they must make it to the application, all the layers exist on a host



Routers act as **L3 gateway**  
as such they implement L2 and L3



Switches act as **L2 gateway**  
as such they only implement L2



Let's see how it looks like in practice

on a host, using **Wireshark**

<https://www.wireshark.org>





# Communication Networks and Internet Technology

## Part 1: General overview

What is a network made of?

How is it shared?


How is it organized?

How does communication happen?

#5

How do we characterize it?

A network *connection* is characterized by its delay, loss rate and throughput




How long does it take for a packet to reach the destination

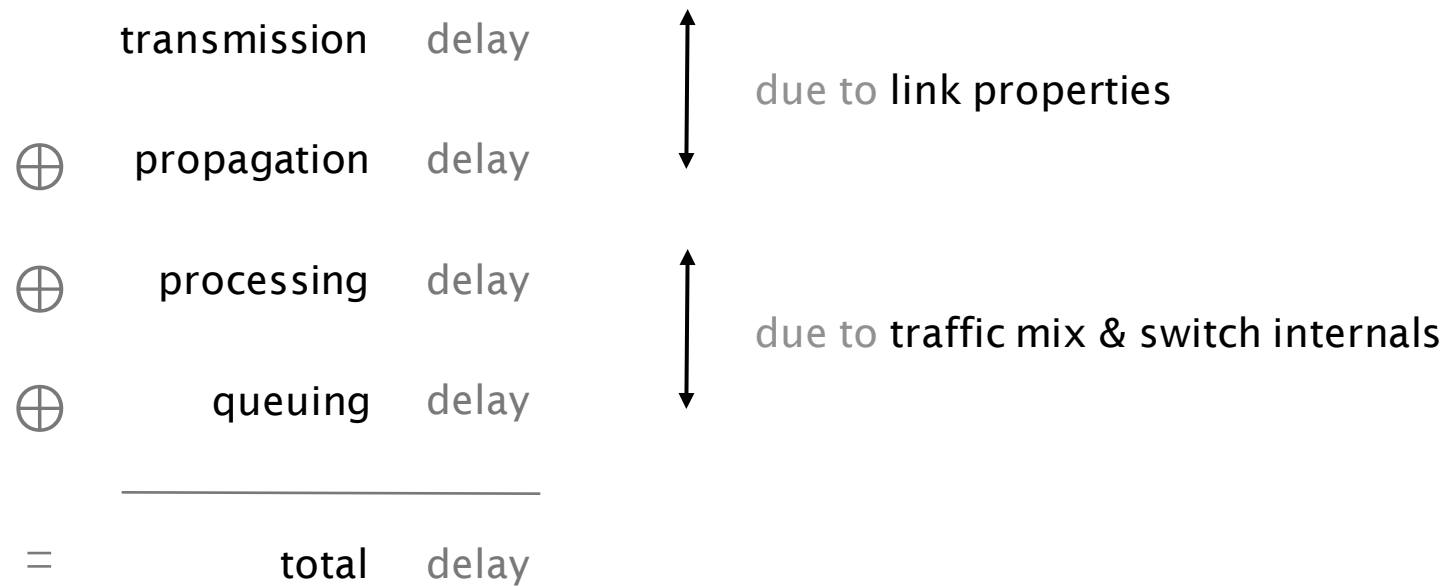
What fraction of packets sent to a destination are dropped?

At what rate is the destination receiving data from the source?

A network *connection* is characterized by its delay, loss rate and throughput

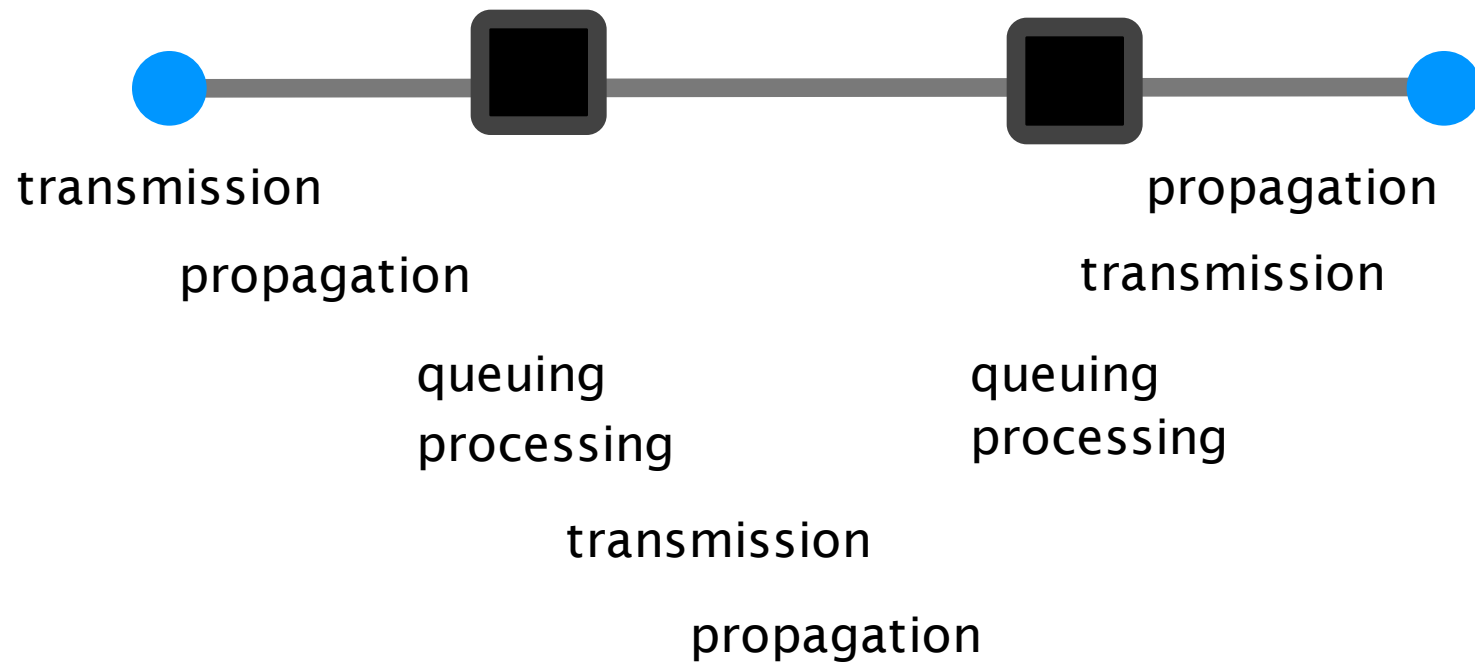


Each packet suffers from several types of delays  
at *each node* along the path



Overall, the main culprits for the overall delay are the transmission, propagation and queuing delays

$$\begin{array}{rcll} & \text{transmission} & \text{delay} & \\ \oplus & \text{propagation} & \text{delay} & \\ \oplus & \text{processing} & \text{delay} & \textit{tend to be tiny} \\ \oplus & \text{queuing} & \text{delay} & \\ \hline = & \text{total} & \text{delay} & \end{array}$$



The transmission delay is the amount of time required to push all of the bits onto the link

$$\begin{array}{lcl} \text{Transmission delay} & = & \frac{\text{packet size}}{\text{link bandwidth}} \\ \text{[sec]} & & \begin{array}{l} \text{[#bits]} \\ \text{[#bits/sec]} \end{array} \end{array}$$

$$\begin{array}{lcl} \text{Example} & & \frac{1000 \text{ bits}}{100 \text{ Gbps}} \\ & & 10 \text{ ns} \end{array}$$

The propagation delay is the amount of time required for a bit to travel to the end of the link


$$\begin{array}{lcl} \text{Propagation delay} & = & \frac{\text{link length}}{\text{propagation speed}} \\ [\text{sec}] & & \begin{array}{l} \text{[m]} \\ \text{[m/sec]} \\ \text{(fraction of speed of light)} \end{array} \end{array}$$

$$\begin{array}{lcl} \text{Example} & & \frac{30\,000\text{ m}}{2 \times 10^8\text{ m/sec}} \\ & & \begin{array}{l} \text{(speed of light in fiber)} \\ 150\text{ }\mu\text{sec} \end{array} \end{array}$$



A network *link* is characterized by its bandwidth and propagation delay

link bandwidth	# bits sent per sec
propagation delay	time for 1 bit to move through the link (sec)
bandwidth $\otimes$ delay	# of bits in-flight at any time



## In practice, BDP can be huge

same city, over slow link

bandwidth      100 Mbps

delay            0.1 ms

BDP             10,000 bits  
                    1.25 KBytes

cross continent, over fast link


bandwidth      10 Gbps

delay            10 ms


BDP              $10^8$  bits  
                    12.5 MBytes

How long does it take for a packet to travel from A to B?  
(not considering queuing for now)

How long does it take to exchange 100 Bytes packet?



If we have a 1 Gbps link,  
the total time decreases to **1.0008ms**



If we now exchange a 1GB file  
split in 100B packets



$10^7 \times 100\text{B}$  packets

The last bit  
reaches B at  
 $(10^7 \times 800 \times 10^{-9})$   
 $+ 10^{-3}\text{s}$   
 $= 8001\text{ms}$

Different transmission characteristics imply  
different tradeoffs in terms of which delay dominates

$10^7 \times 100\text{B}$	pkt	1 Gbps link	transmission delay dominates
---------------------------	-----	-------------	------------------------------

$1 \times 100\text{B}$	pkt	1 Gbps link	propagation delay dominates
------------------------	-----	-------------	-----------------------------

$1 \times 100\text{B}$	pkt	1 Mbps link	both matter
------------------------	-----	-------------	-------------

In the Internet, we **can't know** in advance which one matters!


The queuing delay is the amount of time a packet waits (in a buffer) to be transmitted on a link


Queuing delay is the hardest to evaluate  
as it varies from packet to packet

It is characterized with statistical measures  
e.g., average delay & variance, probability of exceeding  $x$




Queuing delay depends on the traffic pattern








Queuing delay depends on the traffic pattern





Queuing delay depends on the traffic pattern














Queues absorb transient bursts,  
but introduce queueing delays



The time a packet has to sit in a buffer before being processed depends on the traffic pattern

Queueing delay depends on:

- arrival rate at the queue
- transmission rate of the outgoing link
- traffic burstiness

average packet arrival rate	$a$	[packet/sec]
-----------------------------	-----	--------------

transmission rate of outgoing link	$R$	[bit/sec]
------------------------------------	-----	-----------

fixed packets length	$L$	[bit]
----------------------	-----	-------

average bits arrival rate	$L \cdot a$	[bit/sec]
---------------------------	-------------	-----------


traffic intensity	$L \cdot a/R$	
-------------------	---------------	--

When the **traffic intensity is  $>1$** , the queue will increase without bound, and so does the queuing delay


Golden rule

Design your queuing system,  
so that it operates far from that point


When the **traffic intensity is  $\leq 1$** ,  
queueing delay depends on the burst size



A network *connection* is characterized by its delay, loss rate and throughput




In practice, queues are not infinite.  
There is an upper bound on queuing delay.




queuing delay upper bound:  $N \cdot L/R$

If the queue is persistently overloaded,  
it will eventually drop packets (loss)






A network *connection* is characterized by its delay, loss rate and throughput



The throughput is the instantaneous rate at which a host receives data

$$\begin{array}{lcl} \text{Average throughput} & = & \frac{\text{data size}}{\text{transfer time}} \\ \text{[#bits/sec]} & & \begin{array}{l} \text{[#bits]} \\ \text{[sec]} \end{array} \end{array}$$

To compute throughput, one has to consider the bottleneck link




Average throughput

$$\min(R_s, R_L)$$


= transmission rate  
of the bottleneck link

To compute throughput, one has to consider the bottleneck link... and the intervening traffic

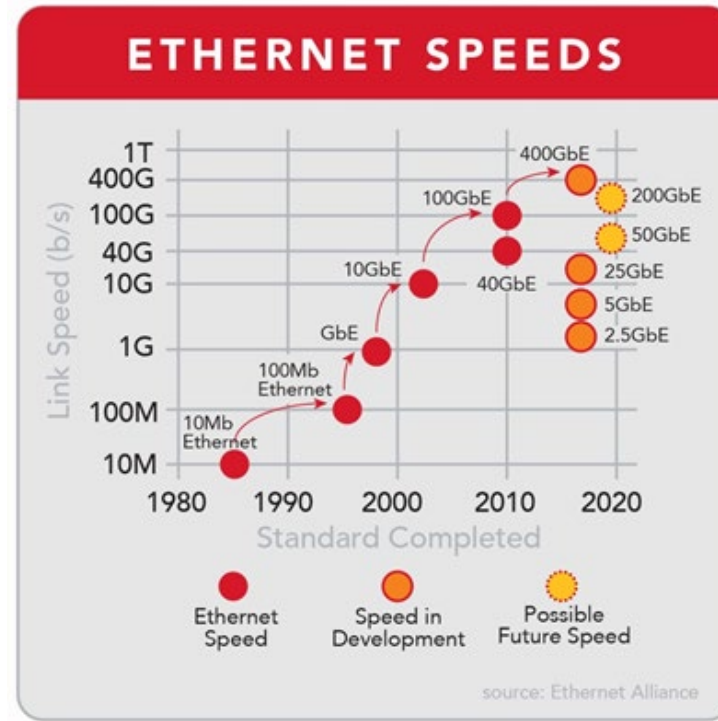


if  $4 \cdot \min(R_S, R_L) > R$  the bottleneck is now in the core,  
providing each download  $R/4$  of throughput

A network *connection* is characterized by its delay, loss rate and throughput

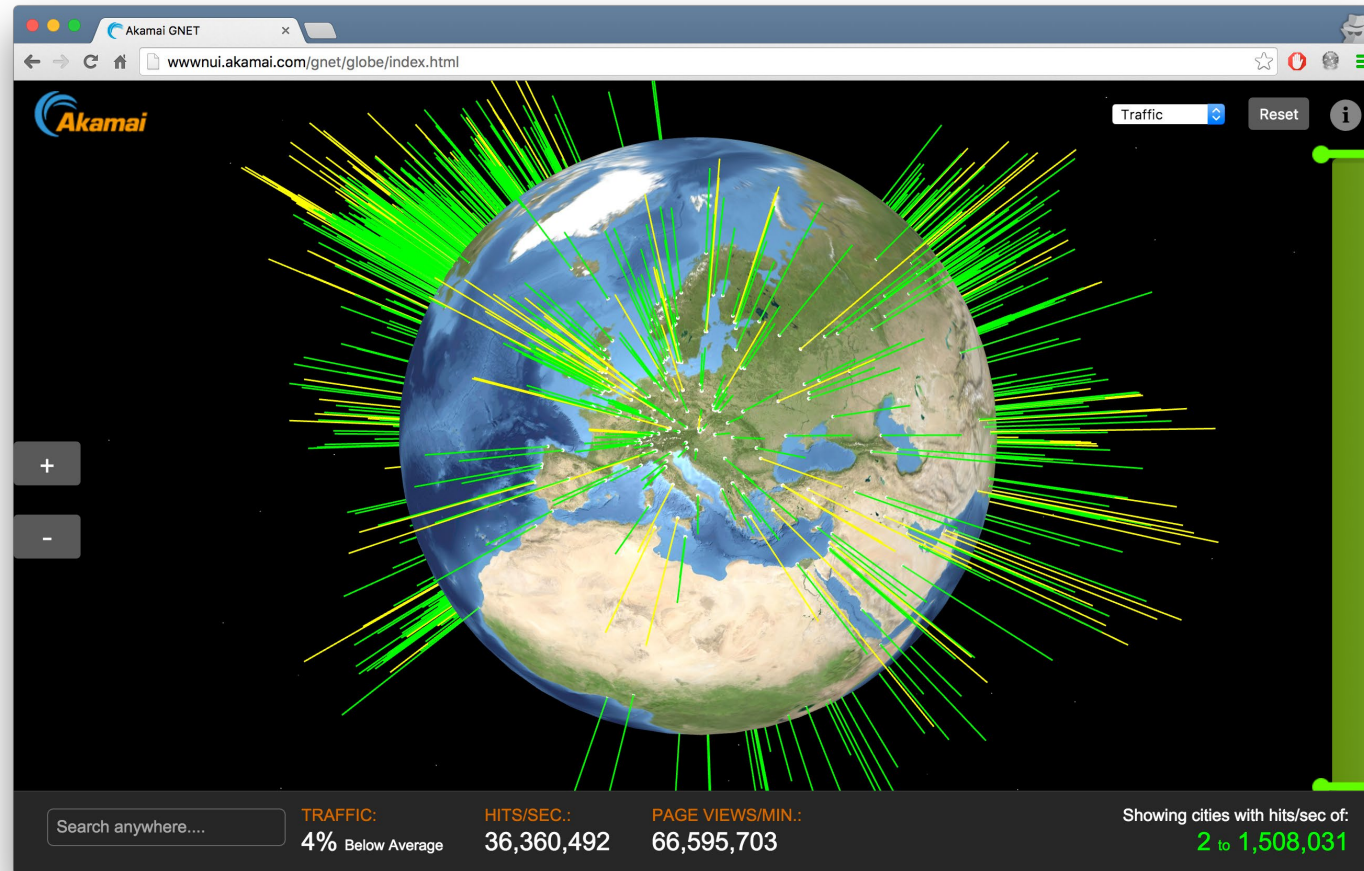


As technology improves, throughput increase & delays are getting lower except for propagation  
(speed of light)



source: ciena.com

Because of propagation delays,  
Content Delivery Networks move content closer to you



# Communication Networks and Internet Technology

## Part 1: General overview

What is a network made of?

How is it shared?

How is it organized?

How does communication happen?

How do we characterize it?



The basic protocols underlying the Internet  
are *intuitive*

The principles behind the Internet are  
*more about architecture than engineering*

Principles

Interconnect many different networks

Ethernet, Optical Fibers, wireless, ...

Scale to the entire world

both geographically and numerically

Tolerate and recover from failures

both constant and inevitable

The principles behind the Internet are  
*more about architecture than engineering*



Architecture

The diagram consists of two orange rectangular boxes side-by-side. Below the left box is the text 'what tasks get done and where', and below the right box is the text 'how tasks get done'. The words 'what' and 'how' are in red, while 'tasks get done' and 'where' are in black.

*what* tasks get done  
and *where*

Engineering

*how* tasks get done

The principles behind the Internet are  
*more about architecture than engineering*

Architecture

*what* tasks get done  
and *where*  
in the network?  
in the hosts?

Engineering

*how* tasks get done  
with what technology?

Network engineering is all about **optimization**  
and **balancing tradeoffs**

Goals

Speed

Quality of Service

Cost

Security

Port density

...

Reliability

Too small timers will cause unnecessary retransmissions,  
too large timers will slow down the communication

The “right” value depends on the network conditions

Protocols have to be flexible and adapt to them

# Communication Networks and Internet Technology

## Part 2: Concepts



routing

reliable  
delivery


# Communication Networks and Internet Technology

## Part 2: Concepts



routing

How do you guide IP packets  
from a source to destination?



reliable  
delivery

How do you ensure reliable transport  
on top of best-effort delivery?



This week

routing

Next week

reliable  
delivery

How do you guide **IP packets**  
from a source to destination?

Think of IP packets as envelopes




Packet

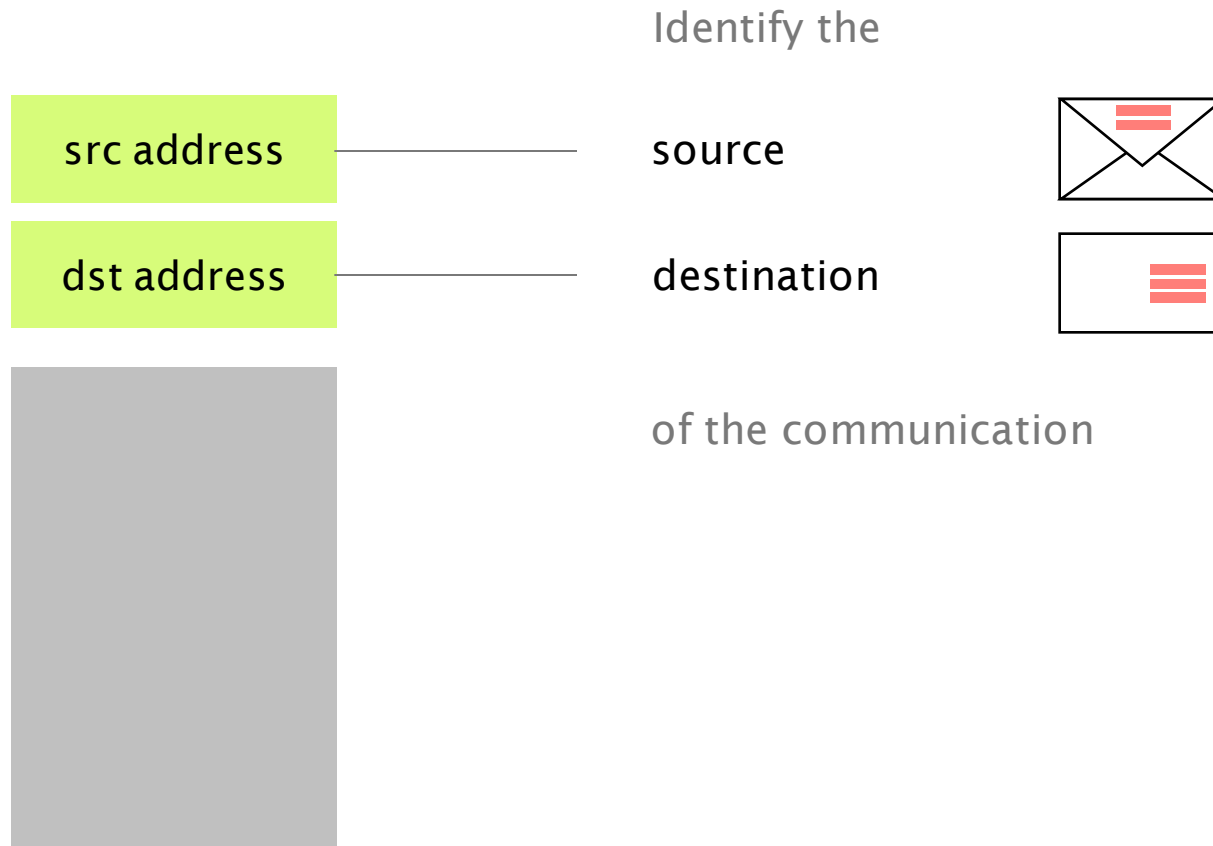
Like an envelope,  
packets have a header



Like an envelope,  
packets have a payload



The header contains the metadata  
needed to forward the packet



The payload contains  
the data to be delivered

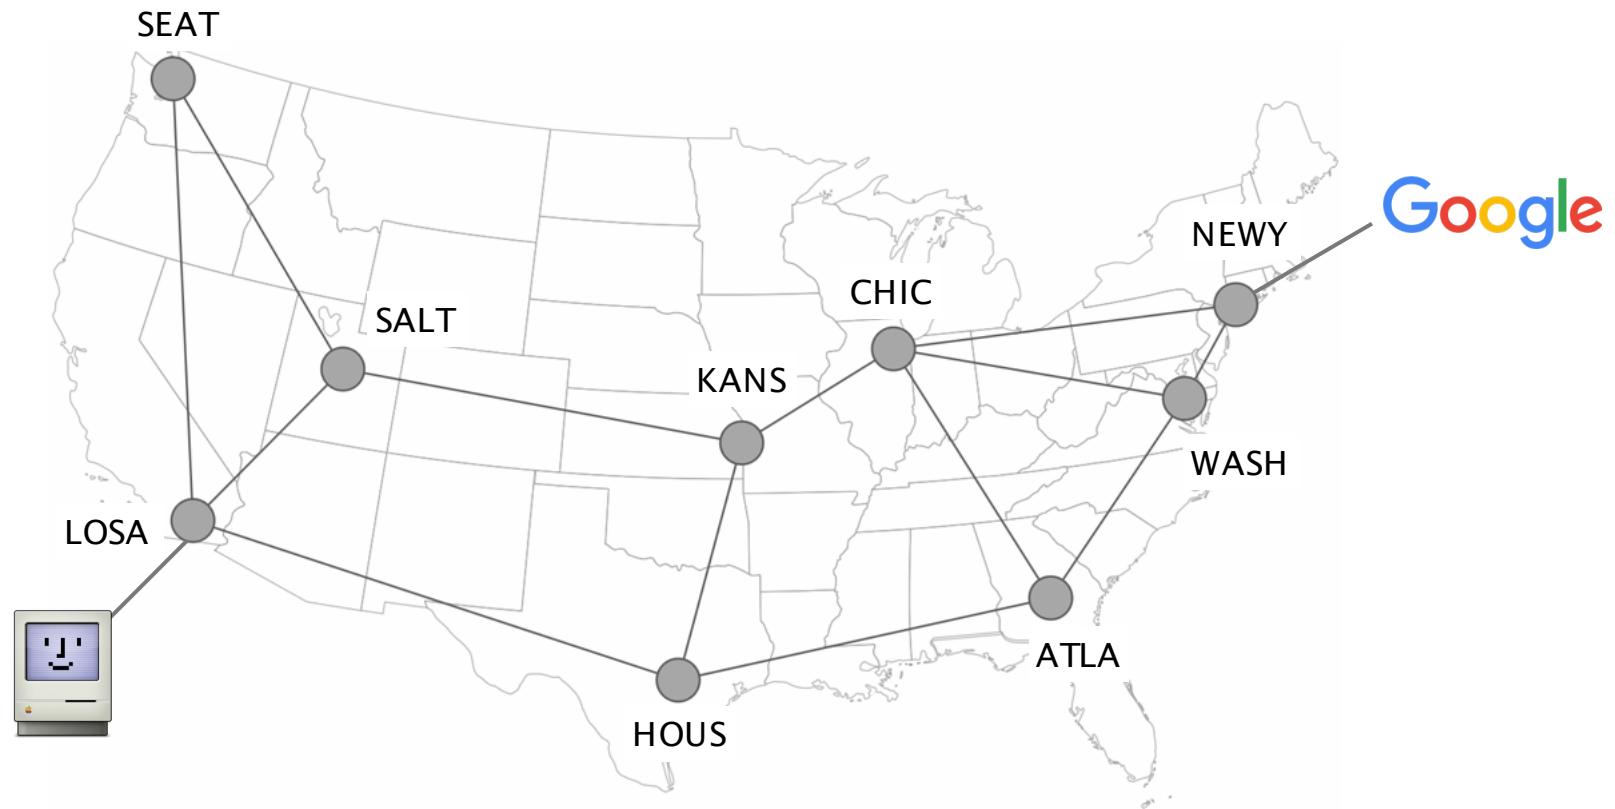


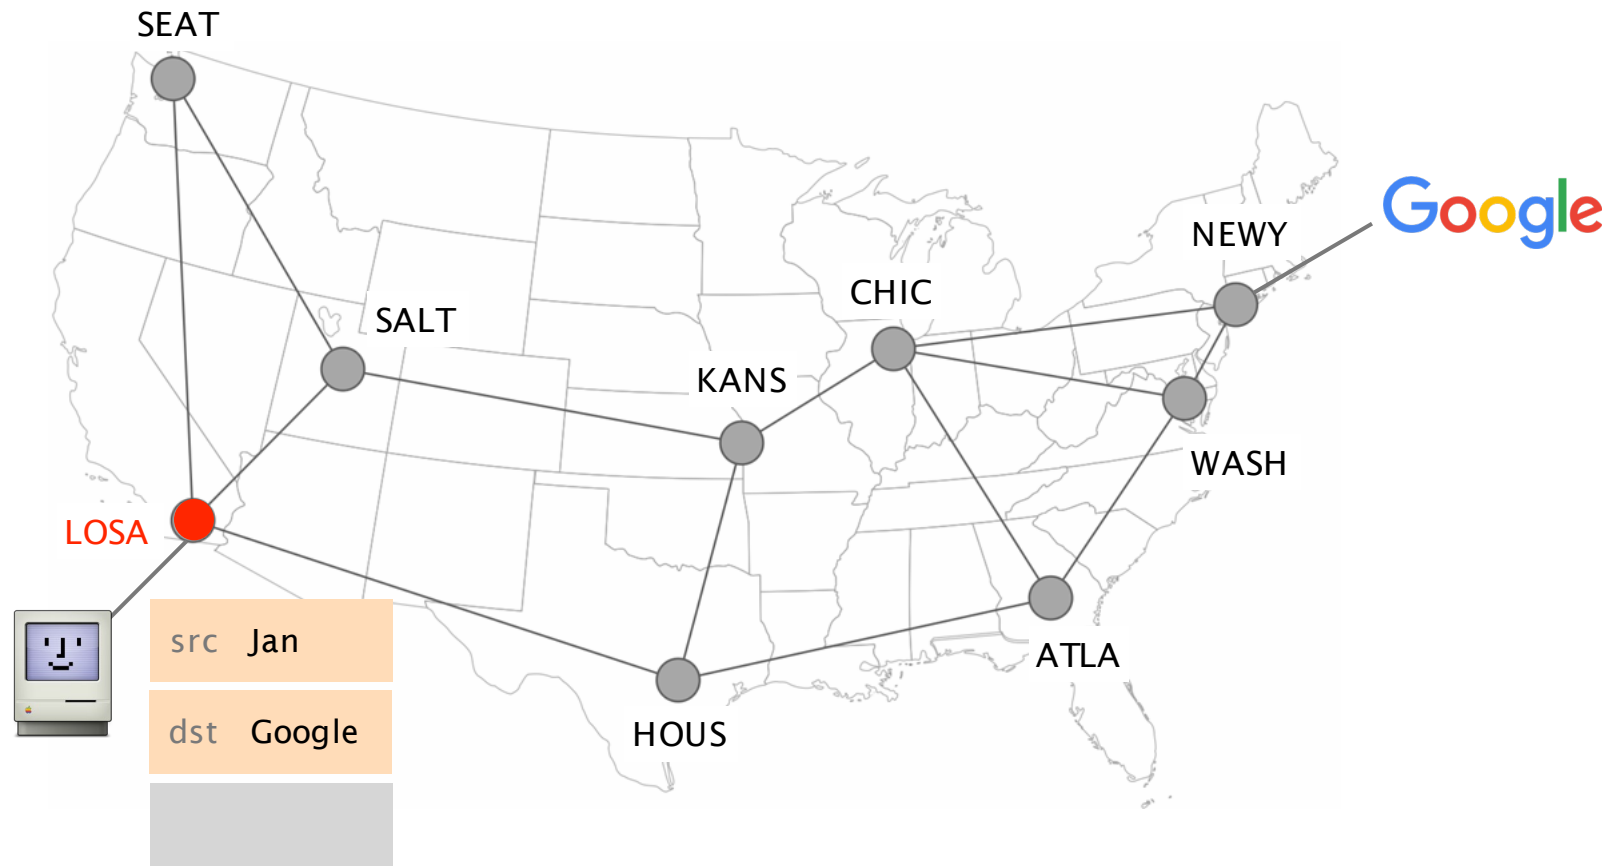
Payload

```
<html><head>
<meta http-equiv="content-type" content="text/html; charset=UTF-8">
<title>Google</title>
</head><body>
  
  <form action="/search" name=f>
    <input name=h1 type=hidden value=en>
    <input name=q size=55 title="Google Search" value="">
    <input name=btnG type=submit value="Google Search">
    <input name=btnI type=submit value="I'm Feeling Lucky">
  </form>
</body></html>
```

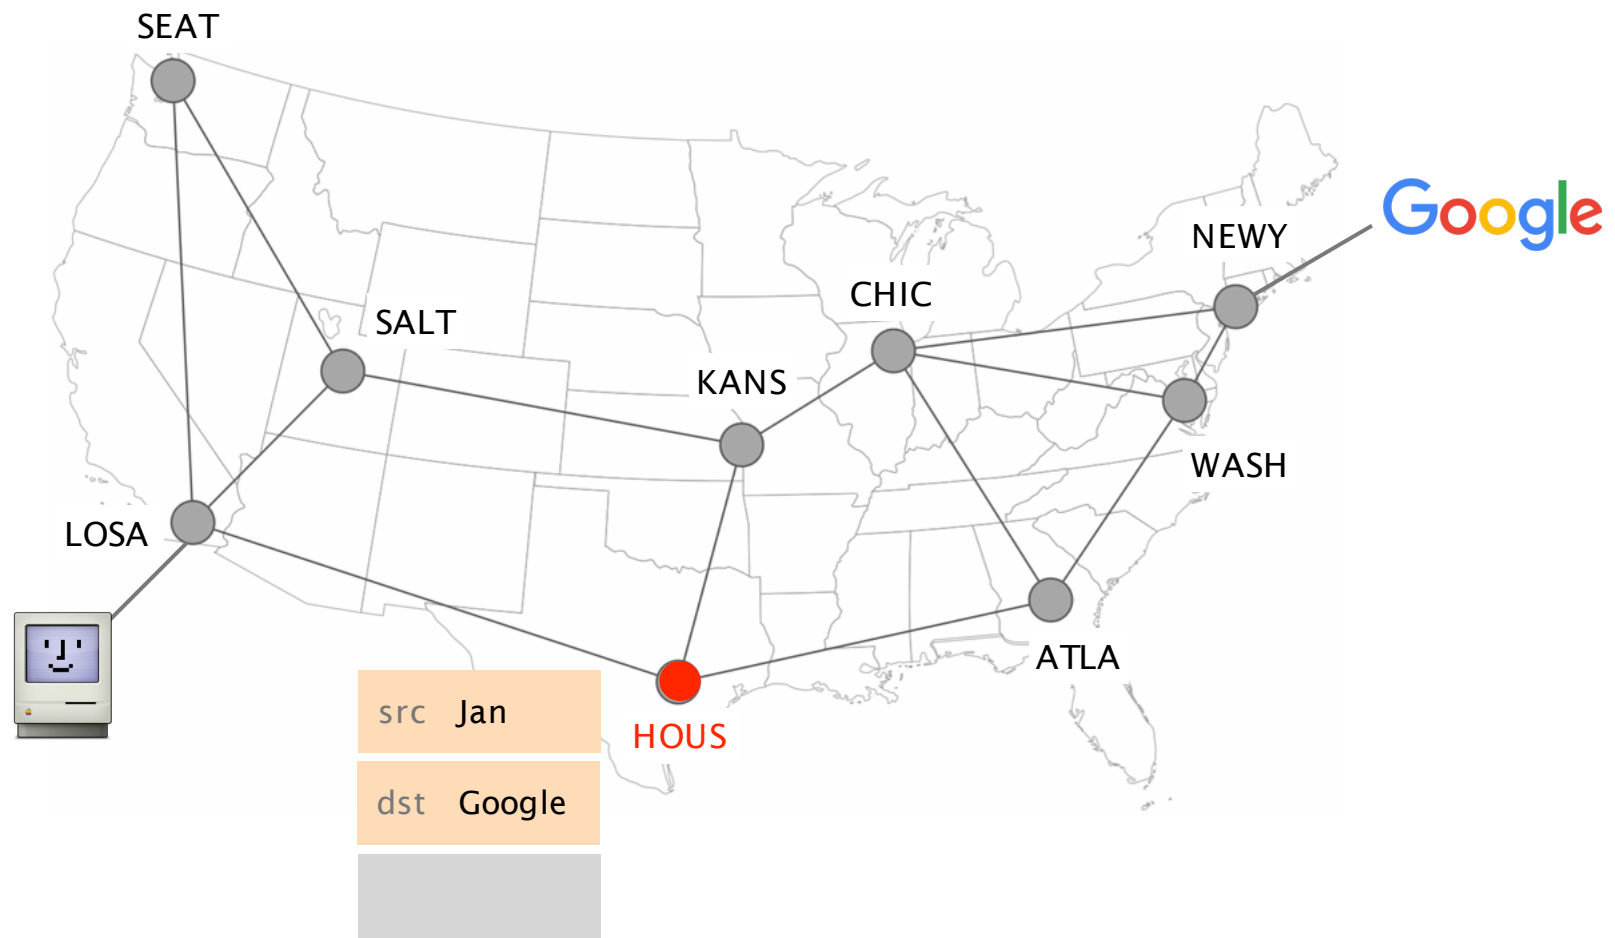


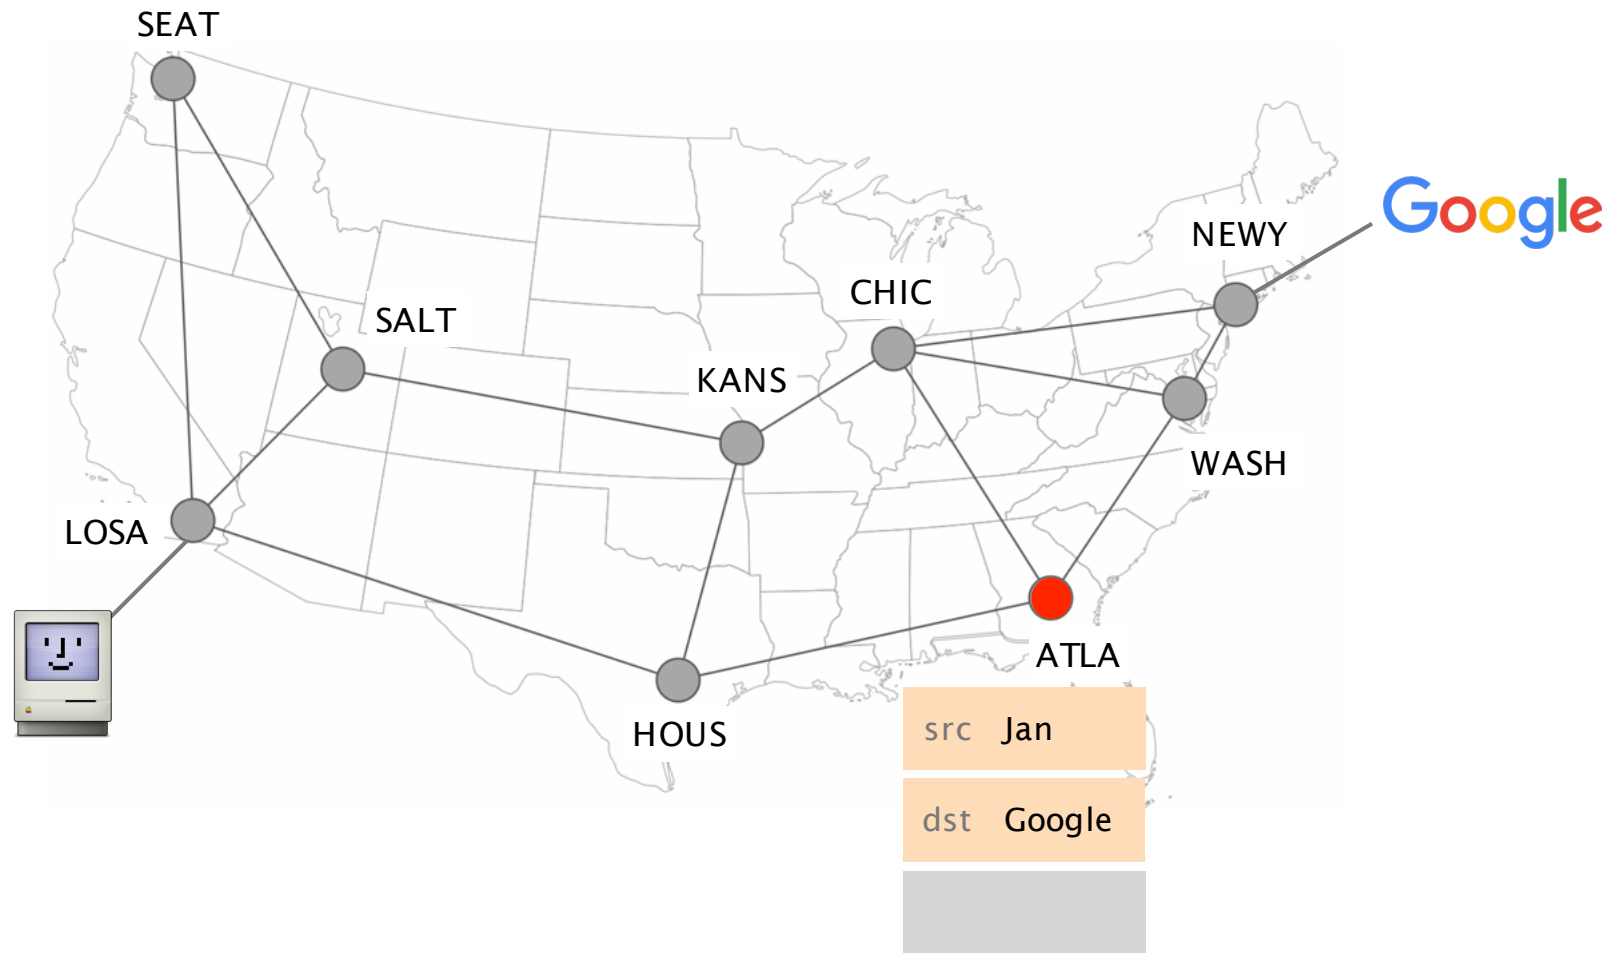
Routers forward IP packets hop-by-hop  
towards their destination

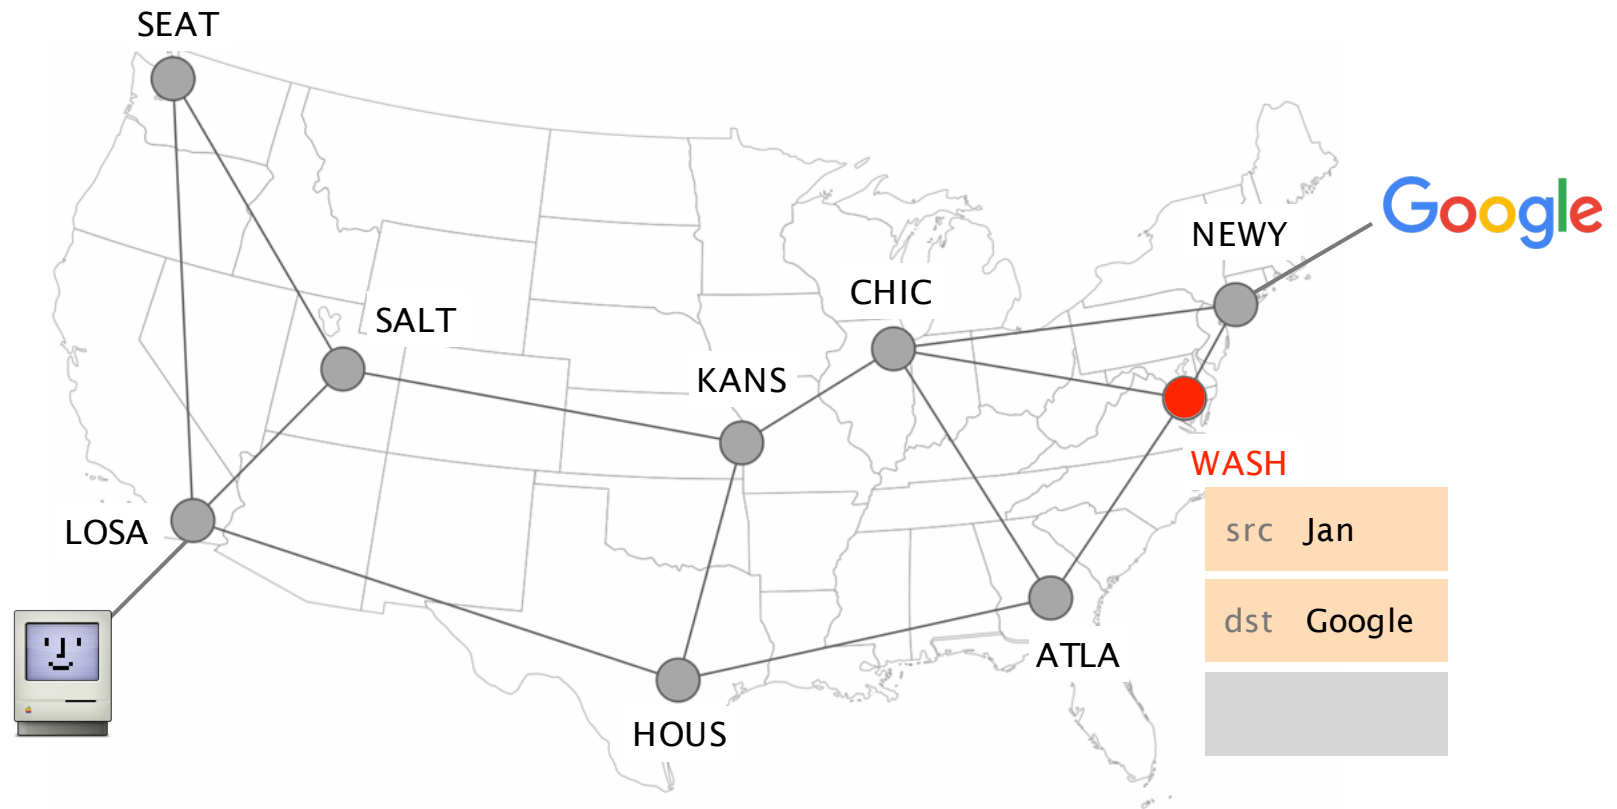


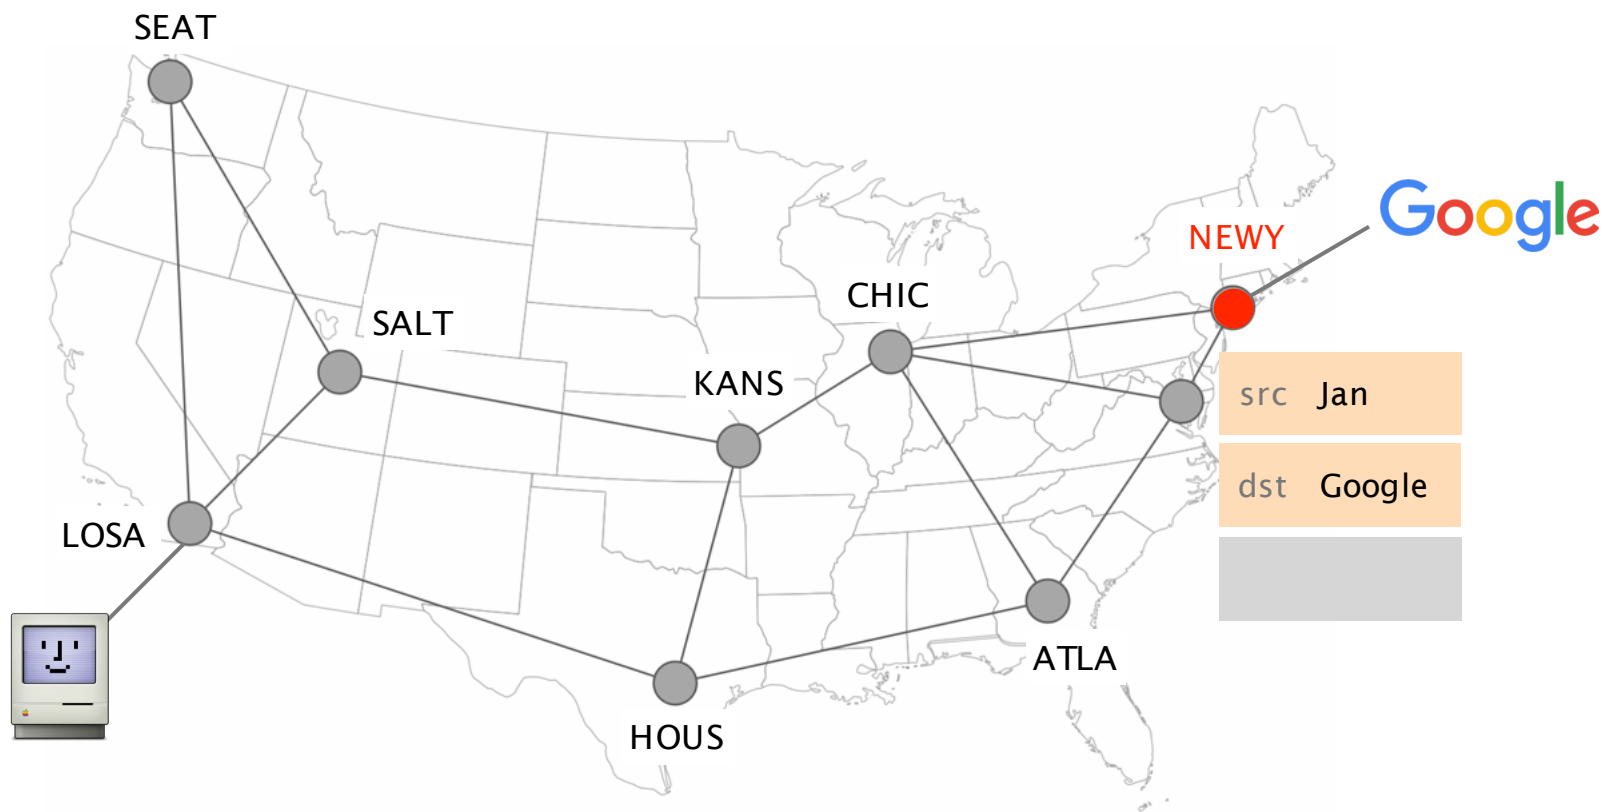


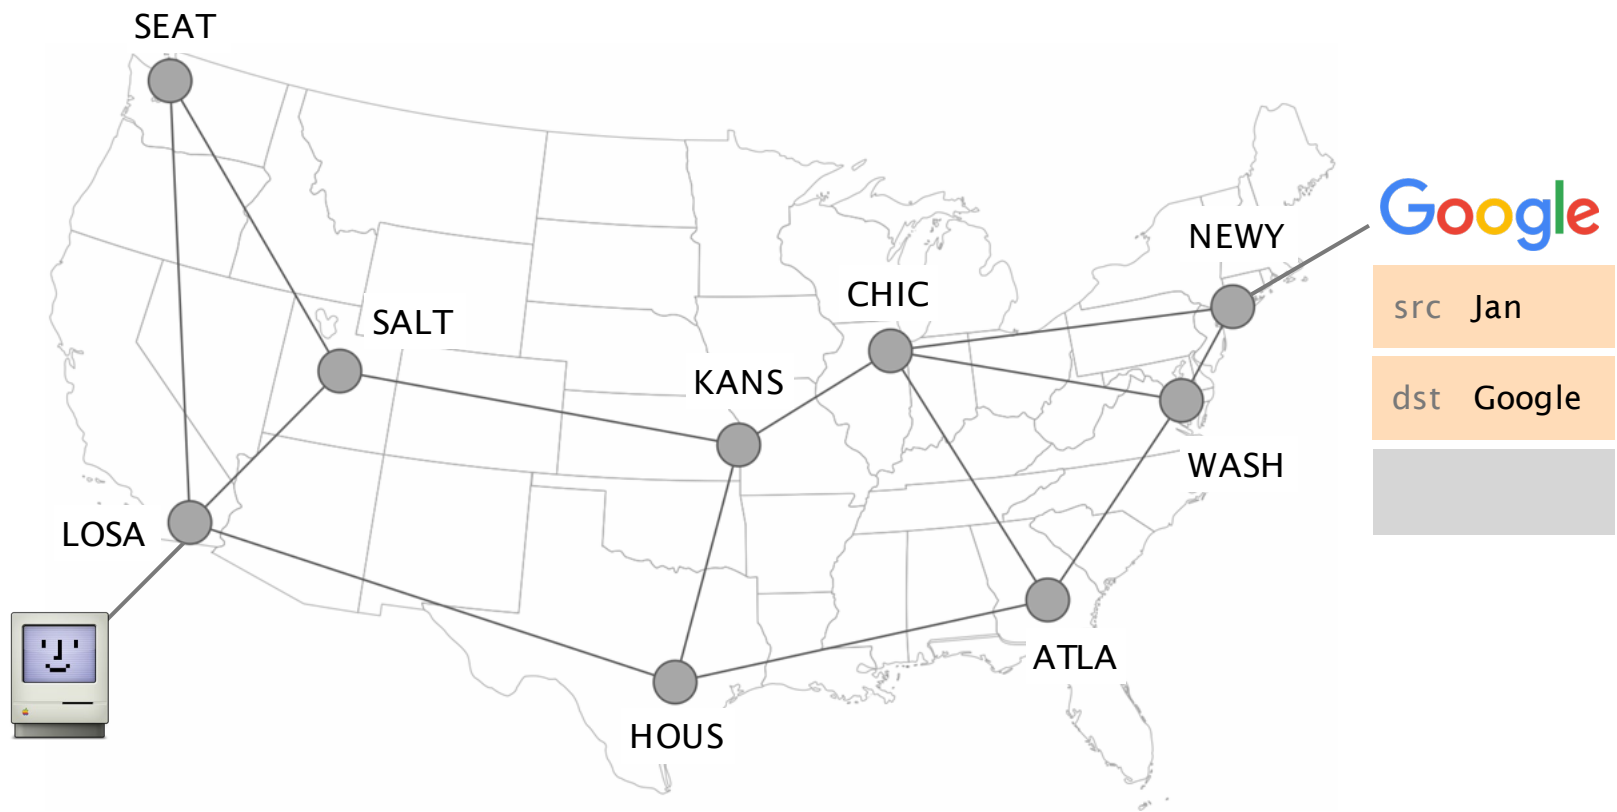




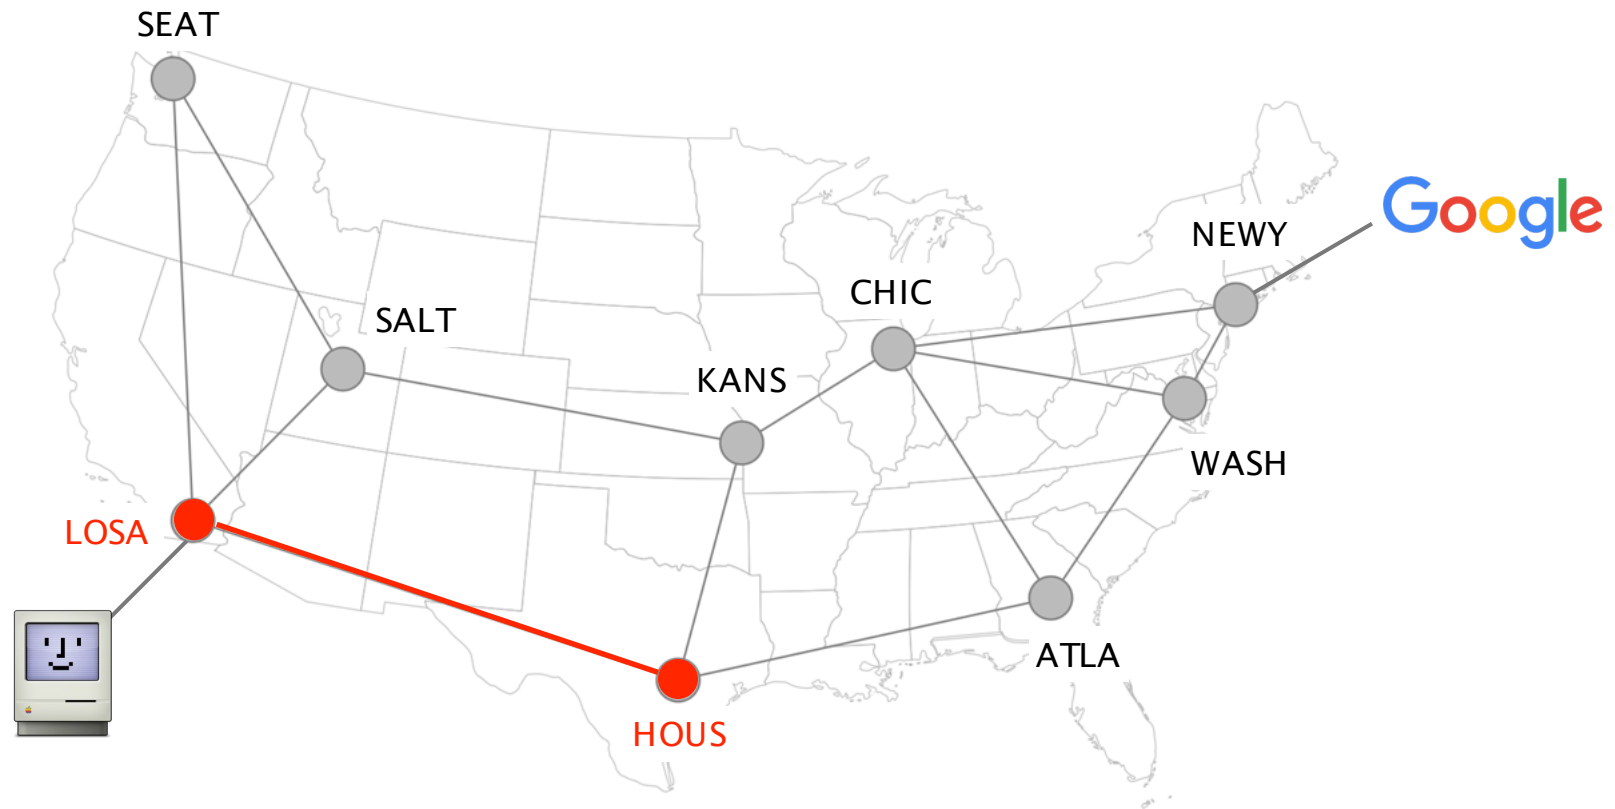







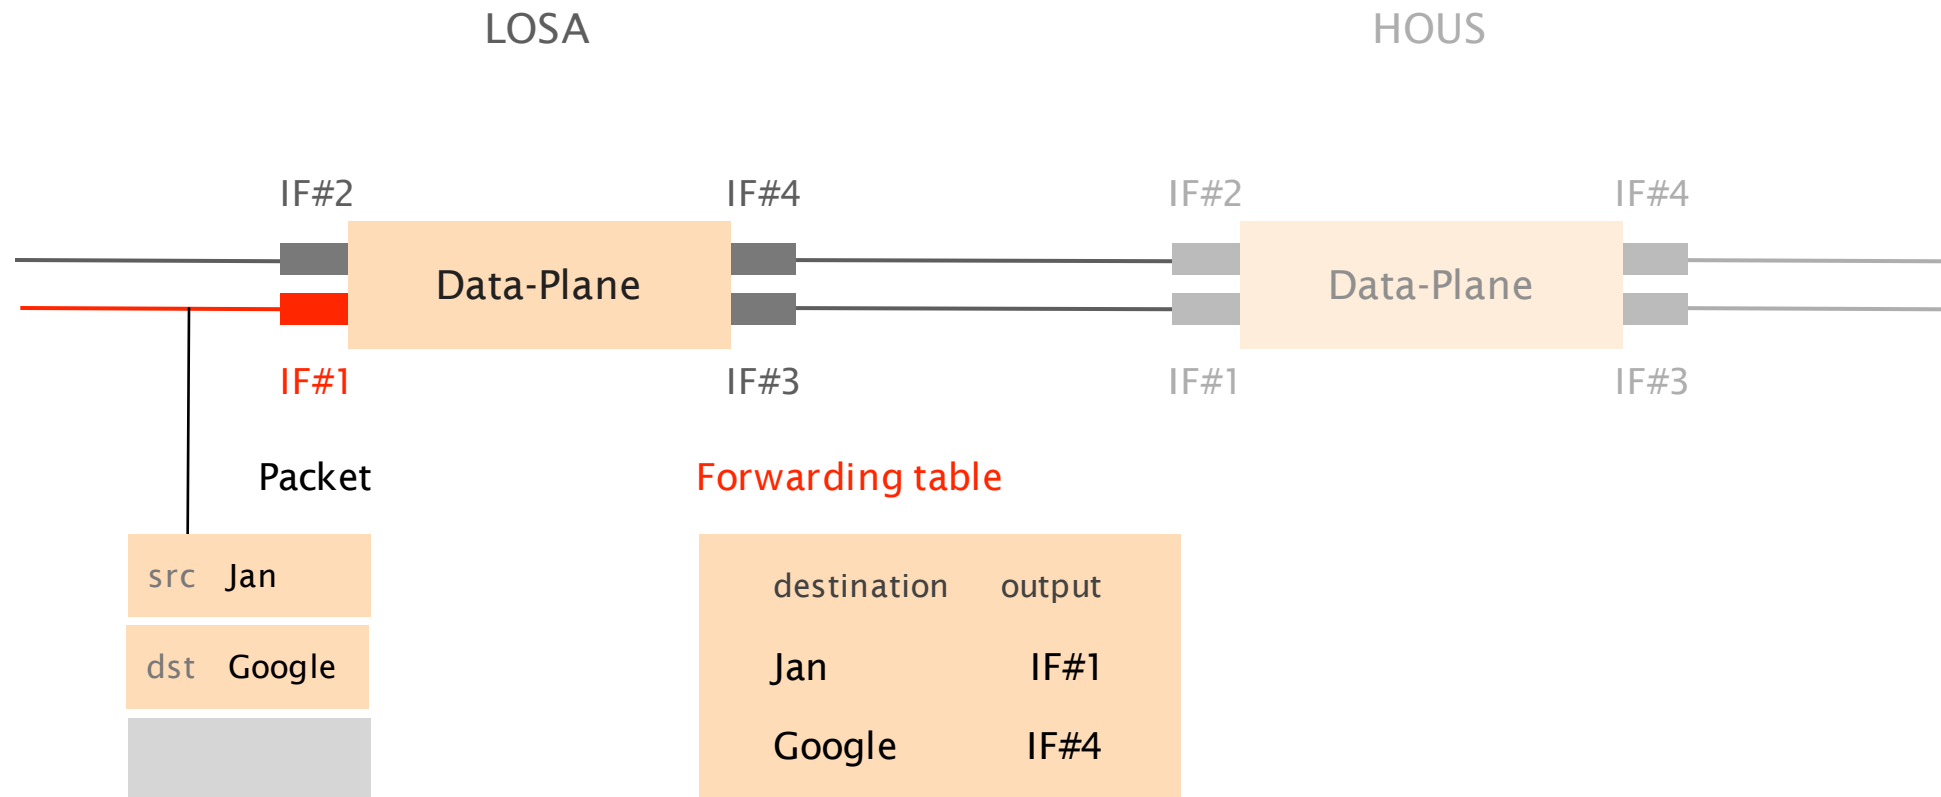


Let's zoom in on what is going on  
between two adjacent routers



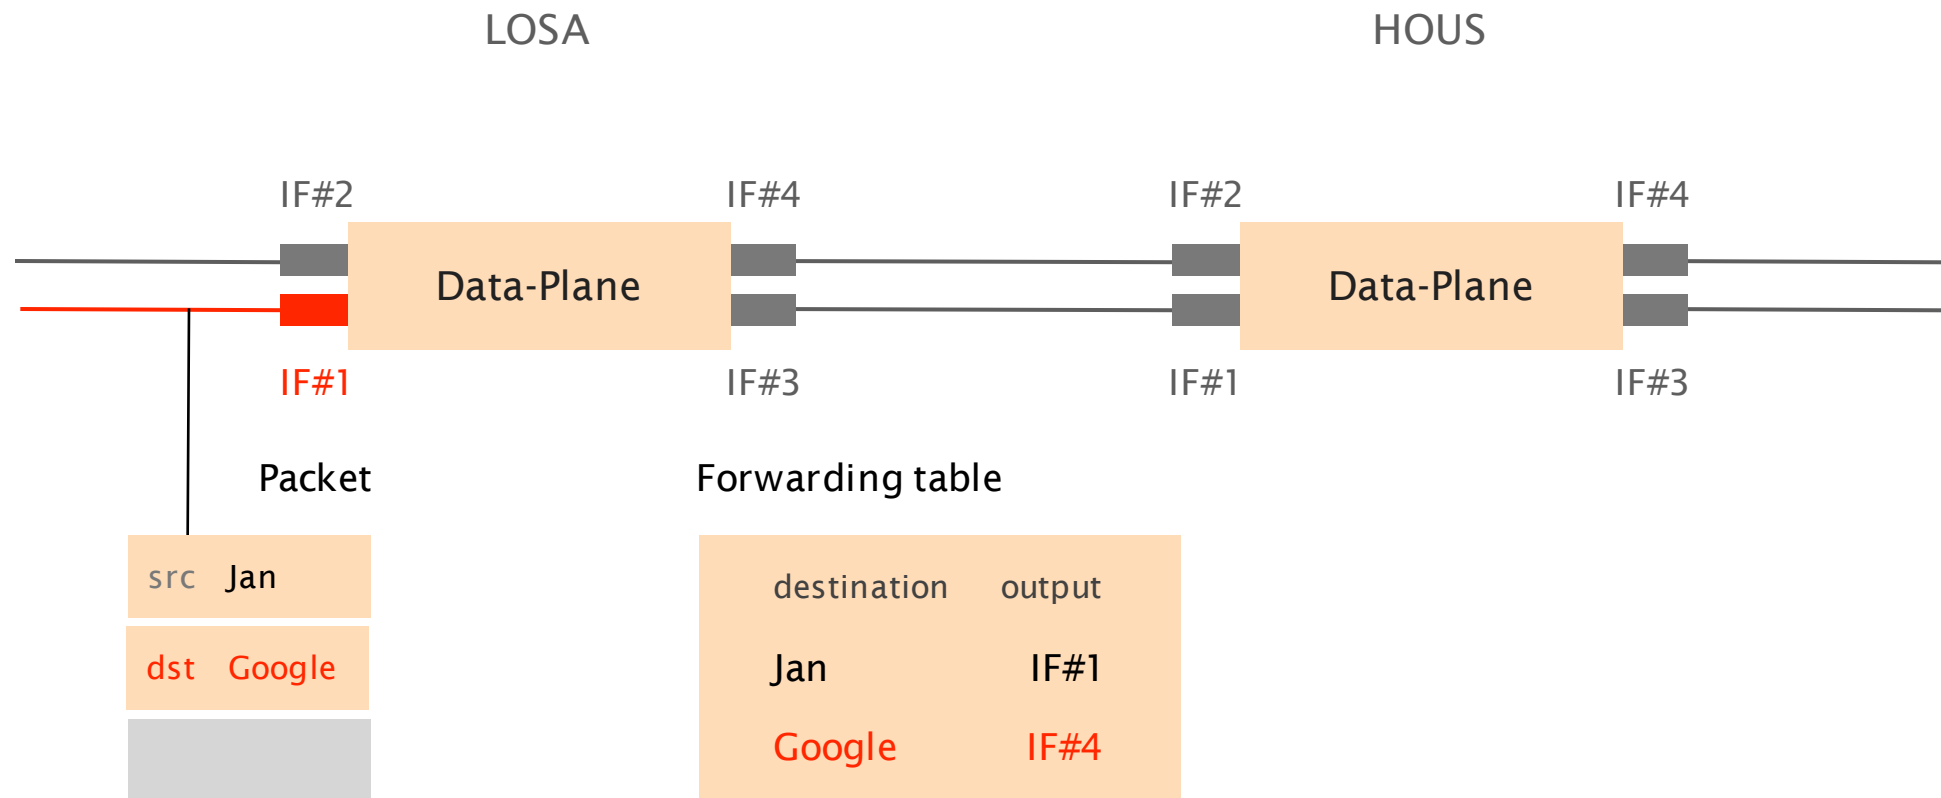


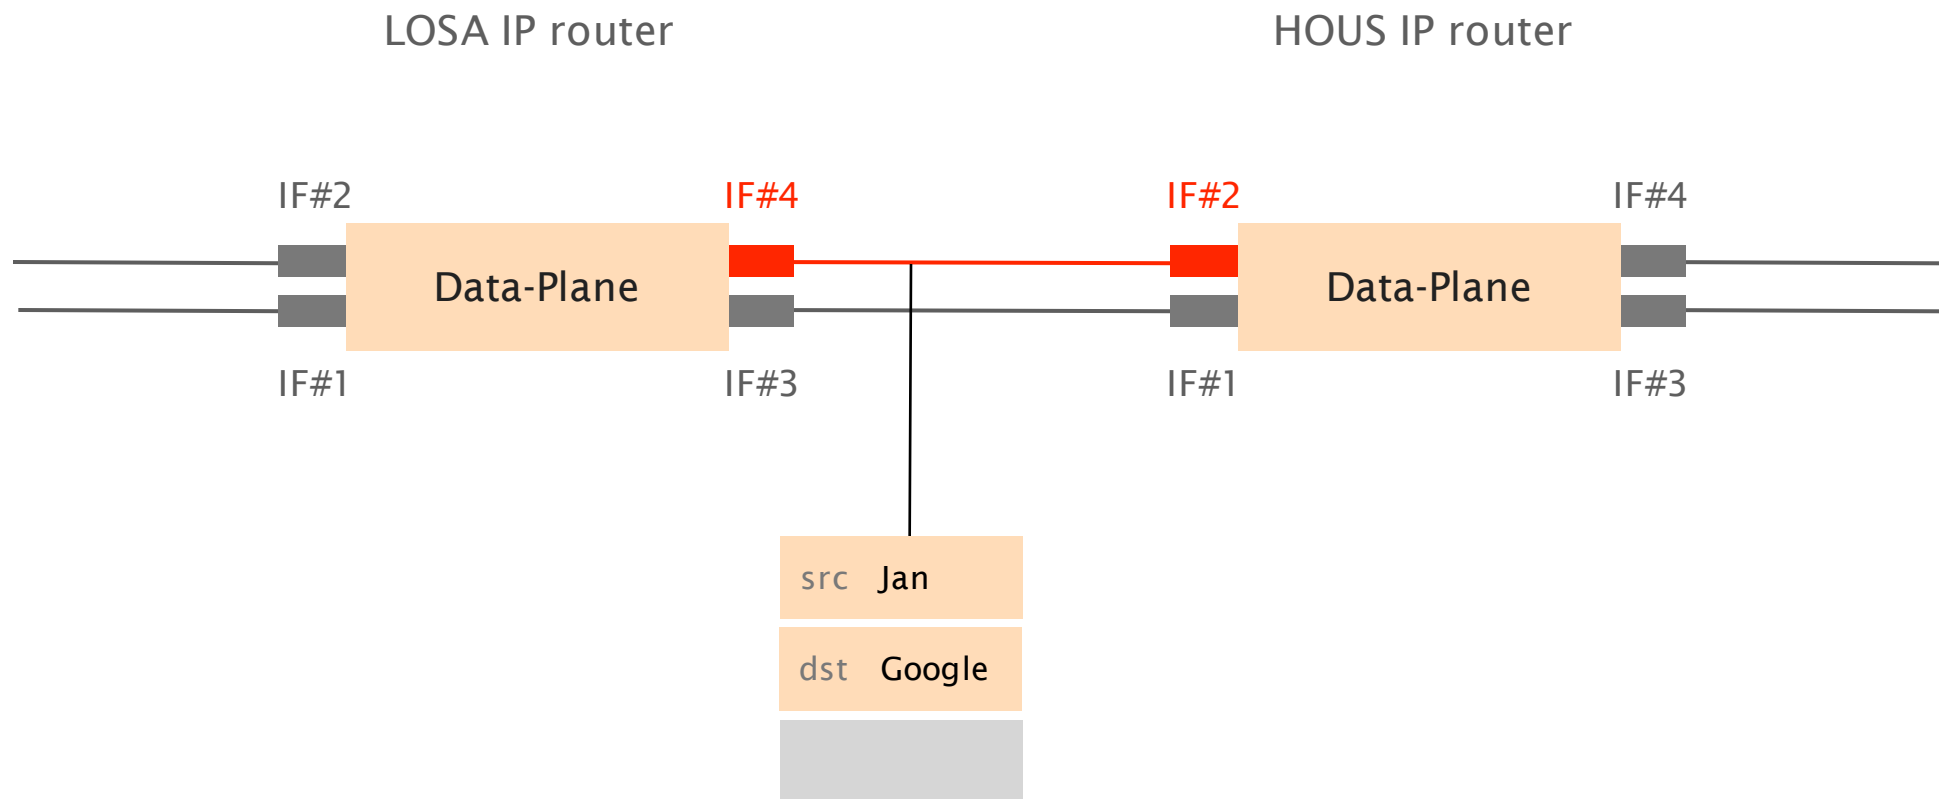
Upon packet reception, routers **locally** look up their forwarding table to know where to send it next



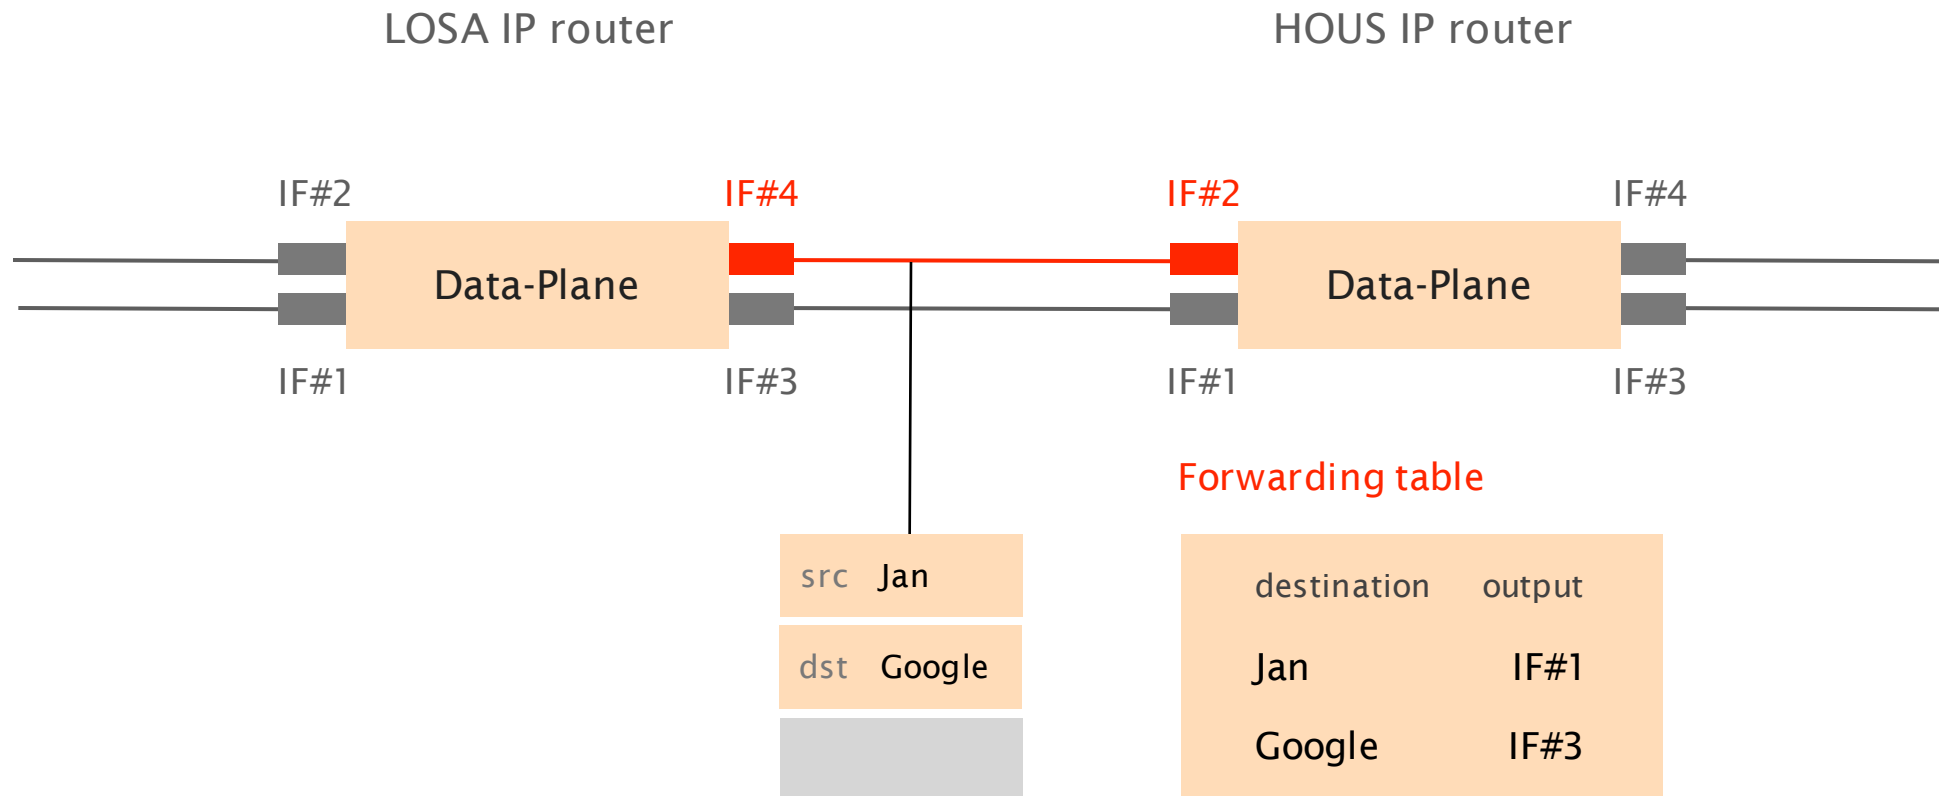


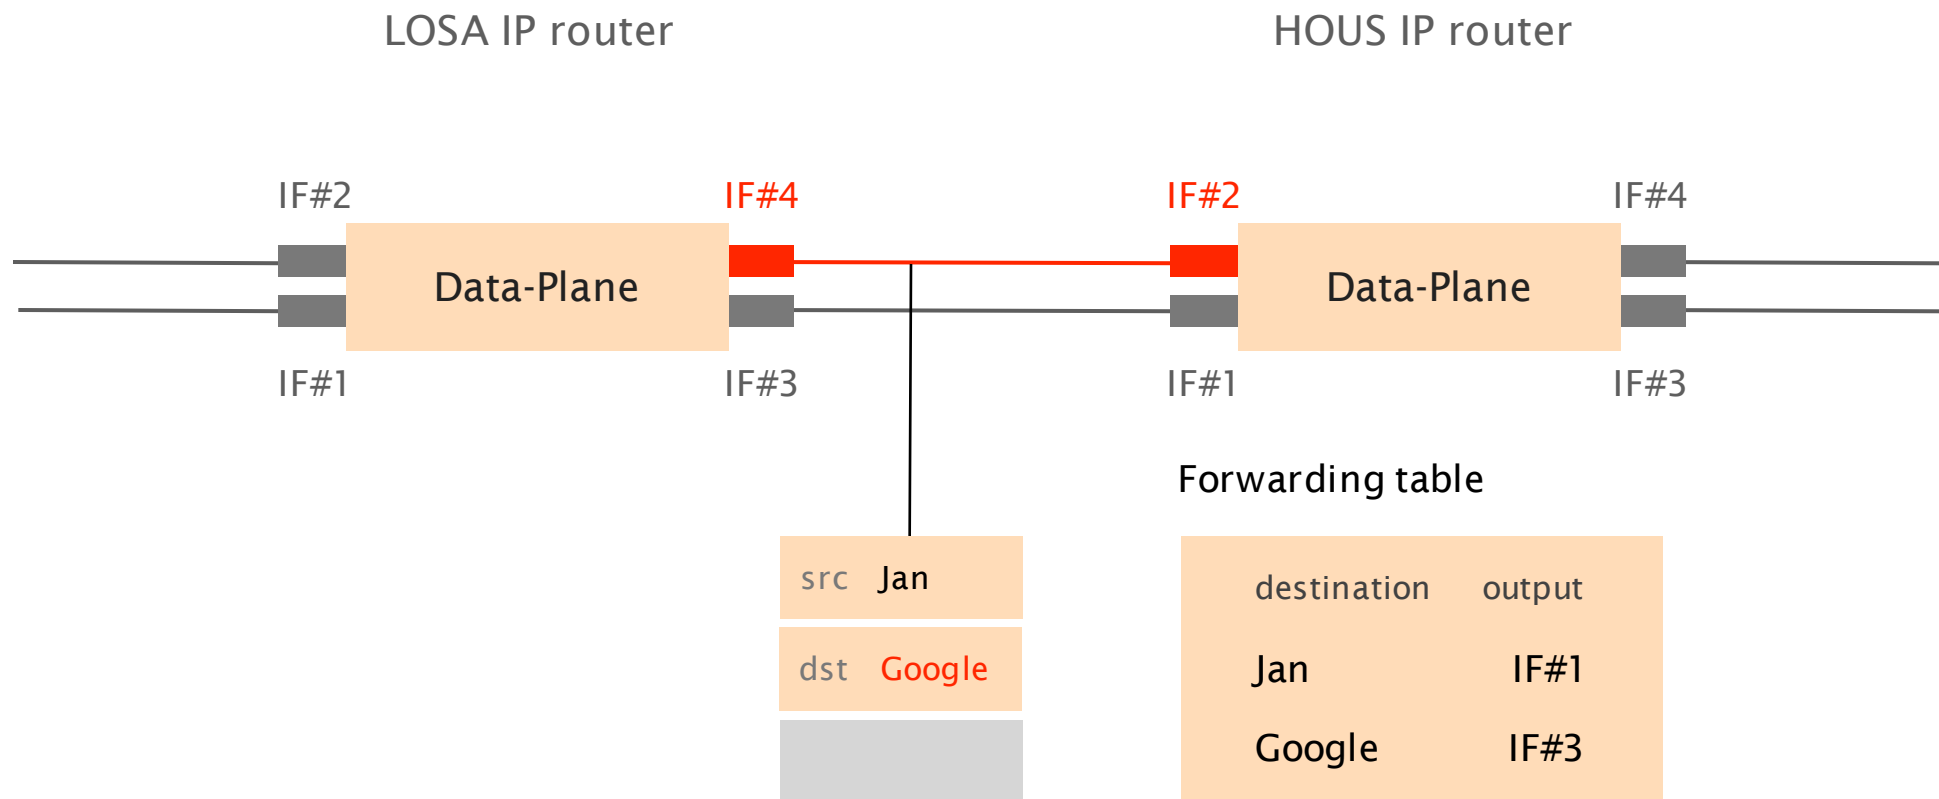
Here, the packet should be directed to **IF#4**

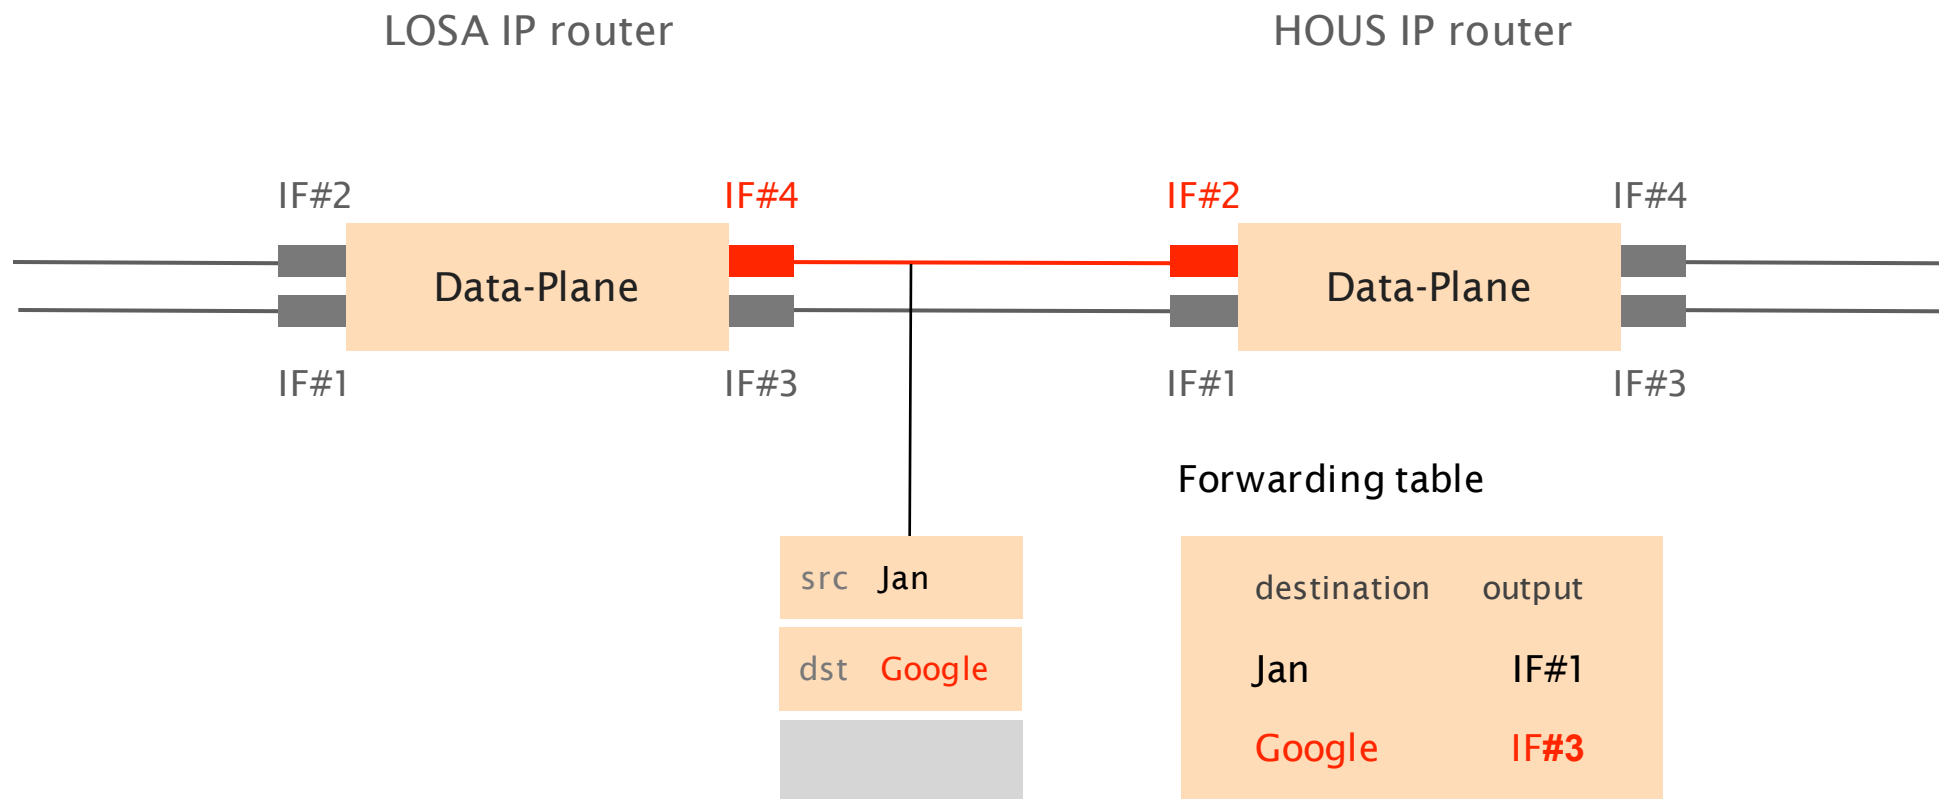


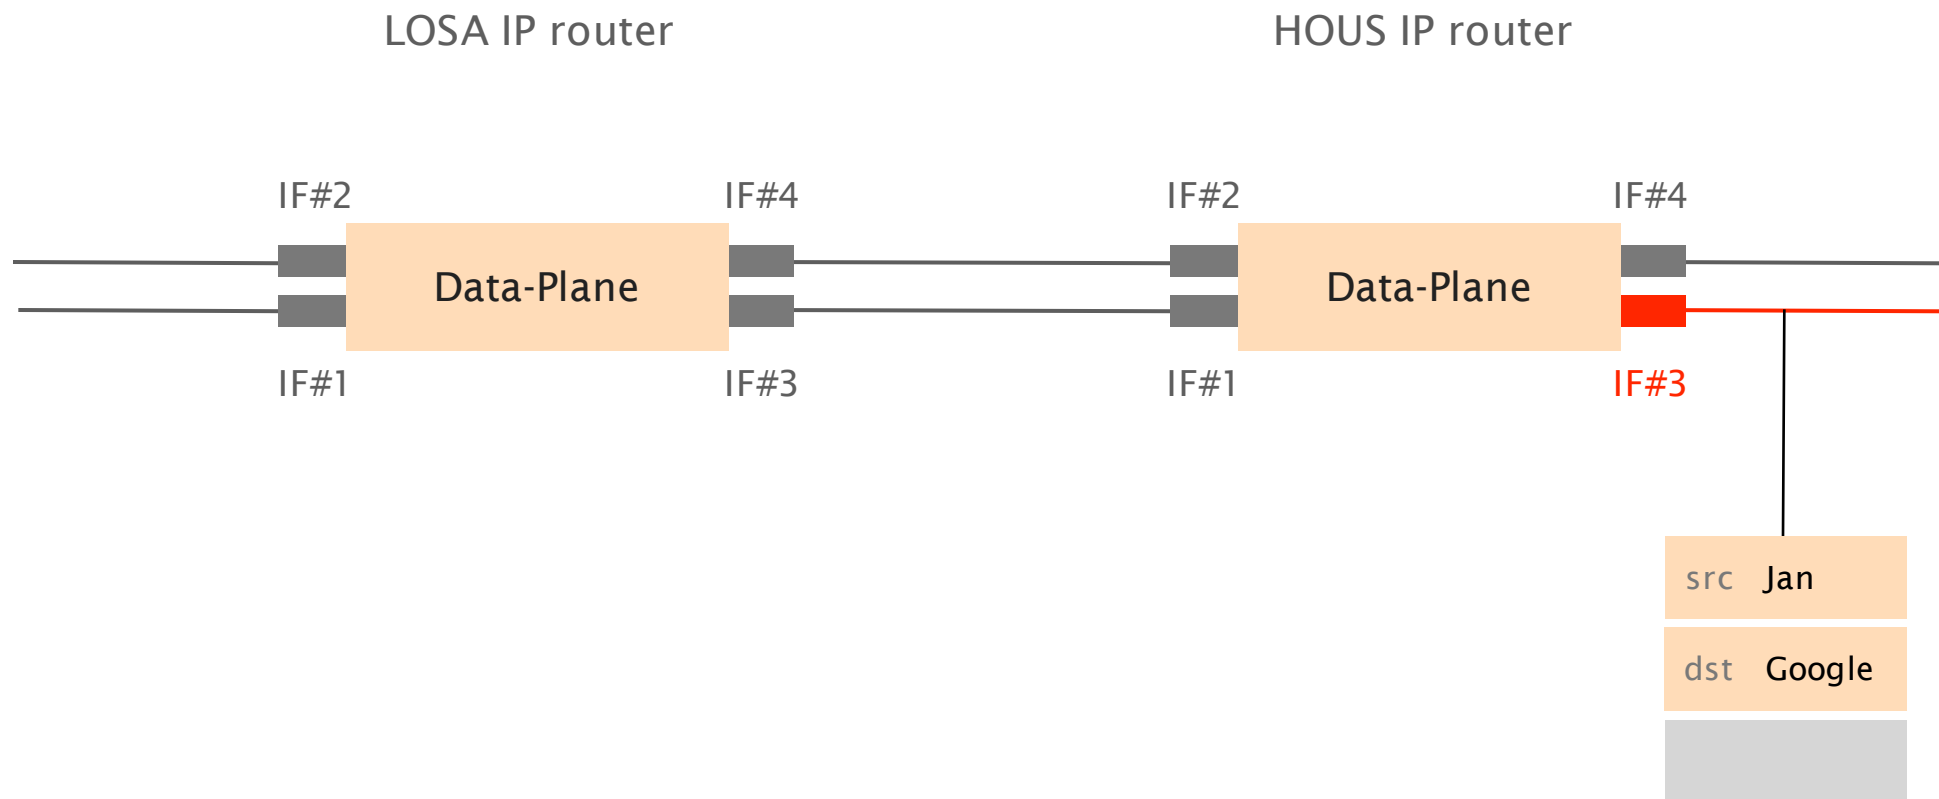


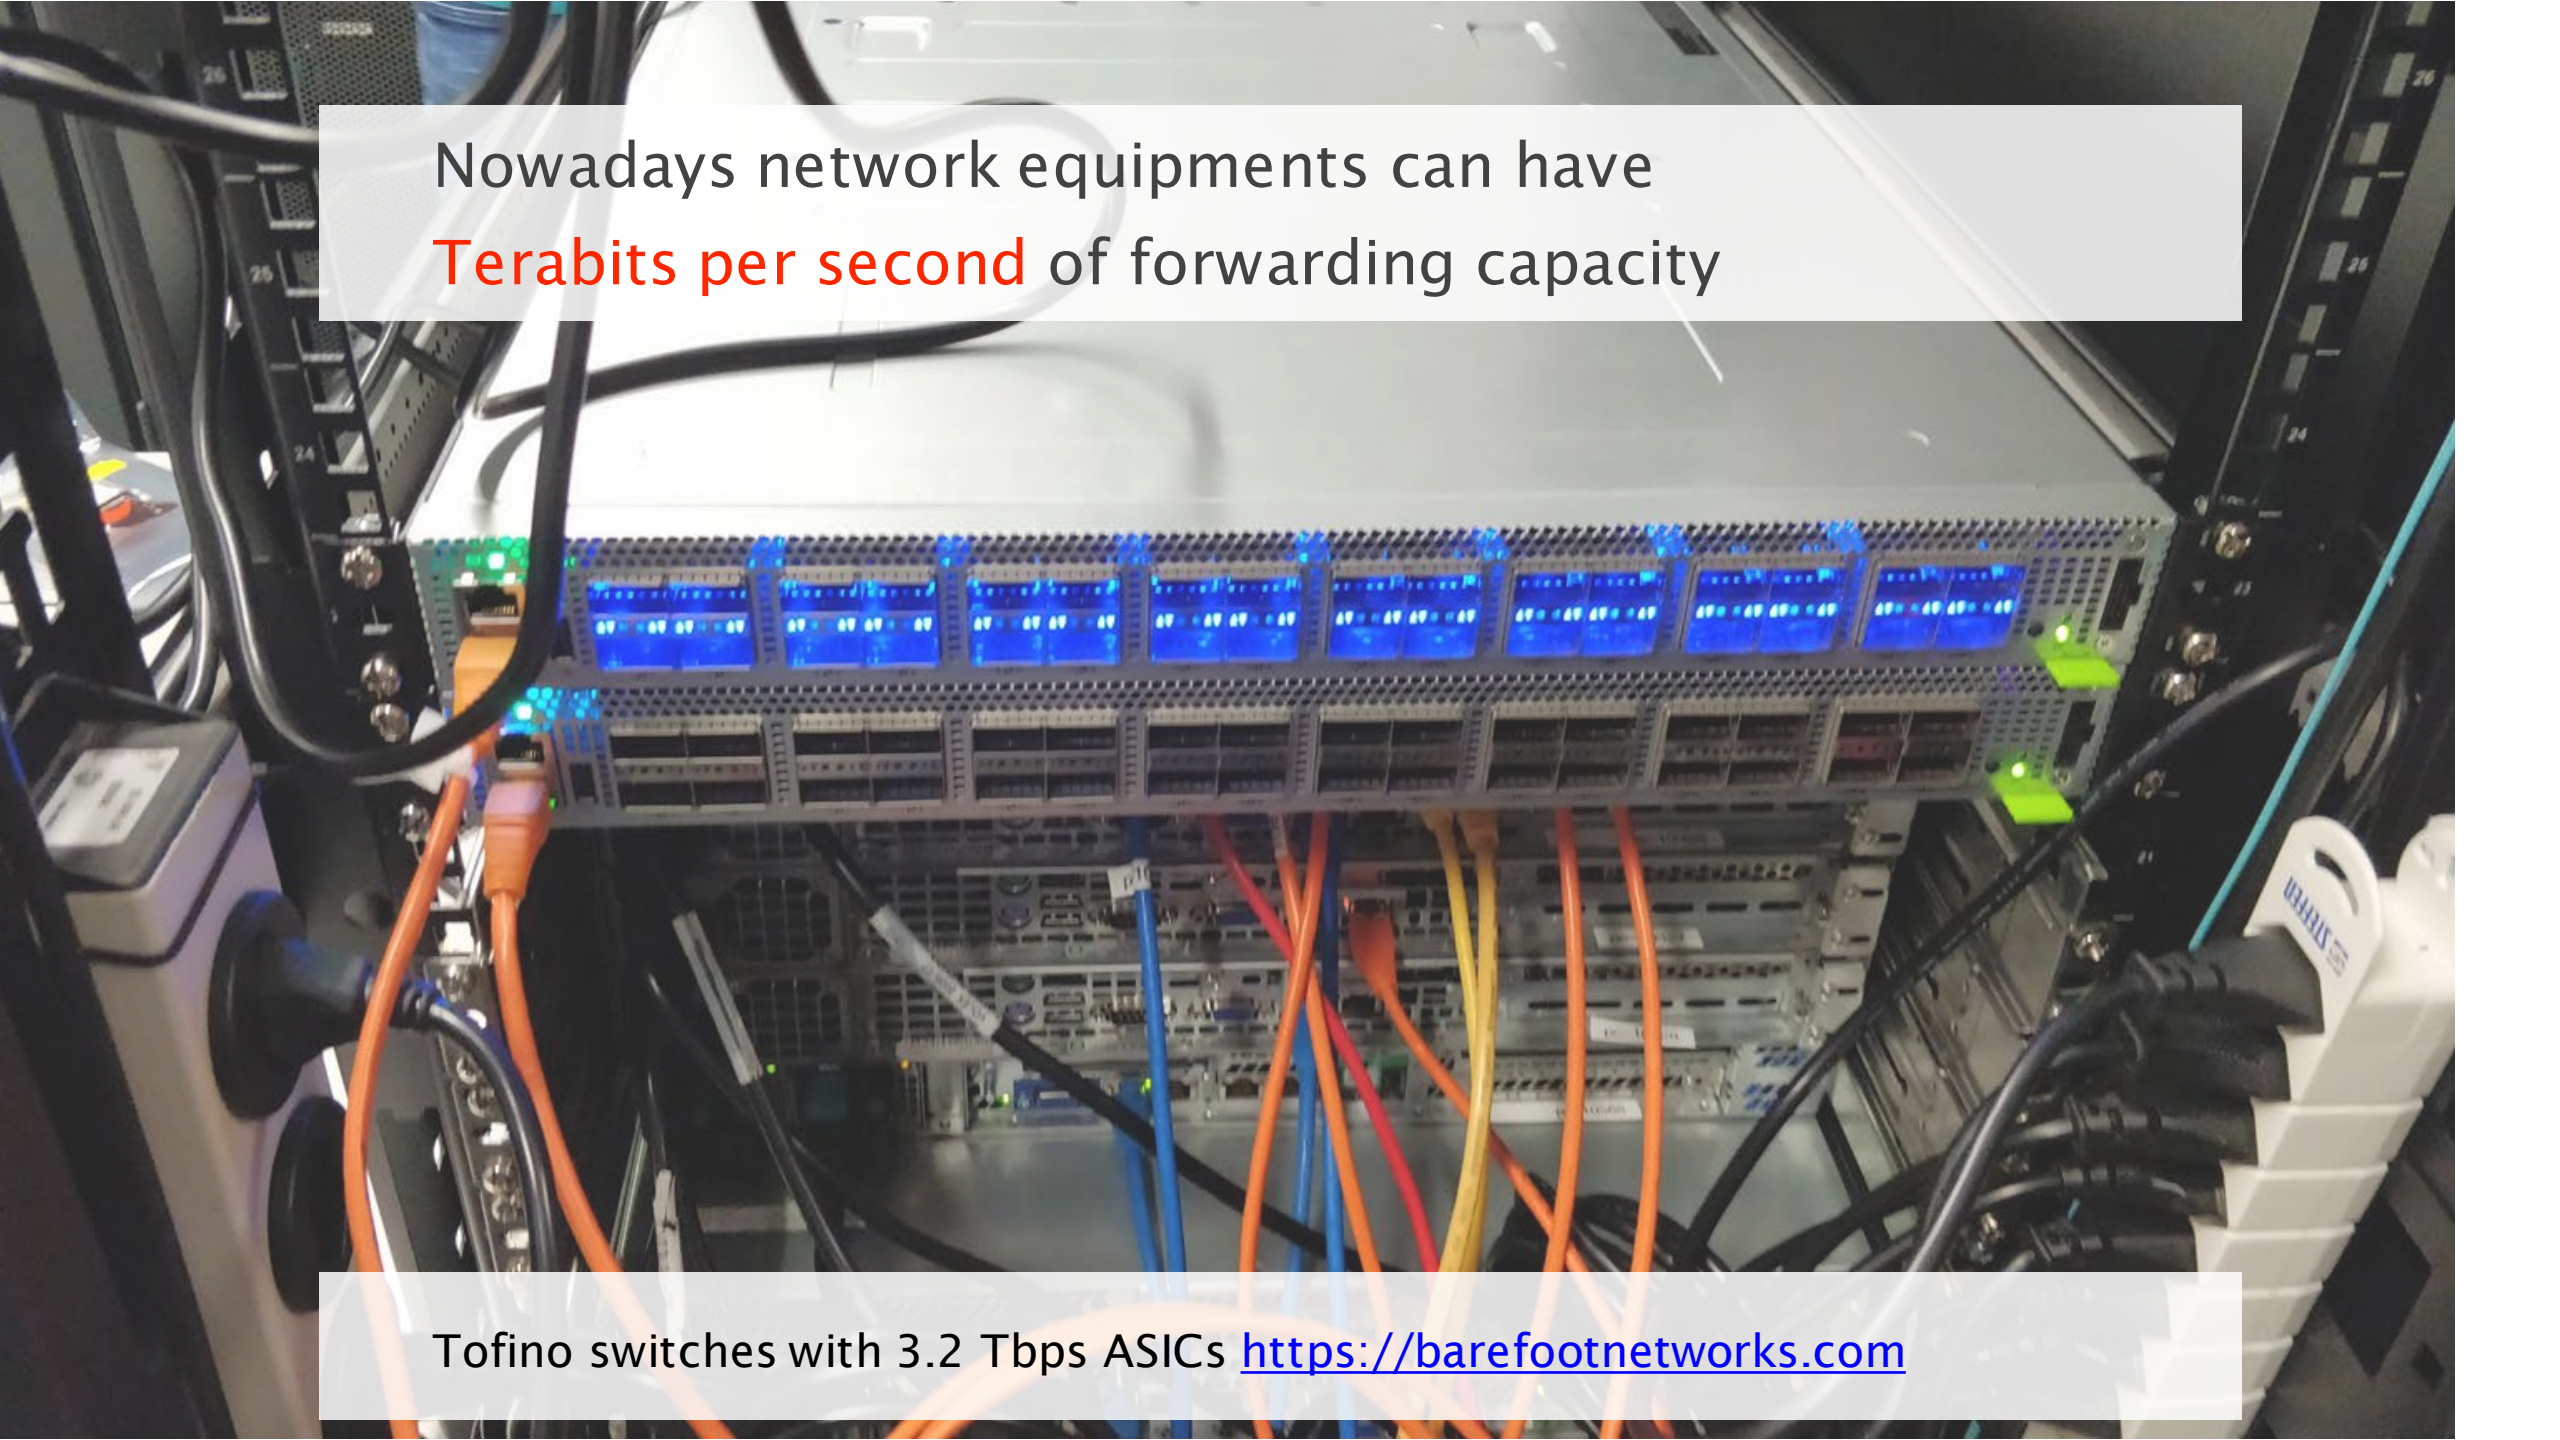
Forwarding is repeated at each router,  
until the destination is reached











Nowadays network equipments can have  
**Terabits per second** of forwarding capacity

Tofino switches with 3.2 Tbps ASICs <https://barefootnetworks.com>

Forwarding decisions necessarily depend on the destination, but can also depend on other criteria

criteria

destination

mandatory (why?)

source

requires  $n^2$  state


input port

traffic engineering


+any other header



With source- & destination-based routing,  
paths from different sources can differ



With destination-based routing,  
paths from different source coincide once they overlap




Once path to destination meet,  
they will *never* split

Set of paths to the destination  
produce a spanning tree rooted at the destination:

- cover every router exactly once
- only one outgoing arrow at each router

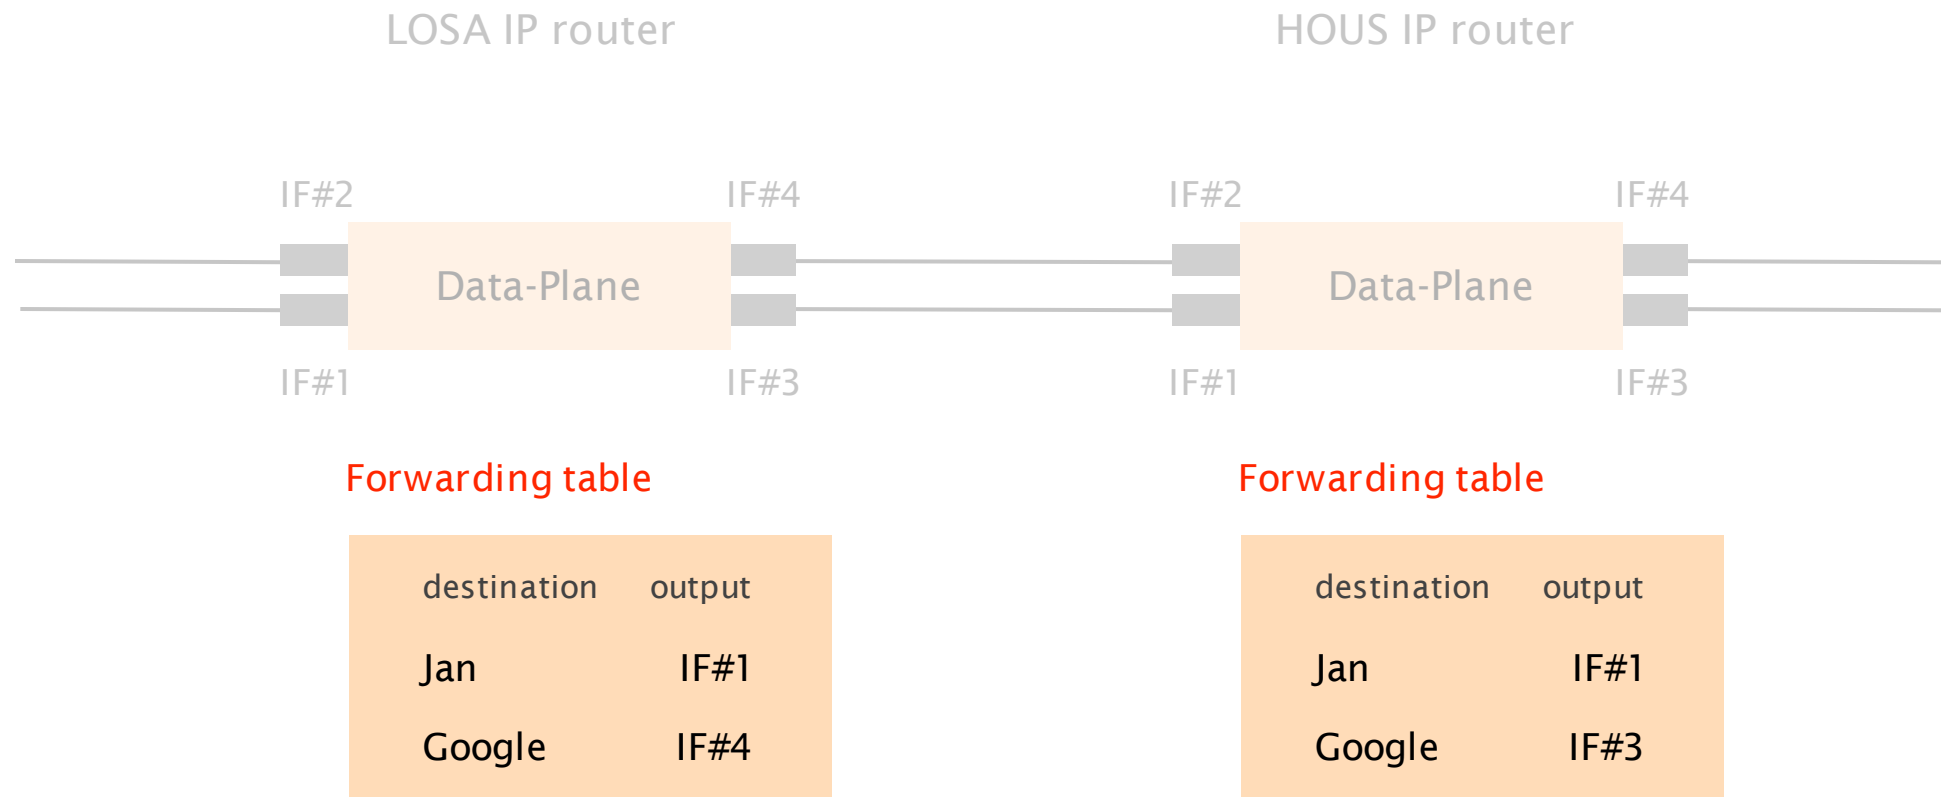
Here is an example of a spanning tree  
for destination  $X$




In the rest of the lecture,  
we'll consider destination-based routing

the default in the Internet

# Where are these forwarding tables coming from?





In addition to a data-plane,  
routers are also equipped with a control-plane





Think of the control-plane as the router's brain

Roles



Routing

Configuration

Statistics

...

**Routing** is the control-plane process that **computes** and **populates** the forwarding tables



While forwarding is a *local* process,  
routing is inherently a *global* process

How can a router know  
where to direct packets  
if it does not know what  
the network looks like?

# Forwarding vs Routing

## summary

	forwarding	routing
goal	directing packet to an outgoing link	computing the paths packets will follow
scope	local	network-wide
implem.	hardware usually	software usually
timescale	nanoseconds	milliseconds (hopefully)

The goal of routing is to compute  
valid global forwarding state

Definition      a global forwarding state is valid if  
  
it **always** delivers packets  
to the correct destination

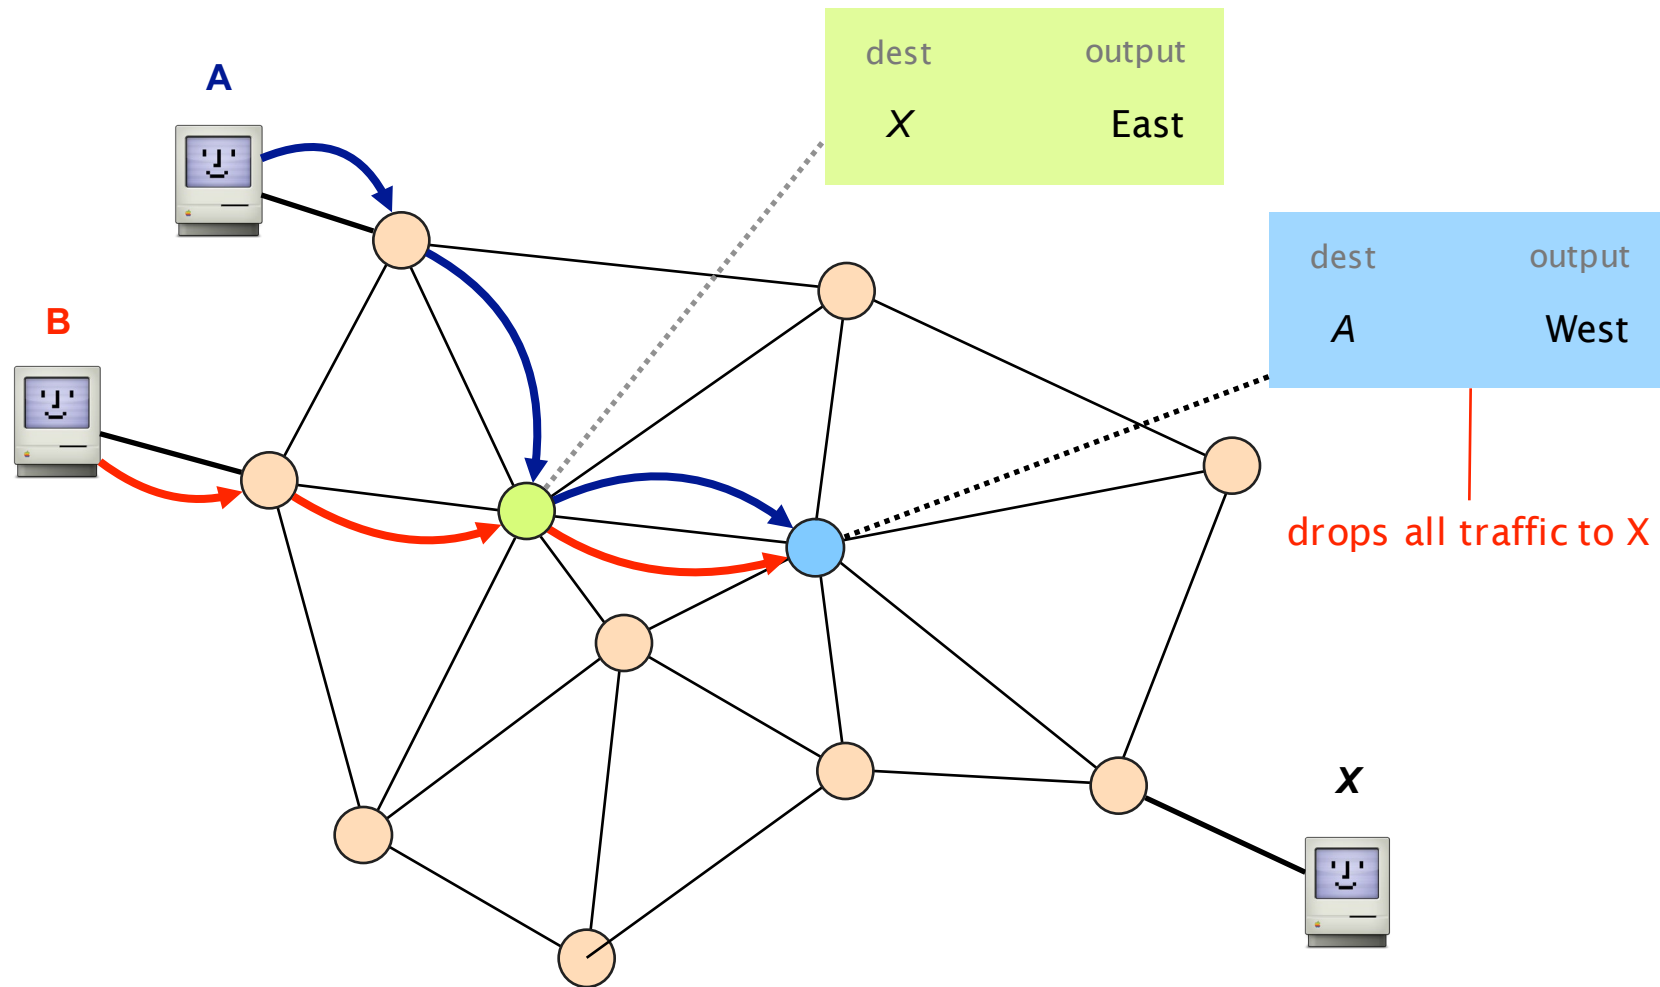
sufficient and necessary condition

Theorem

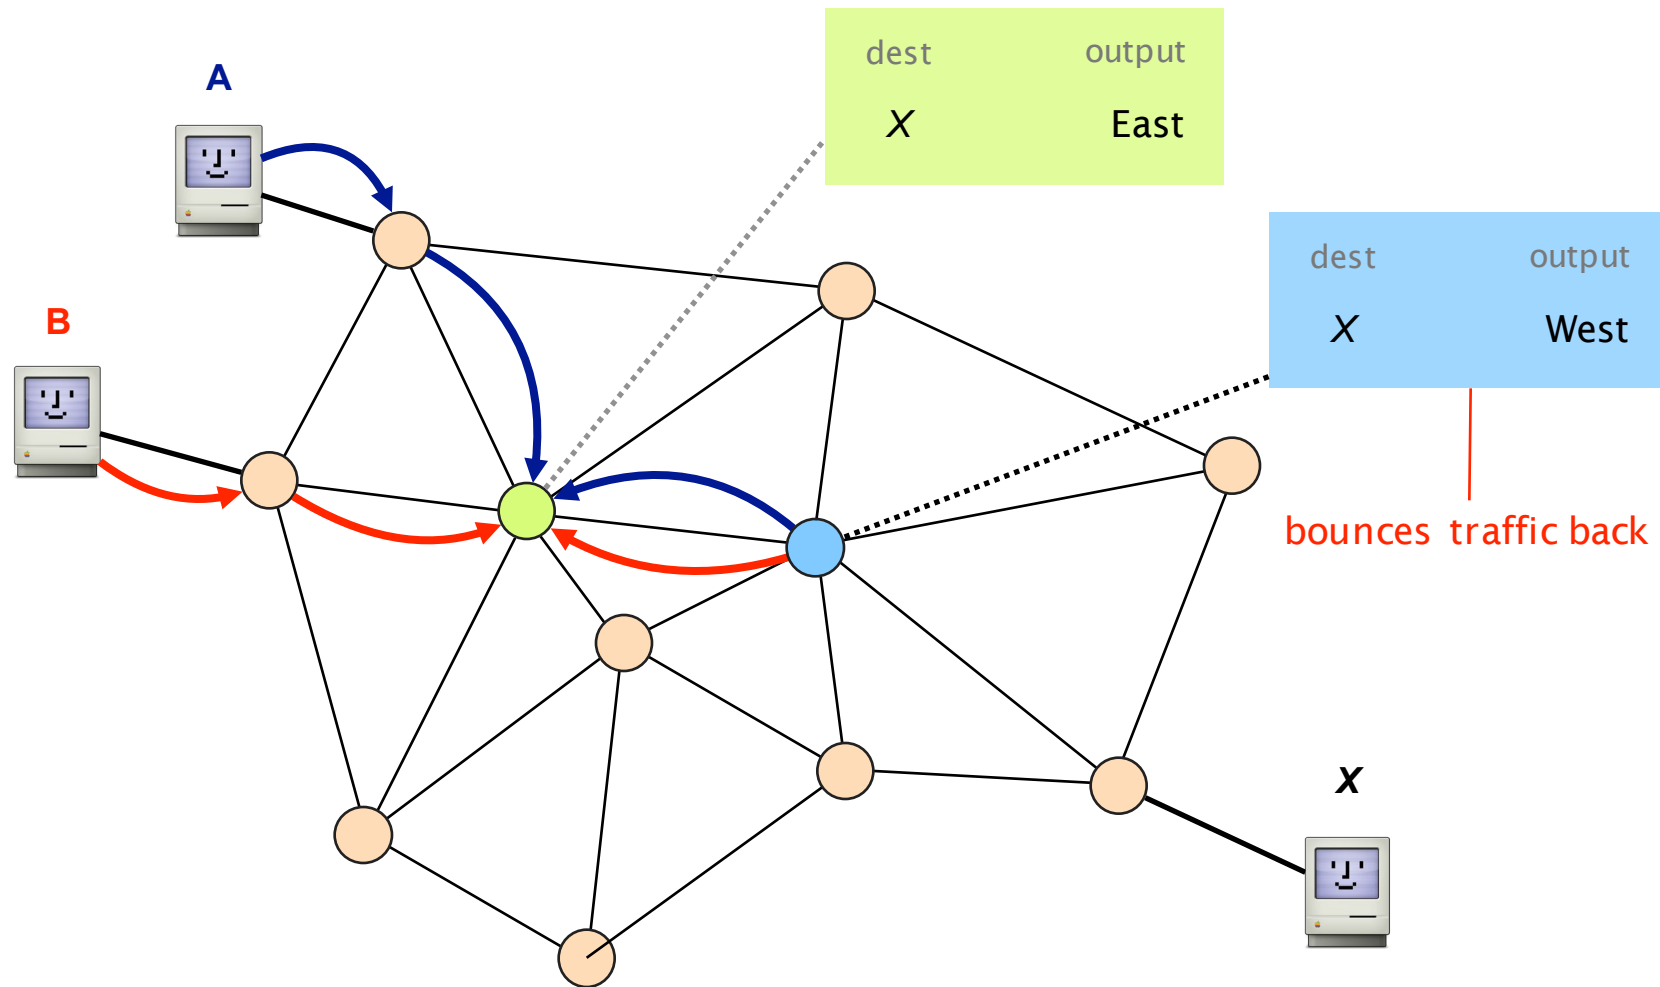
a global forwarding state is valid if and only if

- there are no dead ends  
no outgoing port defined in the table
- there are no loops  
packets going around the same set of nodes

A global forwarding state is valid if and only if there are **no dead ends**



A global forwarding state is valid if and only if there are **no forwarding loops**





sufficient and necessary condition

Theorem

a global forwarding state is valid if and only if

- there are no dead ends  
*i.e.* no outgoing port defined in the table
- there are no loops  
*i.e.* packets going around the same set of nodes

# Proving the necessary condition is easy

Theorem      If a routing state is valid  
then there are no loops or dead-end

Proof          If you run into a dead-end or a loop  
you'll never reach the destination  
so the state cannot be correct (contradiction)

## Proving the sufficient condition is more subtle

### Theorem

If a routing state has no dead end and no loop  
then it is valid

### Proof

There is only a finite number of ports to visit

A packet can never enter a switch via the same port,  
otherwise it is a loop (which does not exist by assumption )

As such, the packet must **eventually** reach the destination

question 1      How do we verify that a forwarding state is valid?

question 2      How do we compute valid forwarding state?

question 1

How do we verify that a forwarding state is valid?

How do we compute valid forwarding state?

# Verifying that a routing state is valid is easy

simple algorithm

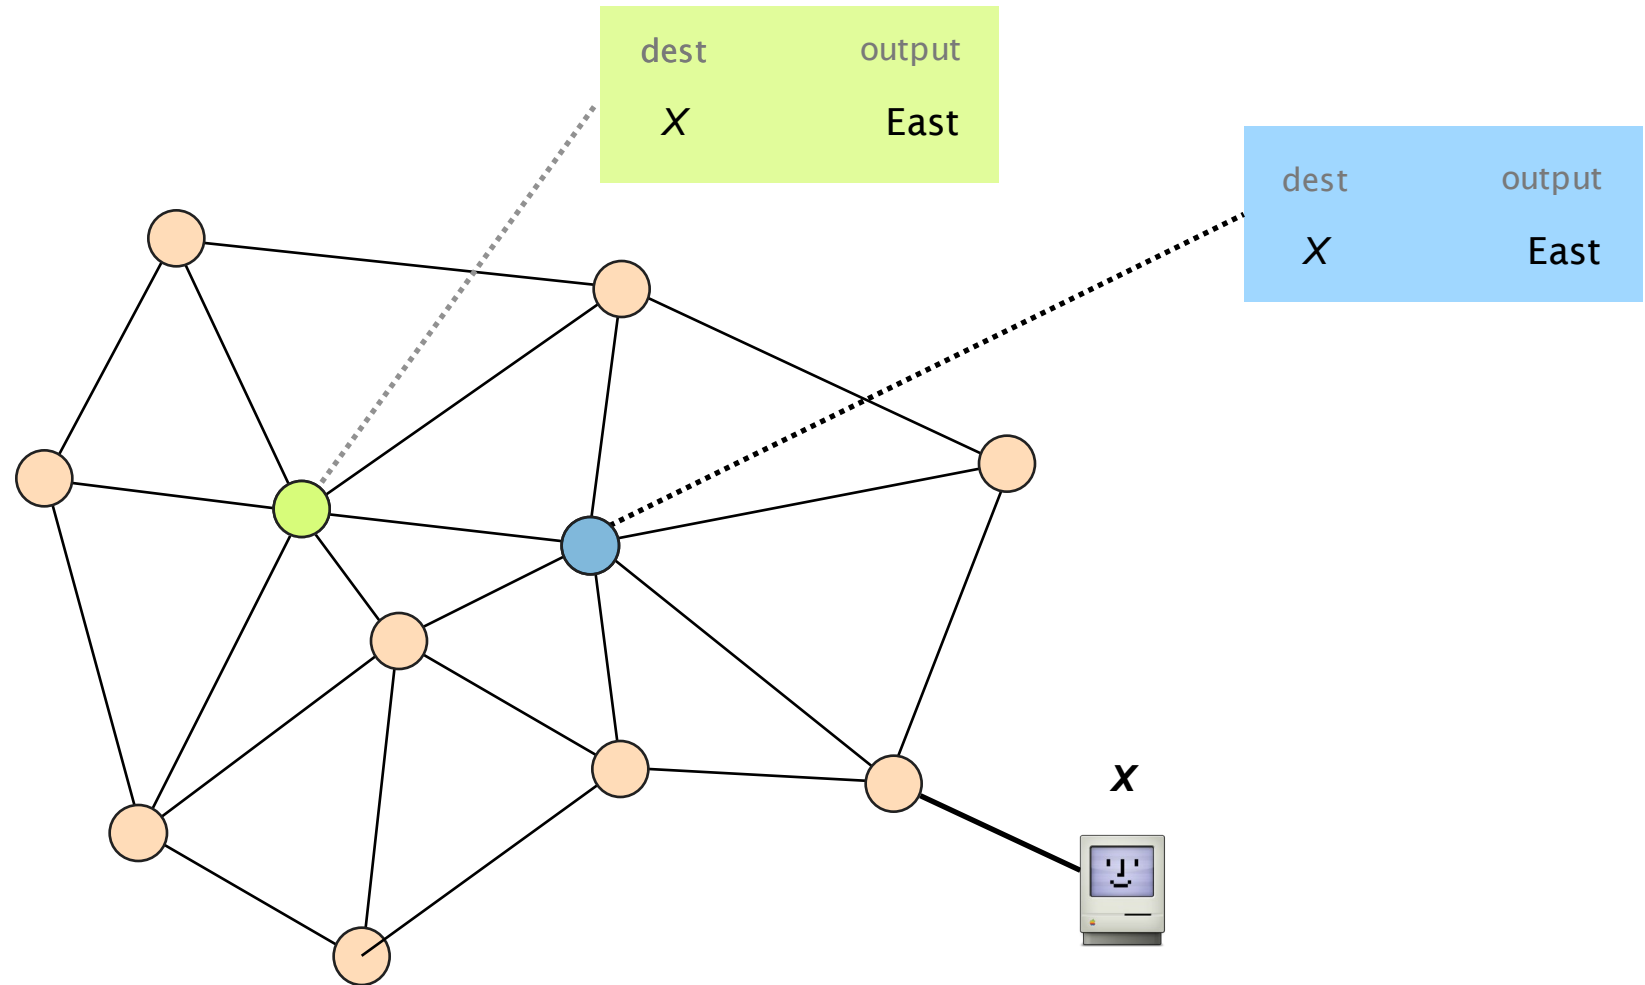
for one destination

Mark all outgoing ports with an arrow


Eliminate all links with no arrow

State is valid *iff* the remaining graph  
is a spanning-tree

Given a graph with the corresponding forwarding state





Mark all outgoing ports with an arrow






Eliminate all links with no arrow






The **result** is a spanning tree.


This is a **valid** routing state



Mark all outgoing ports with an arrow




Eliminate all links with no arrow



The result is **not a spanning-tree**.

The routing state is **not valid**



How do we verify that a forwarding state is valid?

question 2

How do we compute valid forwarding state?

# Producing valid routing state is harder

prevent dead ends  
easy

prevent loops  
hard



Producing valid routing state is harder  
**but doable**

prevent dead ends  
easy

prevent loops  
**hard**

This is the question  
you should focus on

Existing routing protocols differ in  
how they avoid loops

prevent loops

hard

Essentially,  
there are three ways to compute valid routing state

	Intuition	Example
#1	Use tree-like topologies	Spanning-tree
#2	Rely on a global network view	Link-State SDN
#3	Rely on distributed computation	Distance-Vector BGP

Communication Networks and Internet Technology

**Short Recap on this weeks lecture**

# Communication Networks and Internet Technology

## Part 1: General overview

What is a network made of?

How is it shared?

How is it organized?

How does communication happen?

How do we characterize it?

# Communication Networks and Internet Technology

## Part 2: Concepts



routing

How do you guide IP packets  
from a source to destination?

reliable  
delivery

How do you ensure reliable transport  
on top of best-effort delivery?

# Reading: Book Kurose & Ross


Class textbook:

*Computer Networking: A Top-Down Approach (8<sup>th</sup> ed.)*

J.F. Kurose, K.W. Ross

Pearson, 2020

[http://gaia.cs.umass.edu/kurose\\_ross](http://gaia.cs.umass.edu/kurose_ross)



- Week 01
  - 1.1 (The Internet), 1.2 (The Network Edge), 1.3 (The Network Core) and 1.5 (Protocol Layers)
- Week 02
  - 1.4 (Delay, Loss and Throughput), 1.5 (Protocol Layers), 4.2 (What's Inside a Router)


# Check Your Knowledge

[PROBLEM SOLVING HOME](#)[TRY A RANDOM PROBLEM](#)

## INTERACTIVE END-OF-CHAPTER EXERCISES

Supplement to Computer Networking: A Top Down Approach 8th Edition

*"Tell me and I forget. Show me and I remember. Involve me and I understand." Chinese proverb*



The links below will take you to end-of-chapter exercises where you'll be presented with an exercise whose solution can then be displayed (hopefully *after* you've solved the exercise yourself!). Each of the exercises below is similar to an end-of-chapter problem in the text. Most importantly, you can keep generating new instances of each exercise (and hopefully solving each one!) until you've mastered the material.

You may be interested in other supplemental material (online lectures, powerpoint slides, review questions, Wireshark labs) for our book, available [here](#).

This page replaces the earlier interactive problems page, and includes a number of new problems. We're actively adding new problems here in the summer of 2020. If you've got any comments or suggestions - let us know at [kurose@cs.umass.edu](mailto:kurose@cs.umass.edu)

### CHAPTER 1: INTRODUCTION

- Circuit Switching
- Quantitative Comparison of Packet Switching and Circuit Switching (similar to Chapter 1, P8, P9)
- Car - Caravan Analogy
- One-hop Transmission Delay (similar to example on pg. 37)
- Queuing Delay
- End-to-End Delay (similar to Chapter 1, P10)
- End-to-End Throughput (similar to Chapter 1, P20, and Figure 1.20)
- The IP Stack and Protocol Layering

[http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)



# Moving Ahead or Hearing Topics from a Different Voice

- Phil Levis and Nick McKeown  
Stanford


[CS144 Introduction to Computer Networking Fall 2016 on YouTube](#)



- Jim Kurose – UMass

[http://gaia.cs.umass.edu/kurose\\_ross/videos/1/](http://gaia.cs.umass.edu/kurose_ross/videos/1/)

Class textbook:  
*Computer Networking: A Top-Down Approach (8<sup>th</sup> ed.)*  
J.F. Kurose, K.W. Ross  
Pearson, 2020  
[http://gaia.cs.umass.edu/kurose\\_ross](http://gaia.cs.umass.edu/kurose_ross)



No lecture tomorrow 13:15-14:00h